

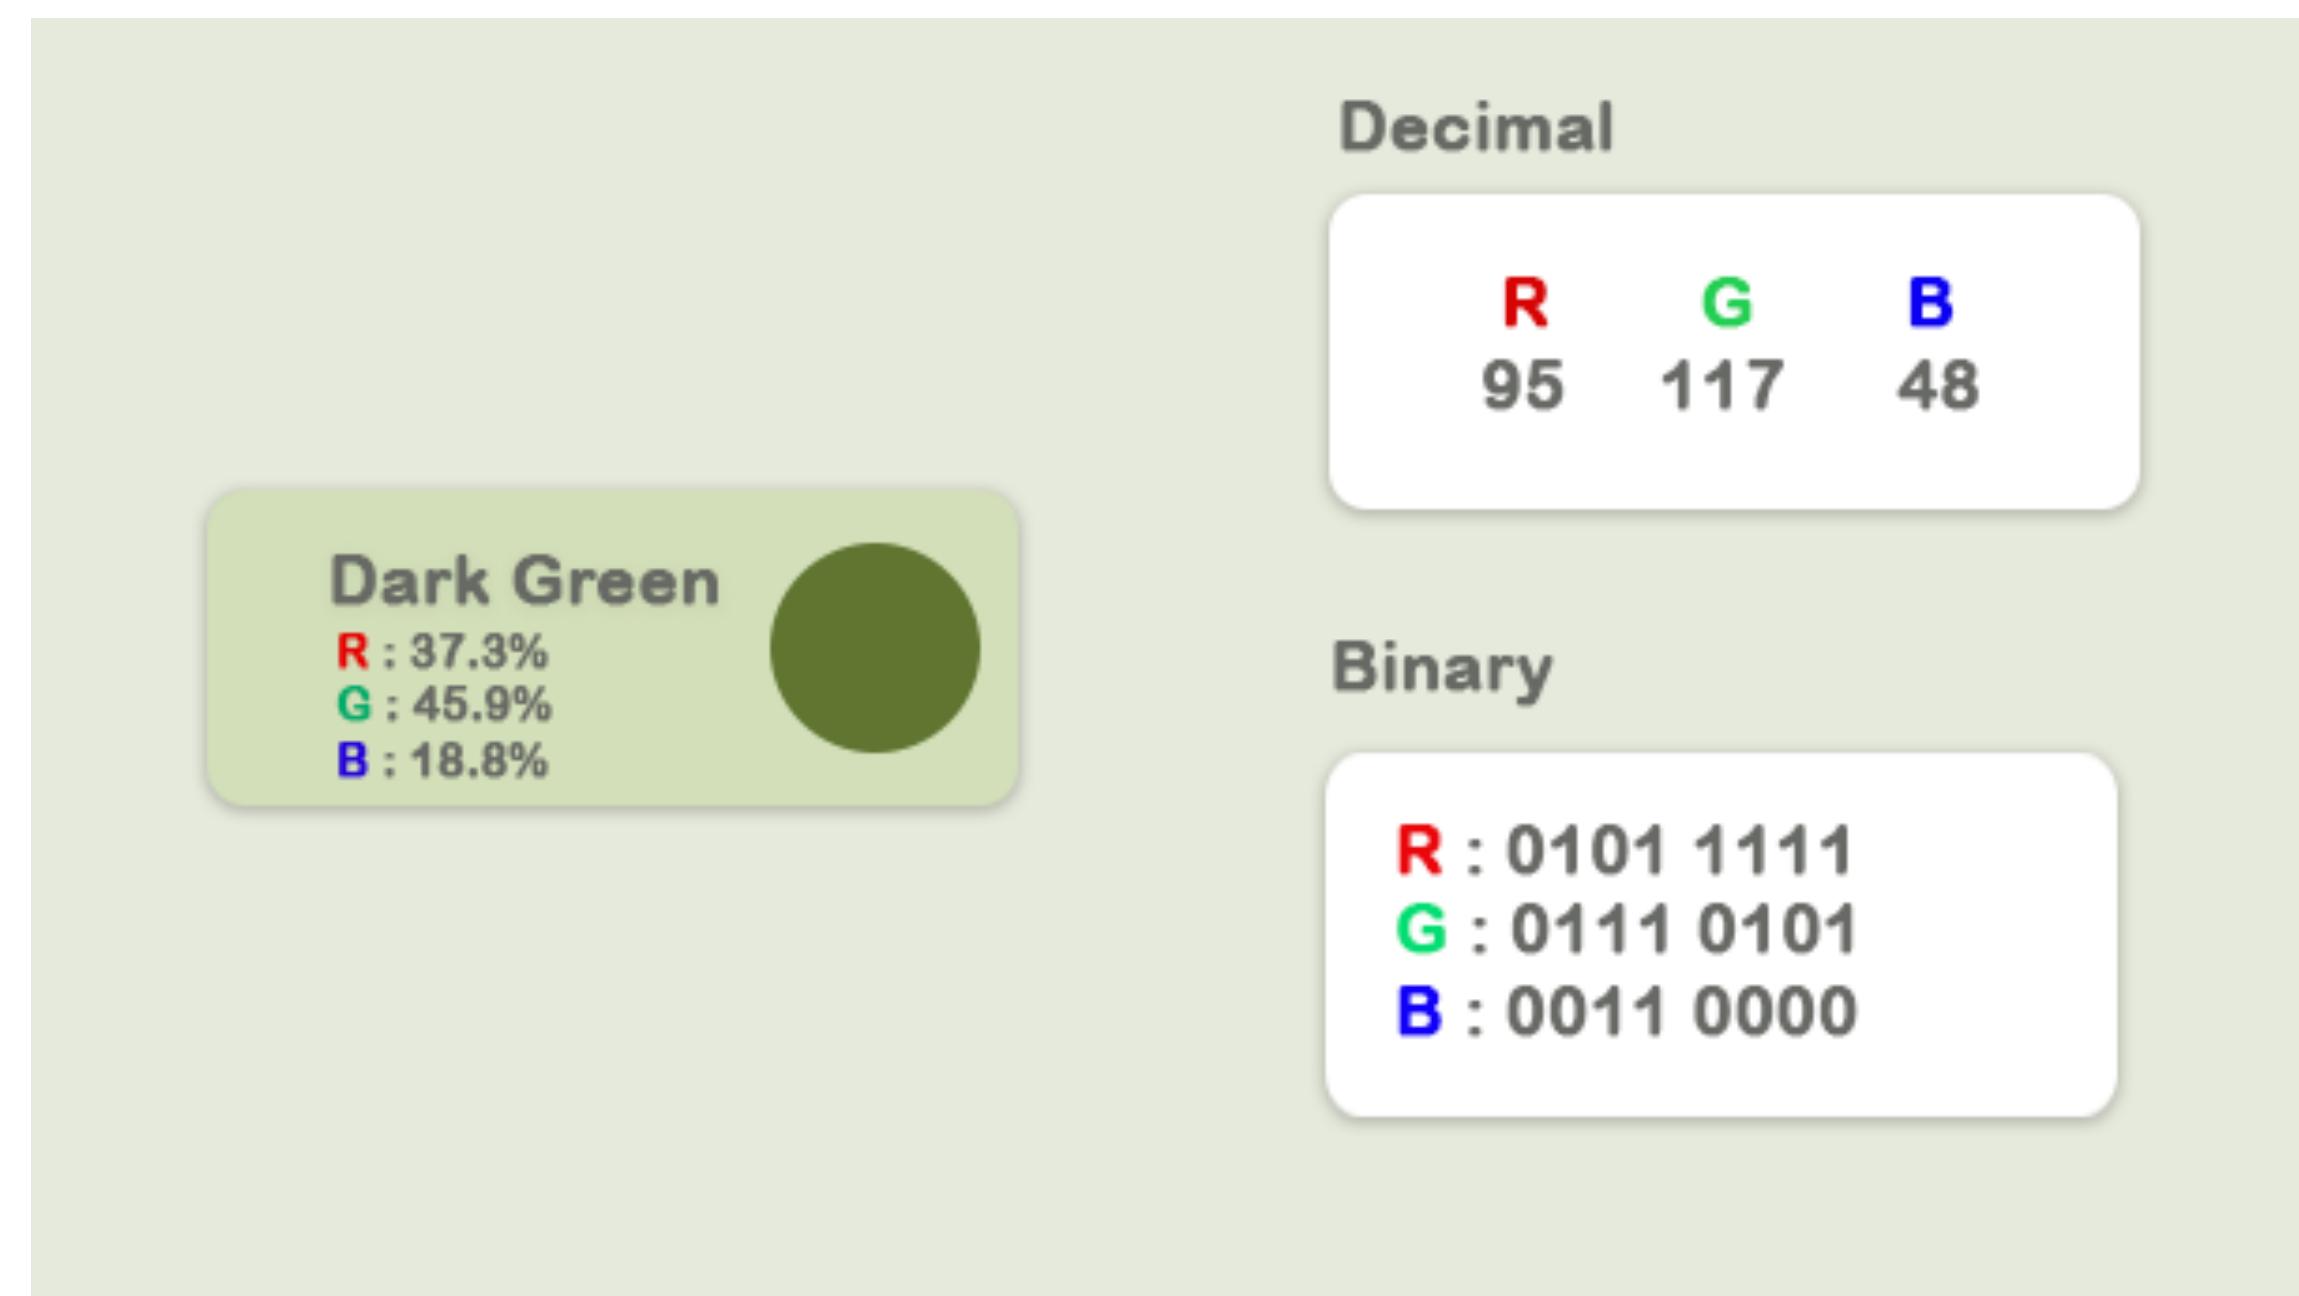
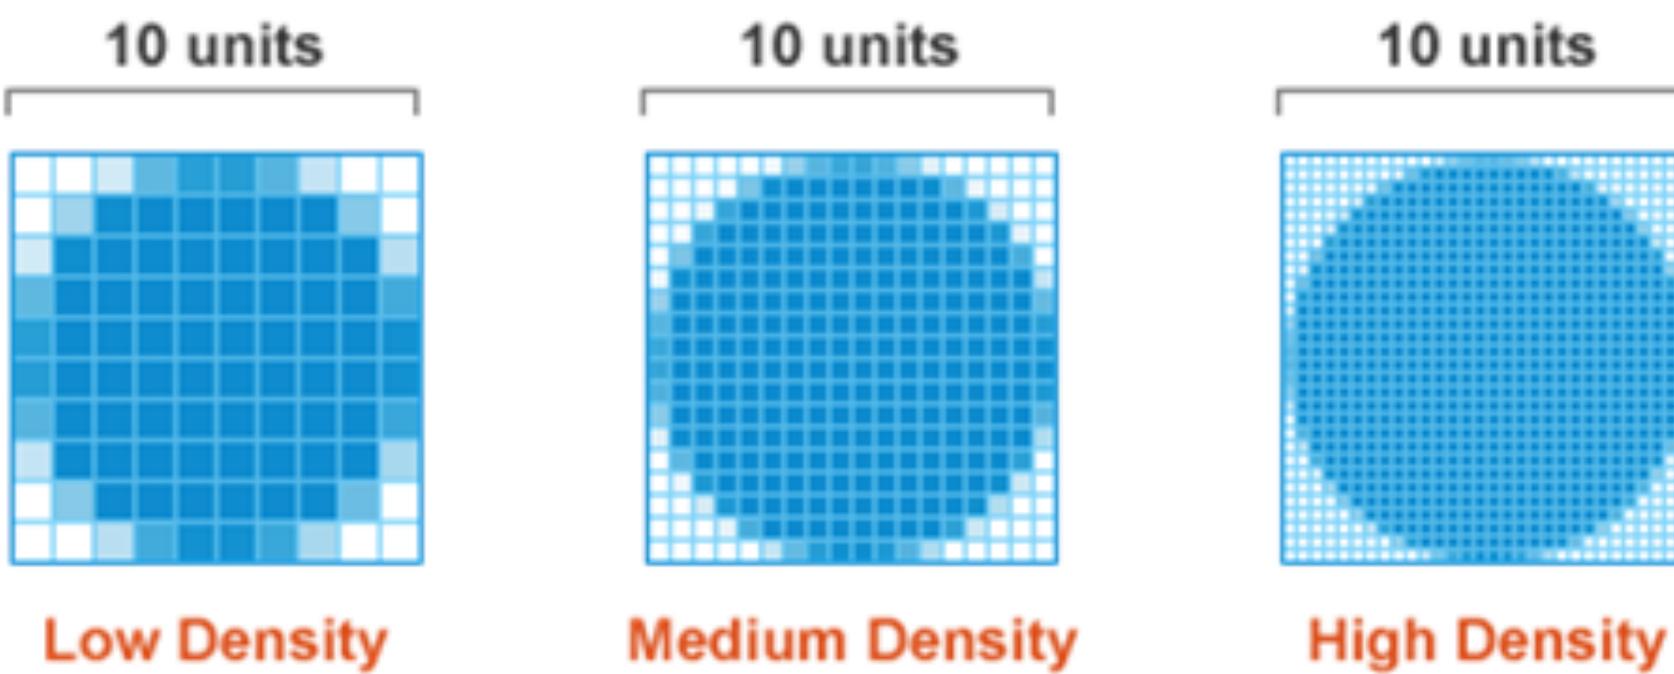
Lecture 12: Image Processing

ENAE 380 Flight Software Systems
October 21, 2024

Examples

- Basic Concepts
- OpenCV Basics
- Optimal Mark Recognition
- Autostereogram
- Automatic segmentation

Pixels and Color



How to grayscale

- Get the red, green, and blue values of a pixel
- Use fancy math to turn those numbers into a single gray value
- Replace the original red, green, and blue values with the new gray value

```
Gray = (Red + Green + Blue) / 3
```

```
For Each Pixel in Image {  
  
    Red = Pixel.Red  
    Green = Pixel.Green  
    Blue = Pixel.Blue  
  
    Gray = (Red + Green + Blue) / 3  
  
    Pixel.Red = Gray  
    Pixel.Green = Gray  
    Pixel.Blue = Gray  
  
}
```

Method 1: Averaging

Gray = (Red + Green + Blue) / 3

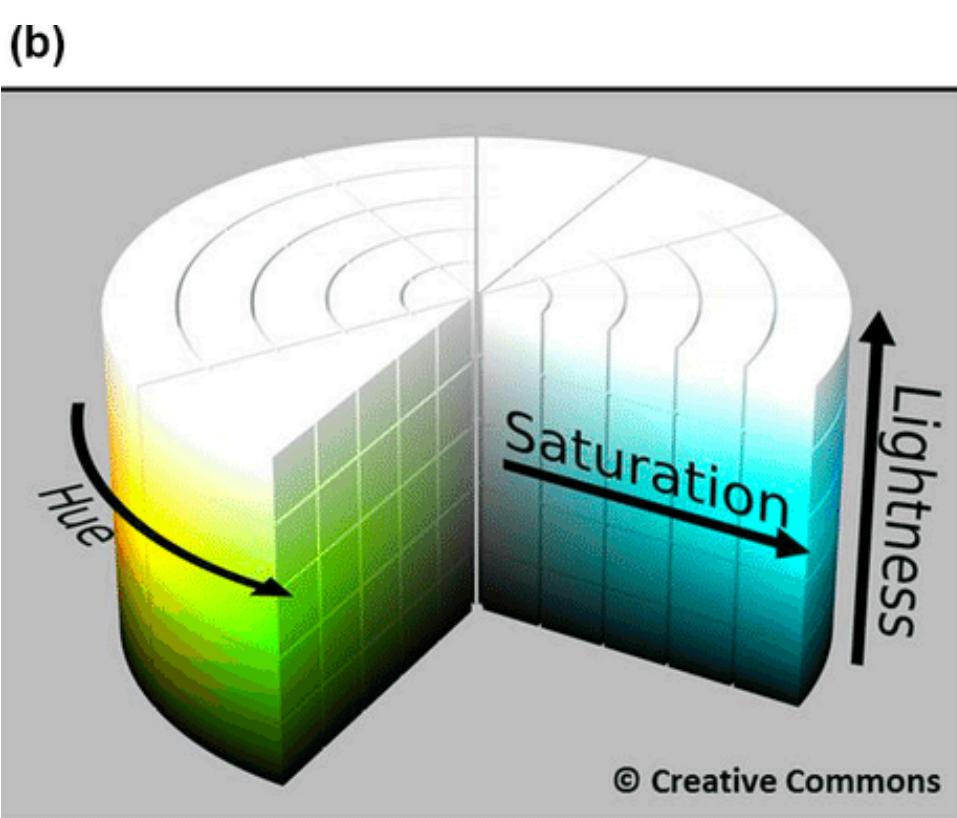
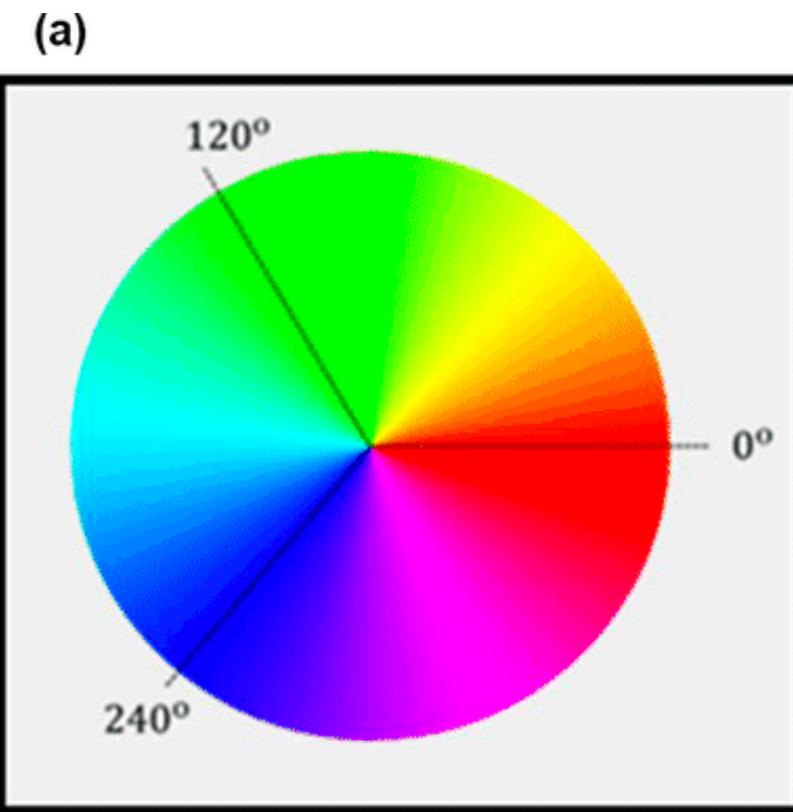


Method 2: “Luminance”



$$\text{Gray} = (\text{Red} * 0.3 + \text{Green} * 0.59 + \text{Blue} * 0.11)$$

Method 3: Desaturation



HSL: hue, saturation, lightness

$$\text{Gray} = (\text{Max(Red, Green, Blue)} + \text{Min(Red, Green, Blue)}) / 2$$

Luminance



Averaging

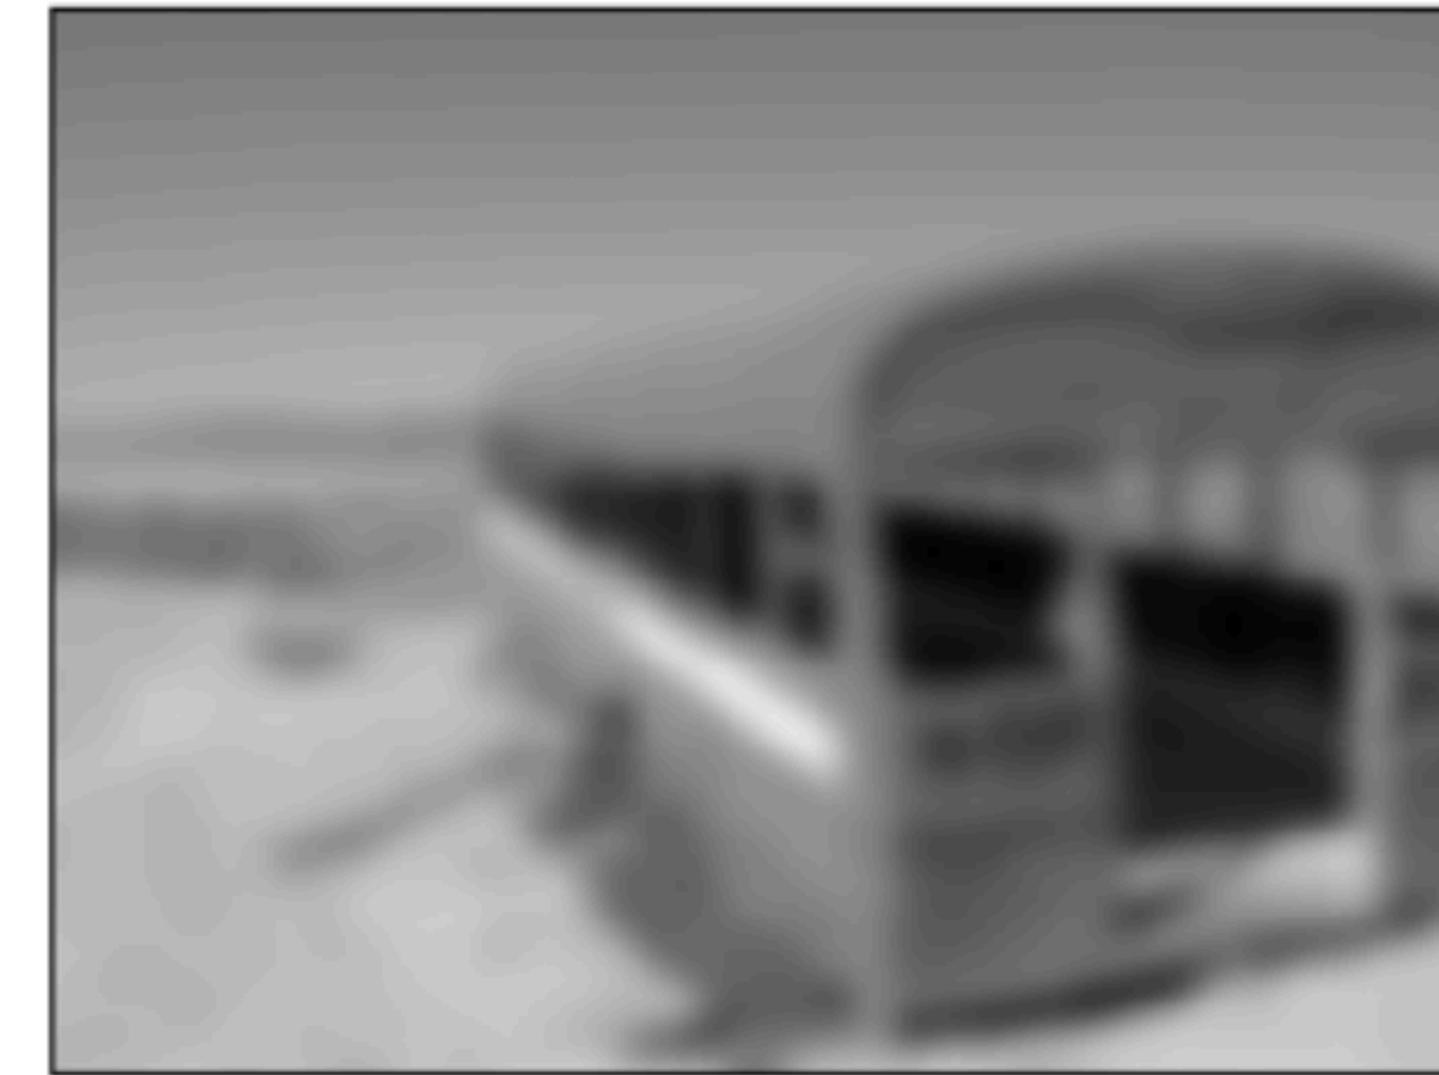


HSL



Filtering

- Point operations are limited
- They cannot accomplish tasks like sharpening or smoothing



Higher Dimensional Convolutions

0	0	0	0	0
0	$1/9$	$1/9$	$1/9$	0
0	$1/9$	$1/9$	$1/9$	0
0	$1/9$	$1/9$	$1/9$	0
0	0	0	0	0



Higher Dimensional Convolutions

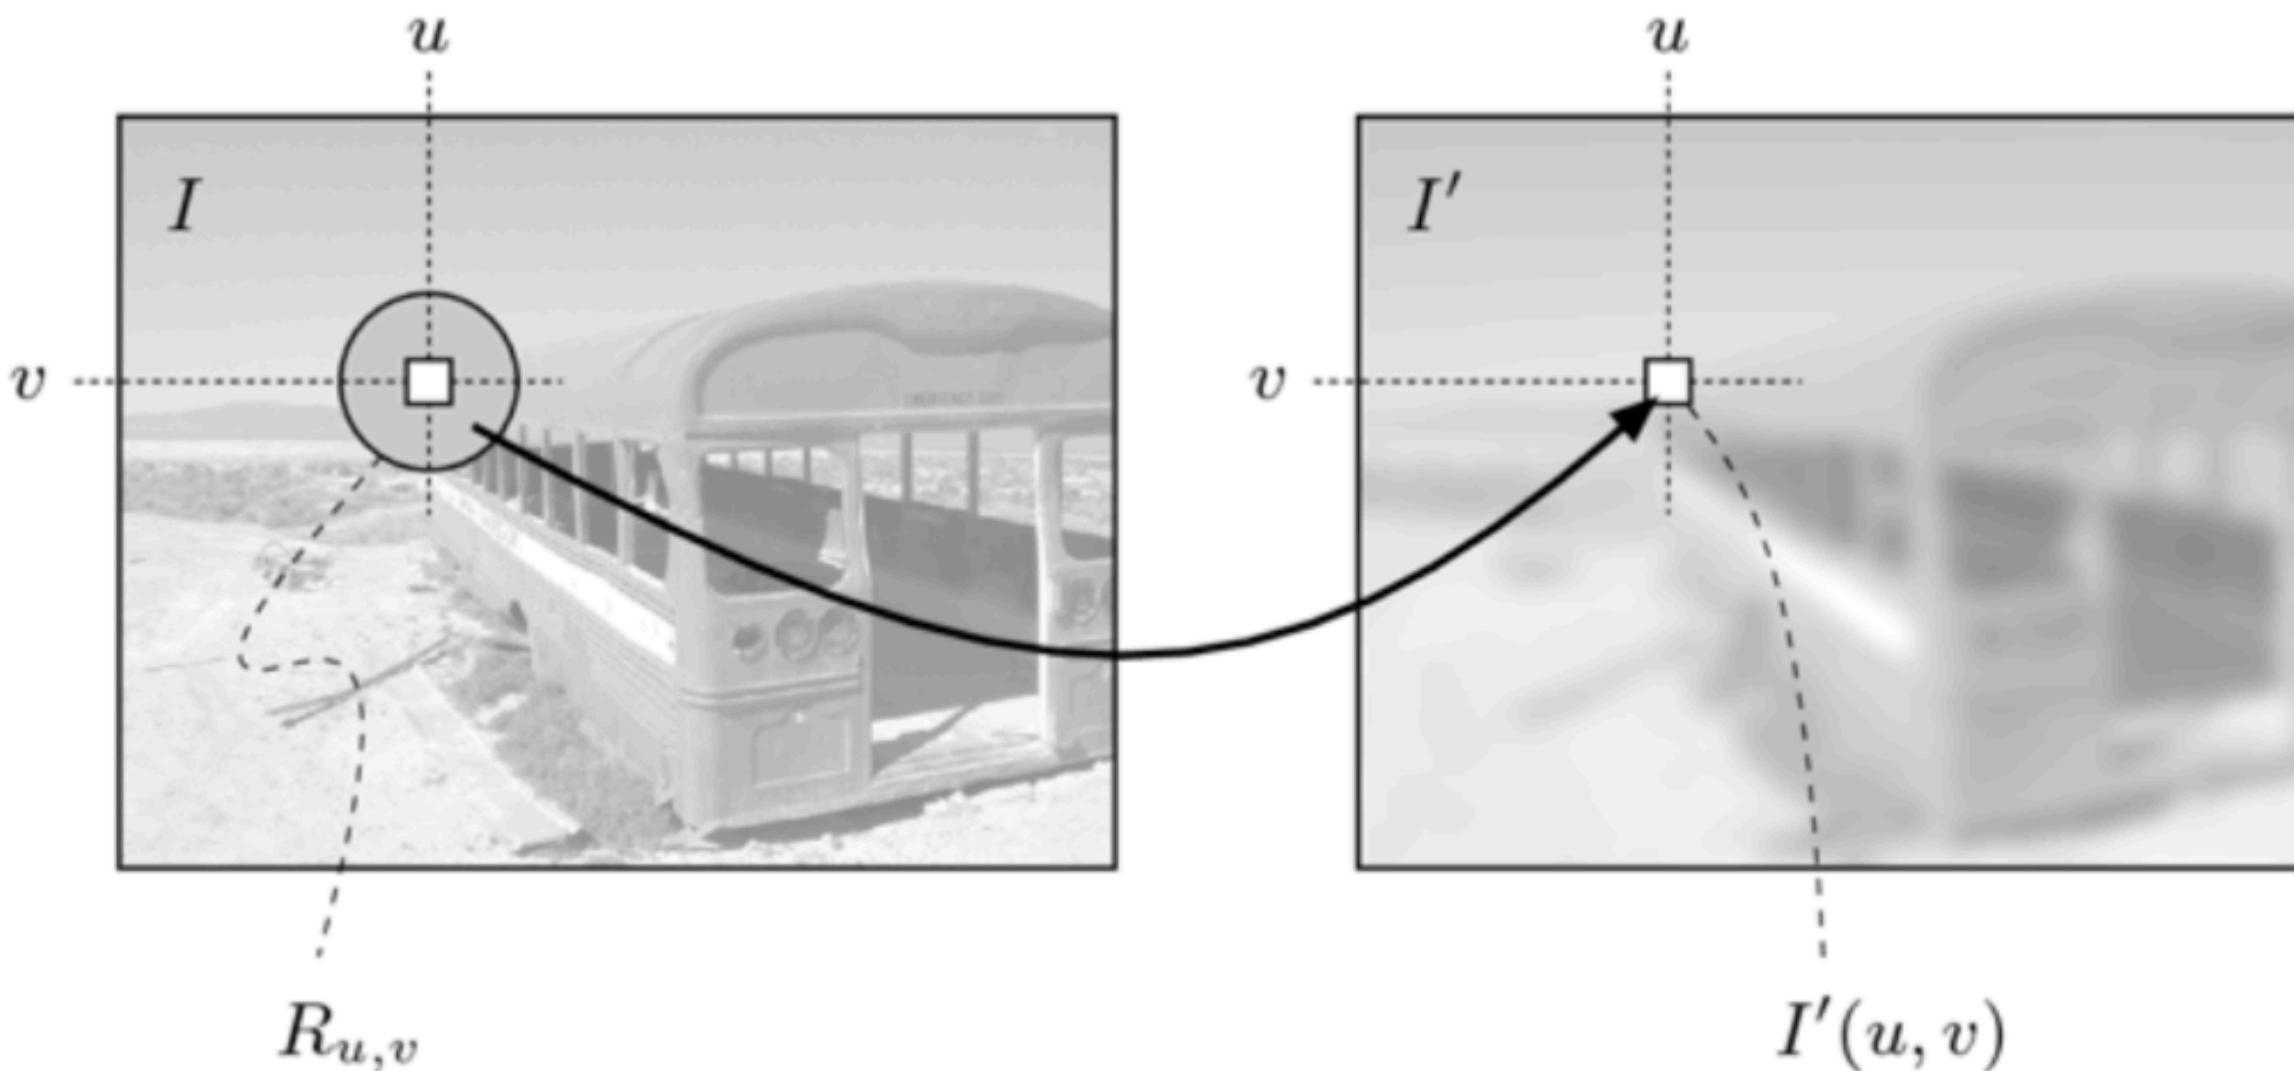
0	0	0	0	0
0	0	0	0	0
0	-1	1	0	0
0	0	0	0	0
0	0	0	0	0



Smoothing by averaging

- Replace each pixel by the average of its neighboring pixels
- Assume a 3x3 neighborhood:

$$I'(u, v) \leftarrow \frac{p_0 + p_1 + p_2 + p_3 + p_4 + p_5 + p_6 + p_7 + p_8}{9}$$



45	60	98	127	132	133	137	133
46	65	98	123	126	128	131	133
47	65	96	115	119	123	135	137
47	63	91	107	113	122	138	134
50	59	80	97	110	123	133	134
49	53	68	83	97	113	128	133
50	50	58	70	84	102	116	126
50	50	52	58	69	86	101	120

*

0.1	0.1	0.1
0.1	0.2	0.1
0.1	0.1	0.1

=

69	95	116	125	129	132
68	92	110	120	126	132
66	86	104	114	124	132
62	78	94	108	120	129
57	69	83	98	112	124
53	60	71	85	100	114

$f(x,y)$

$h(x,y)$

$g(x,y)$

Blur

$$\begin{pmatrix} 0.0625 & 0.125 & 0.0625 \\ 0.125 & 0.25 & 0.125 \\ 0.0625 & 0.125 & 0.0625 \end{pmatrix}$$

Below, for each 3x3 block of pixels in the image on the left, we multiply each pixel by the corresponding entry of the kernel and then take the sum. That sum becomes a new pixel in the image on the right. Hover over a pixel on either image to see how its value is computed.

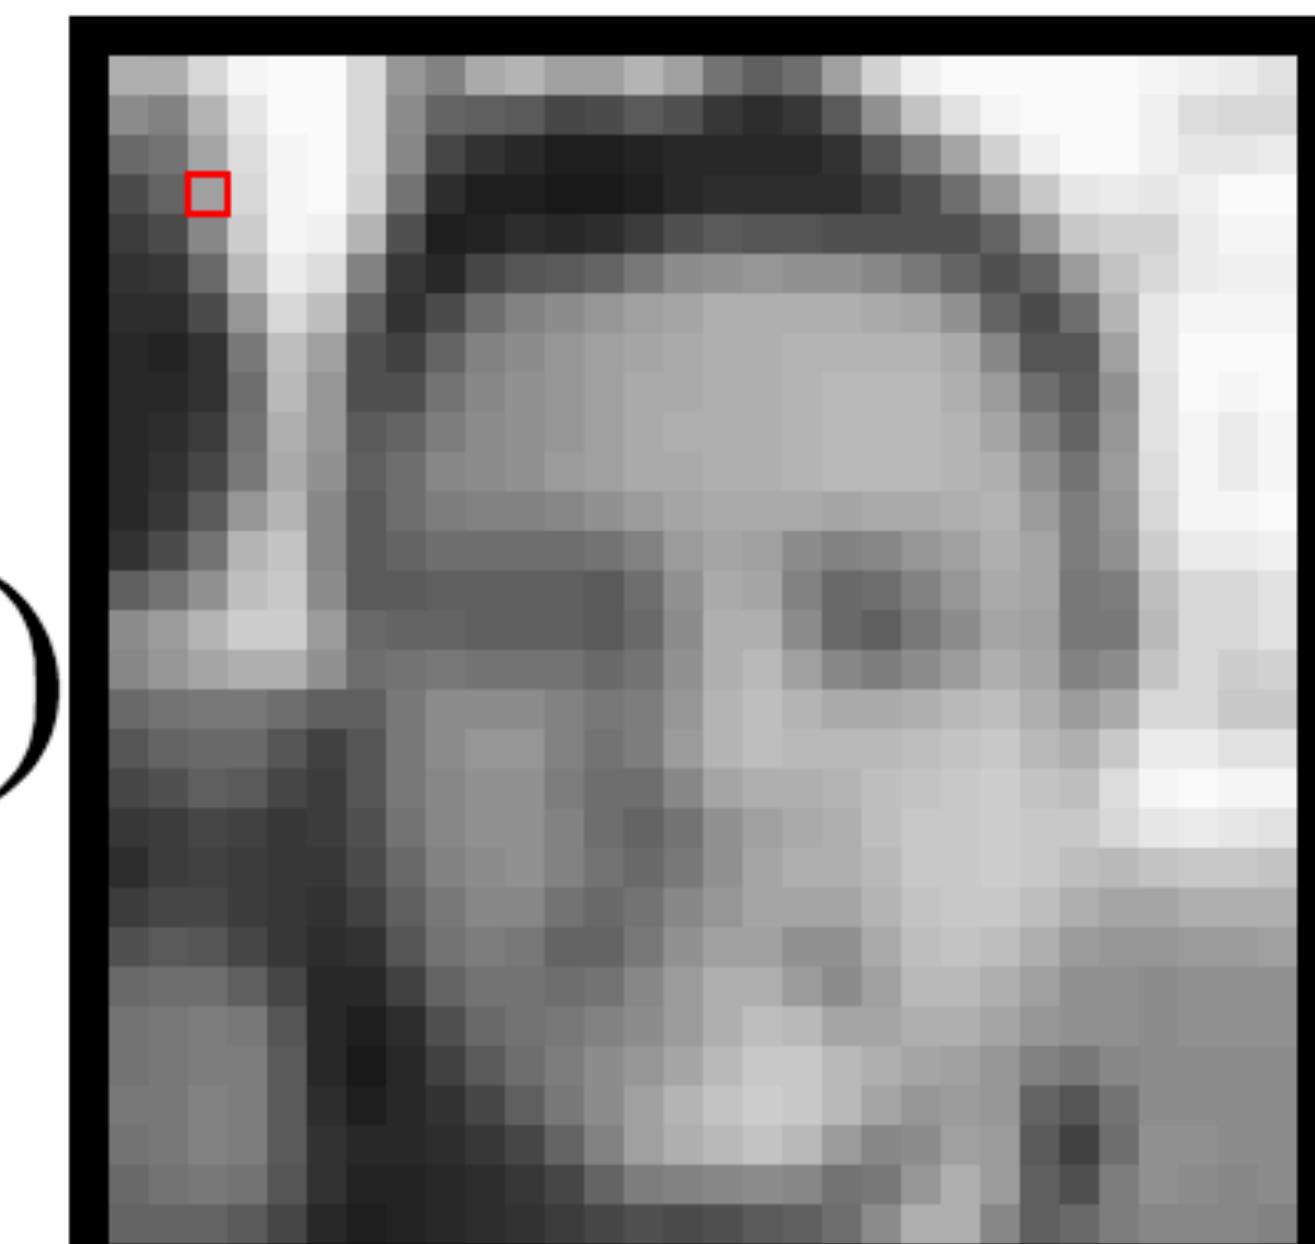


input image

$$\begin{aligned} & (\begin{array}{ccc} 96 & 143 & 223 \\ \times 0.0625 & \times 0.125 & \times 0.0625 \\ + 107 & 196 & 236 \\ \times 0.125 & \times 0.25 & \times 0.125 \\ + 45 & 134 & 218 \\ \times 0.0625 & \times 0.125 & \times 0.0625 \end{array}) \\ & = 163 \end{aligned}$$

kernel:

blur



output image

Sharpen

sharpen

$$\begin{pmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{pmatrix}$$

Below, for each 3x3 block of pixels in the image on the left, we multiply each pixel by the corresponding entry of the kernel and then take the sum. That sum becomes a new pixel in the image on the right. Hover over a pixel on either image to see how its value is computed.



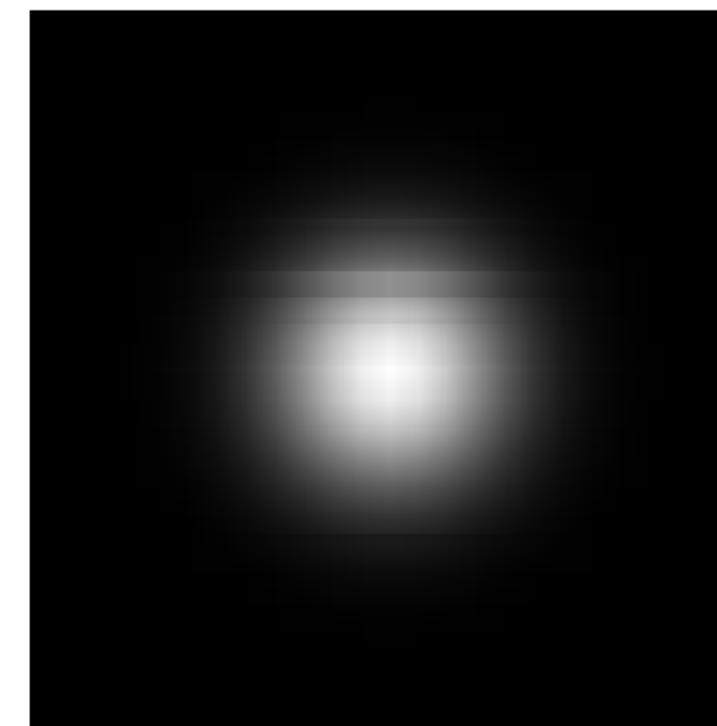
$$\begin{pmatrix} ? & 72 & 44 \\ \times 0 & \times -1 & \times 0 \\ + ? & 55 & 20 \\ \times -1 & \times 5 & \times -1 \\ + ? & 65 & 49 \\ \times 0 & \times -1 & \times 0 \\ = ? \end{pmatrix}$$

kernel:
sharpen



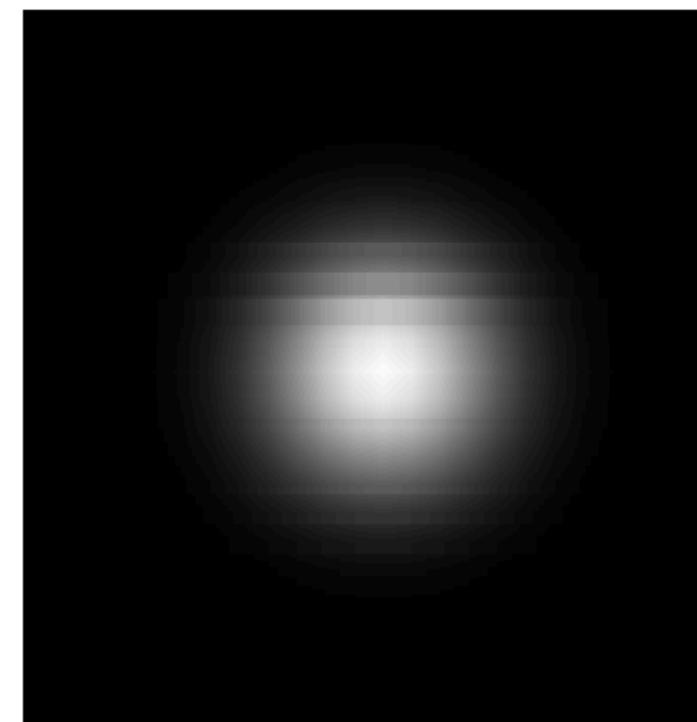
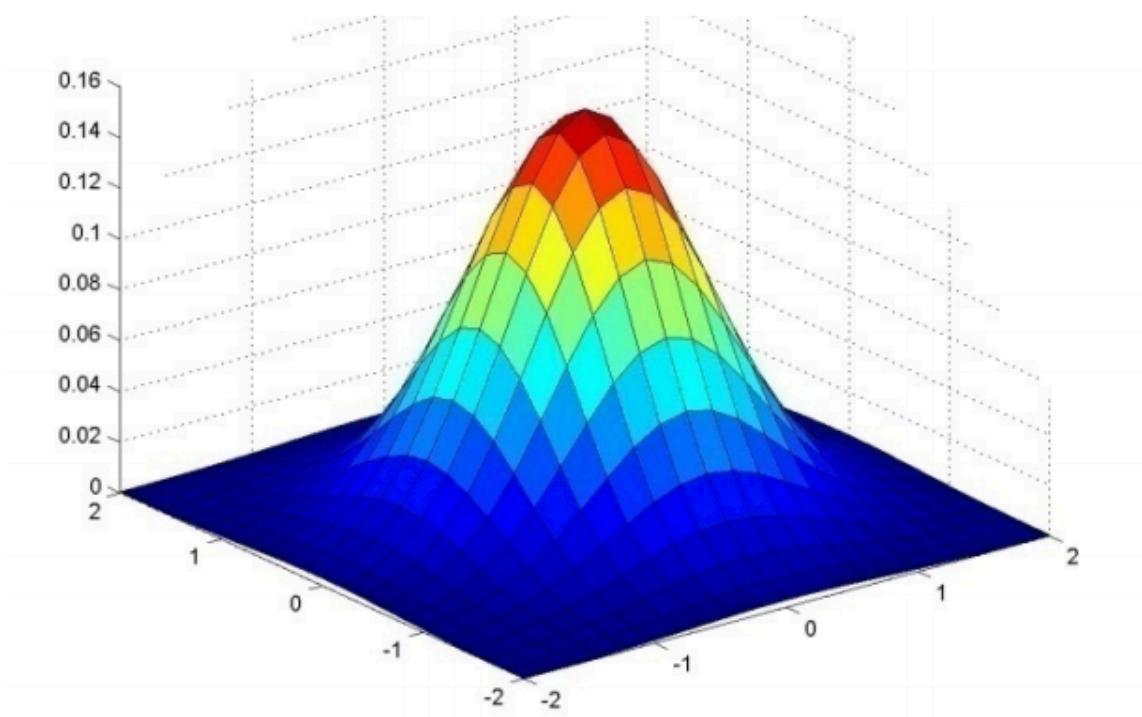
Smoothing with a Gaussian

- Single point of light looks like a fuzzy blob, but averaging process would give a little square
- To eliminate edge effects, weight contribution of neighborhood pixels according to their closeness to center



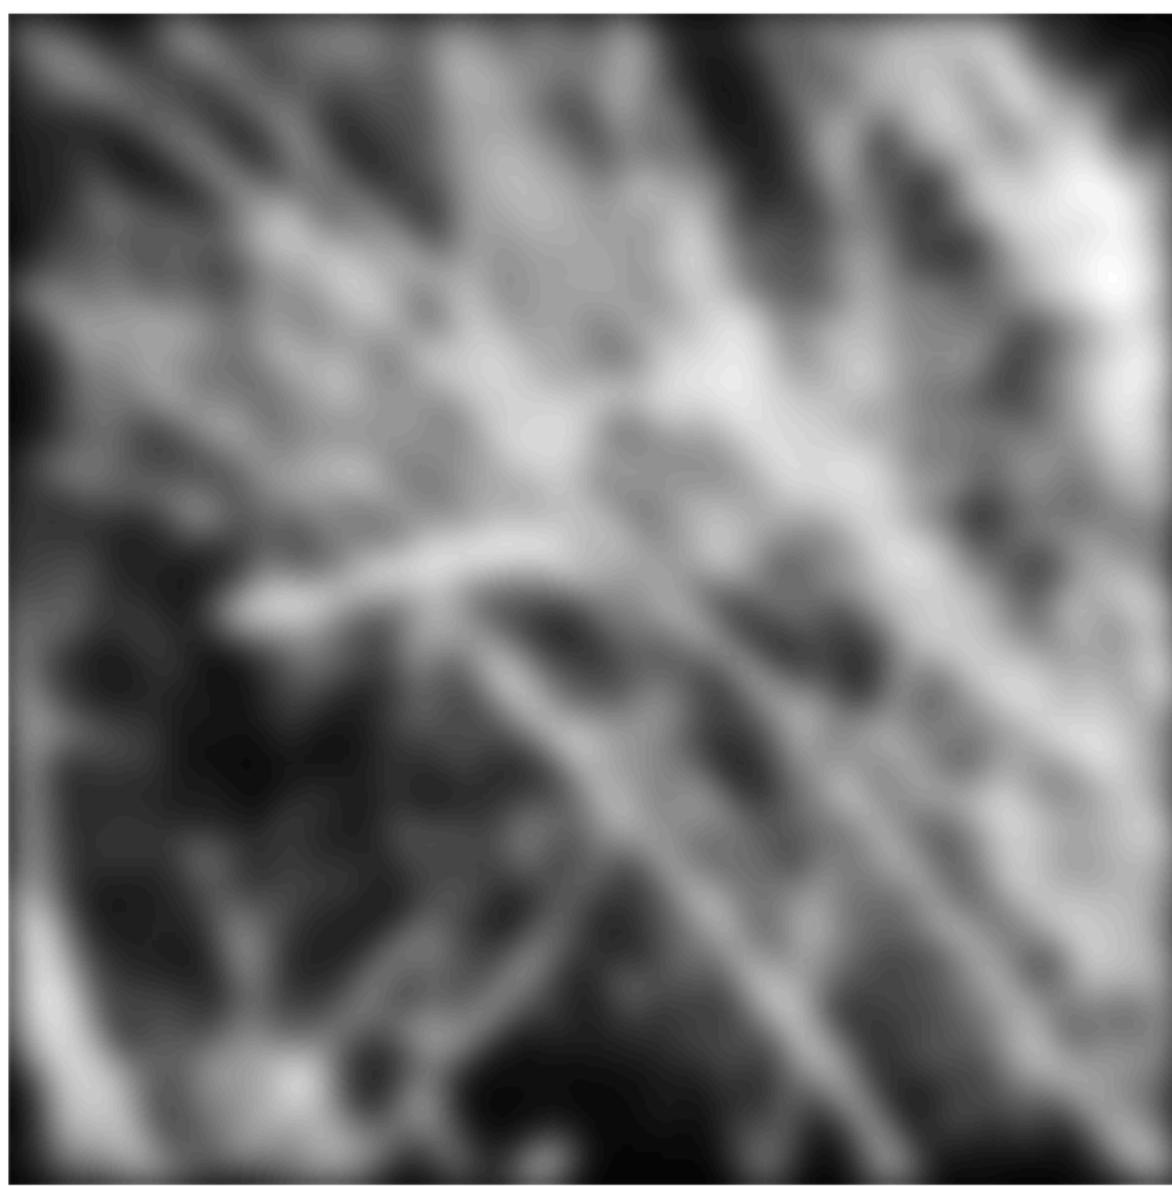
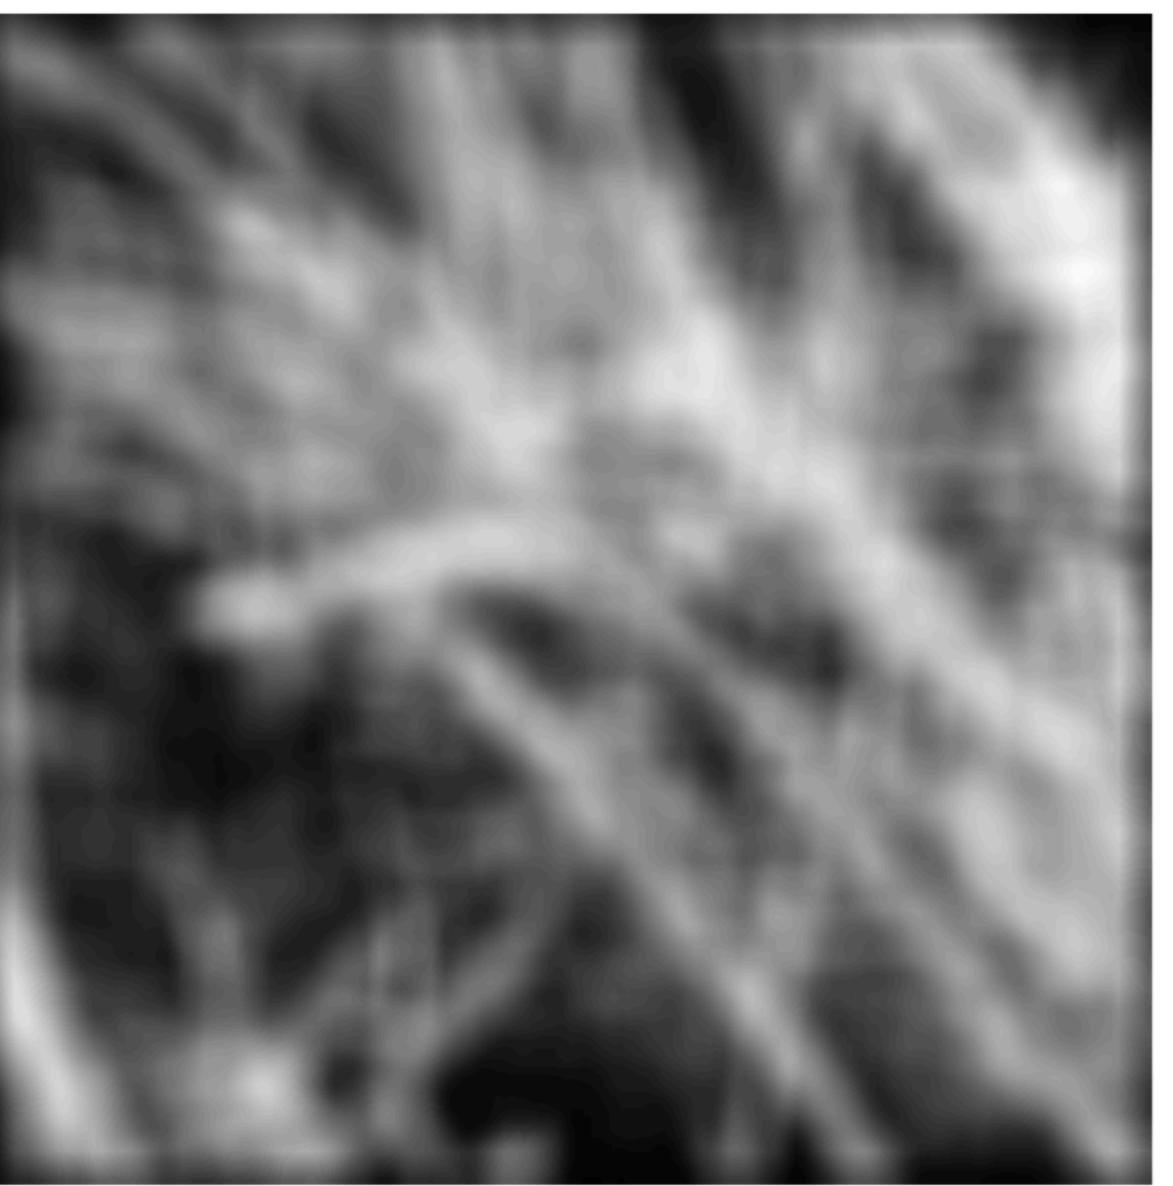
Gaussian kernel

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

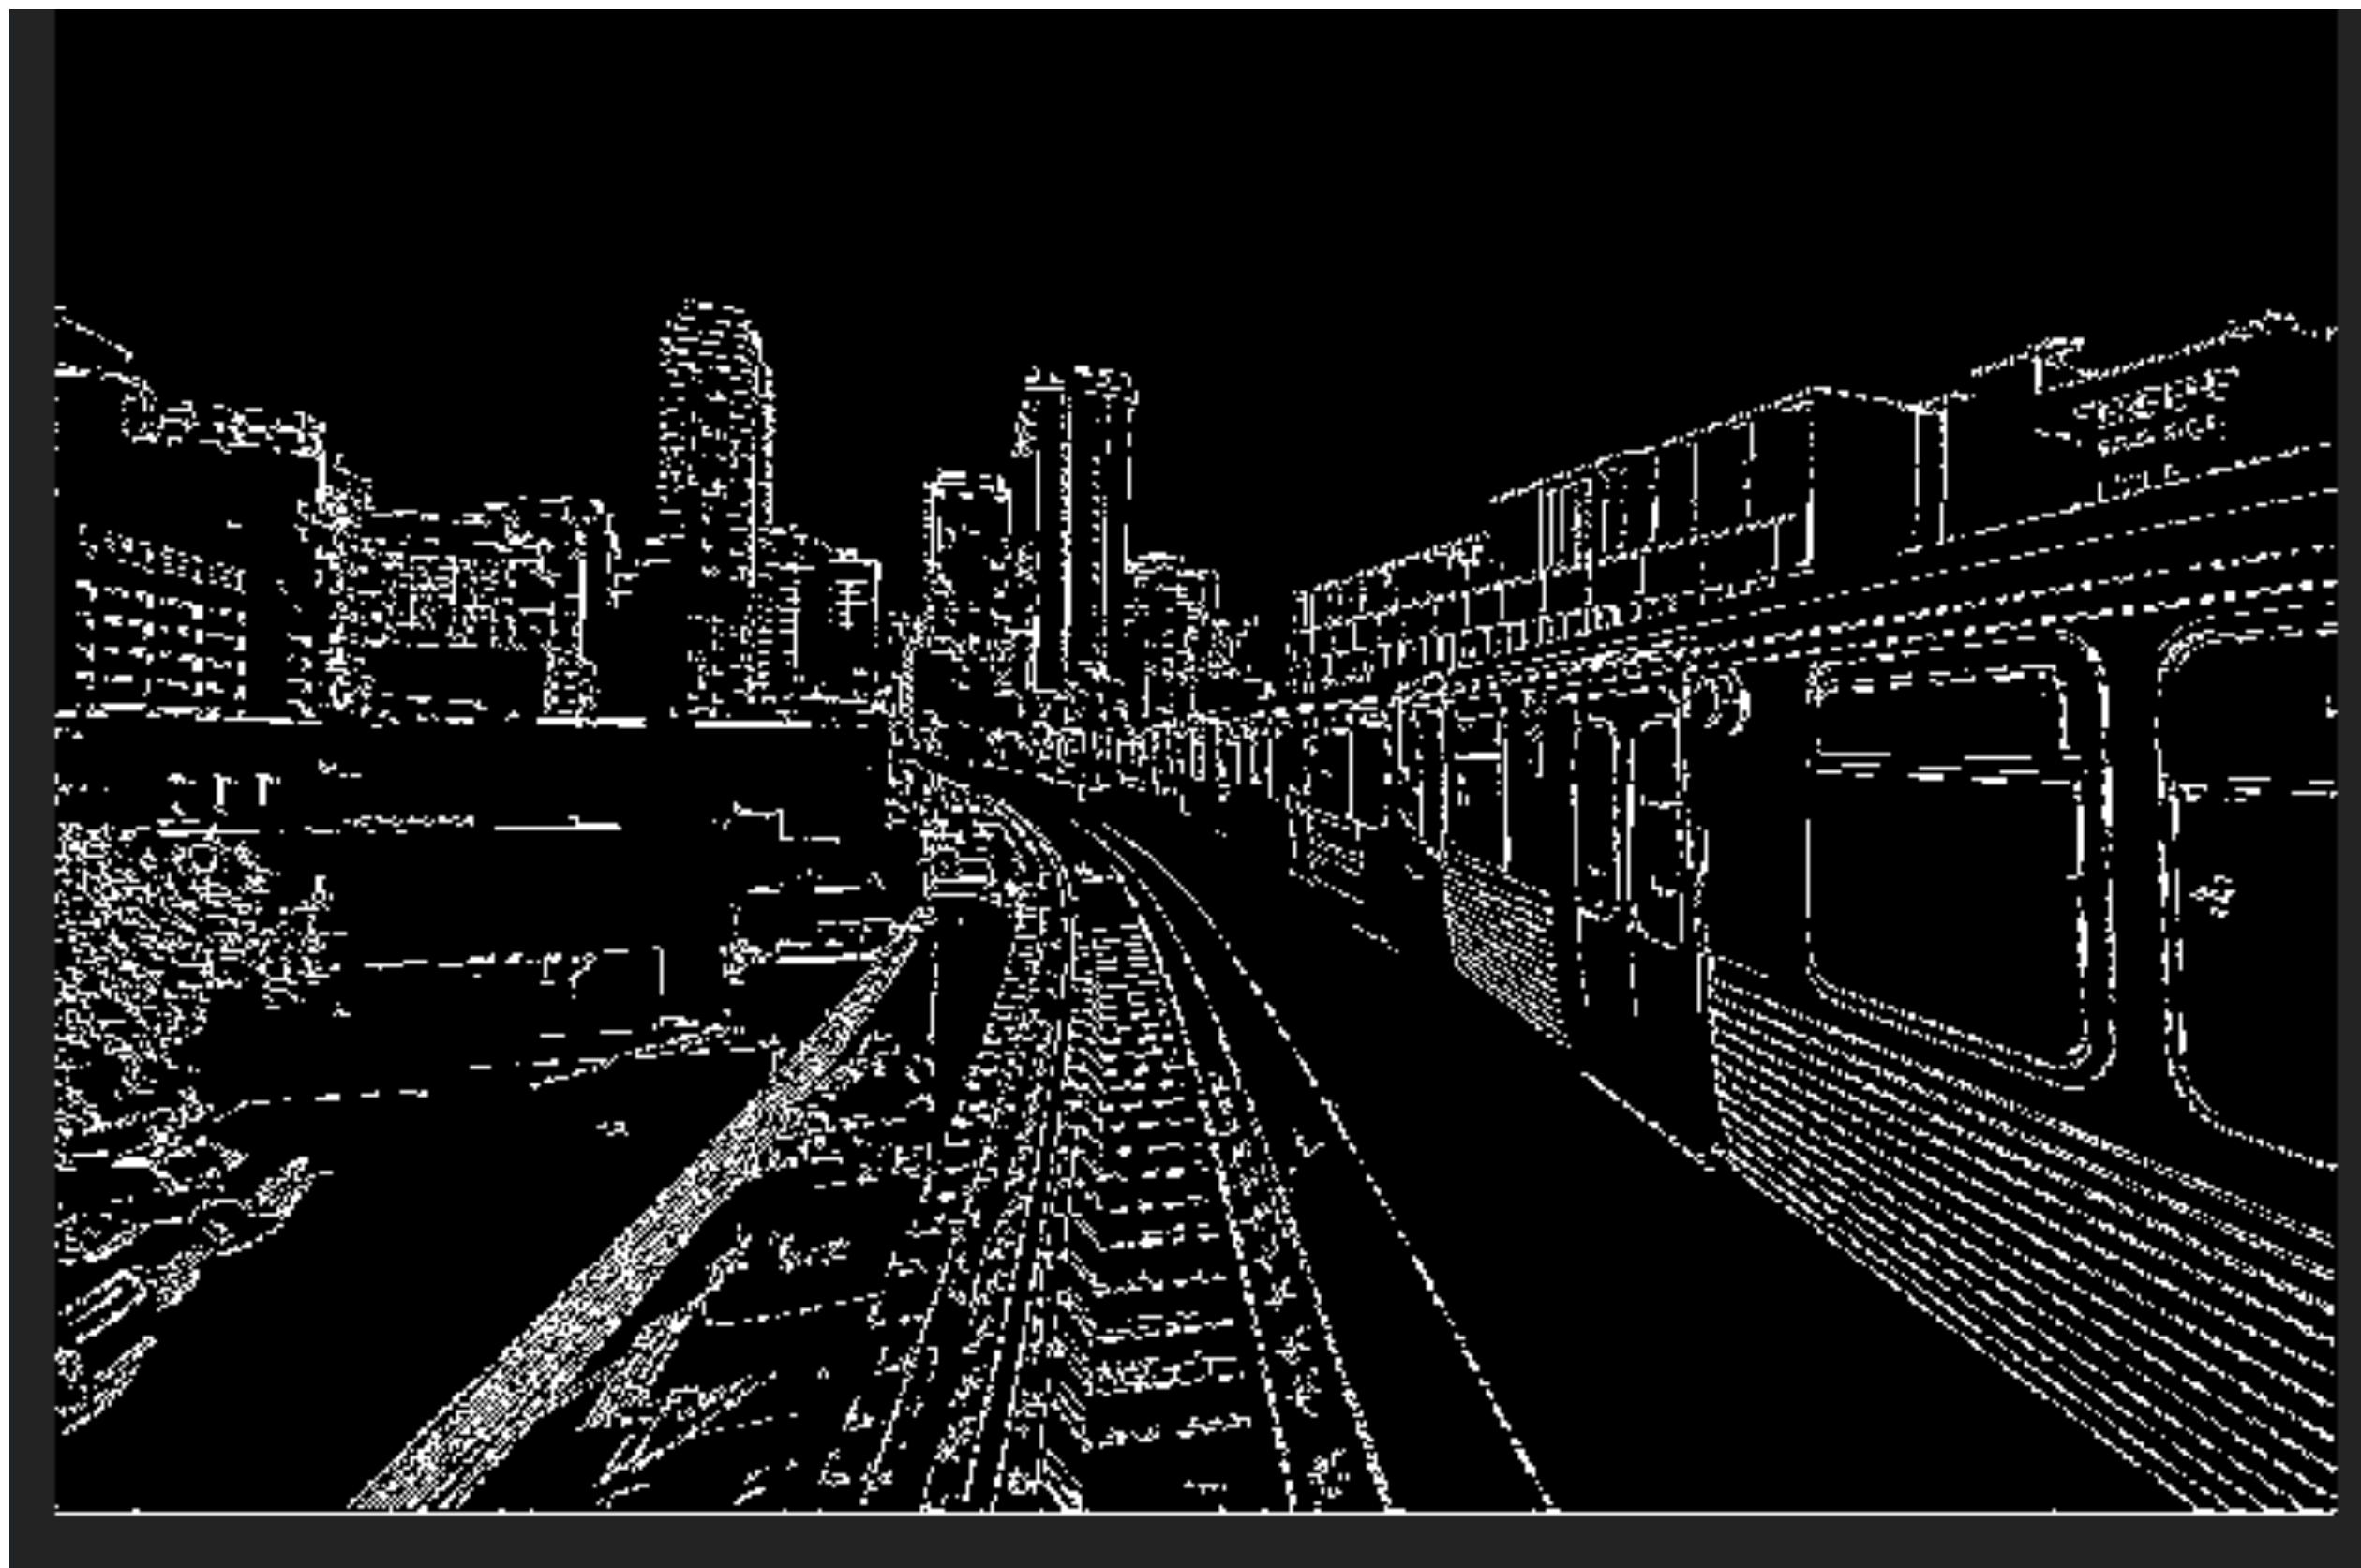


0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

5 x 5, $\sigma = 1$

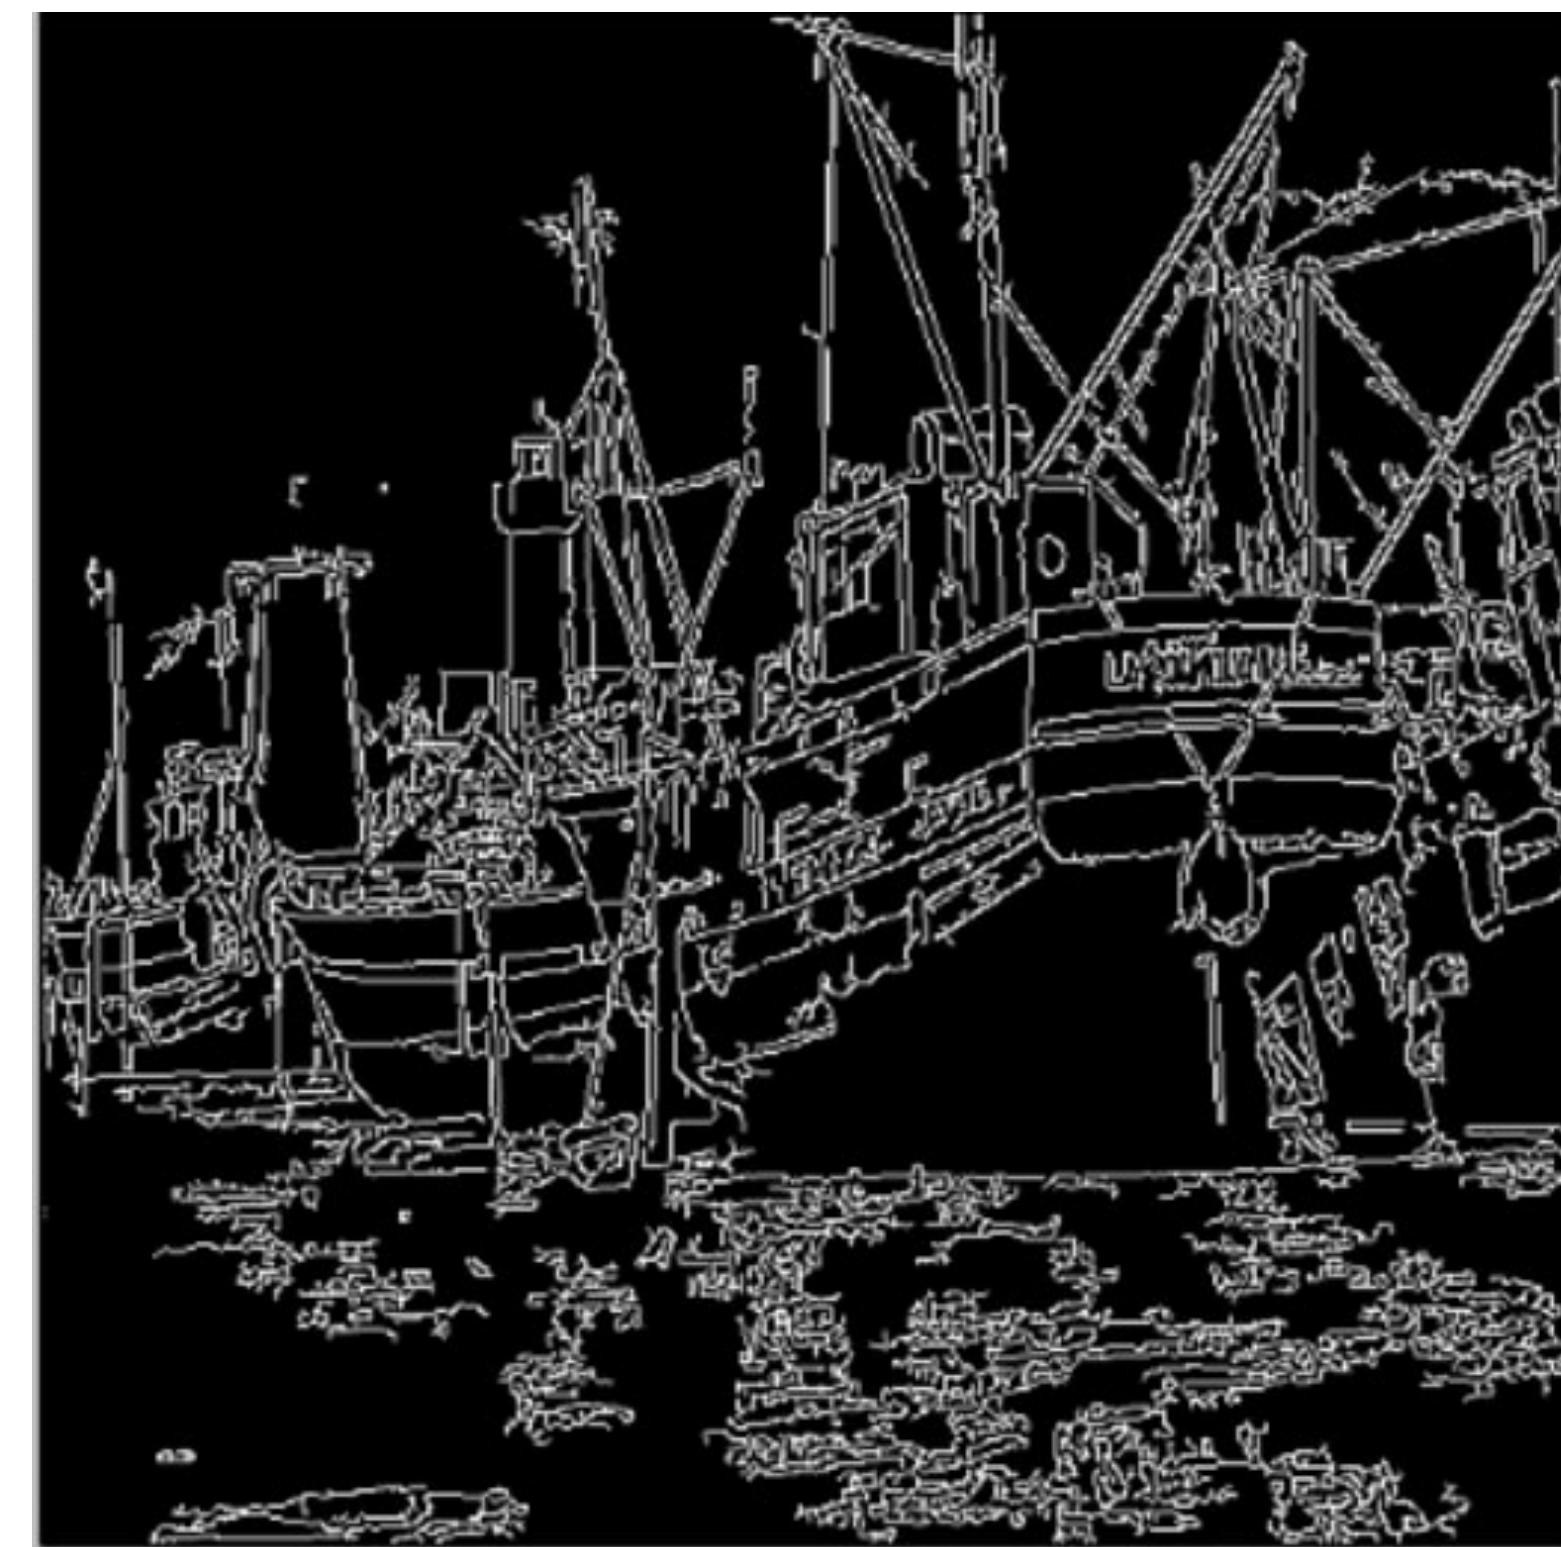


Edge Detection: Canny Edge Detector



Edge Detection: Canny Edge Detector

- Gaussian filter to remove noise
- Find intensity gradients
- Apply non-maximum suppression to remove spurious responses
- Apply double threshold to determine potential edges
- Track edge by hysteresis



```
1 import cv2
2 import matplotlib.pyplot as plt
3
4 im = cv2.imread('boat.png')
5 edges = cv2.Canny(im,25,255,L2gradient=False)
6 plt.imshow(edges,cmap='gray')
7 plt.show()
```

$$\text{Edge Gradient } (G) = \sqrt{G_x^2 + G_y^2}$$

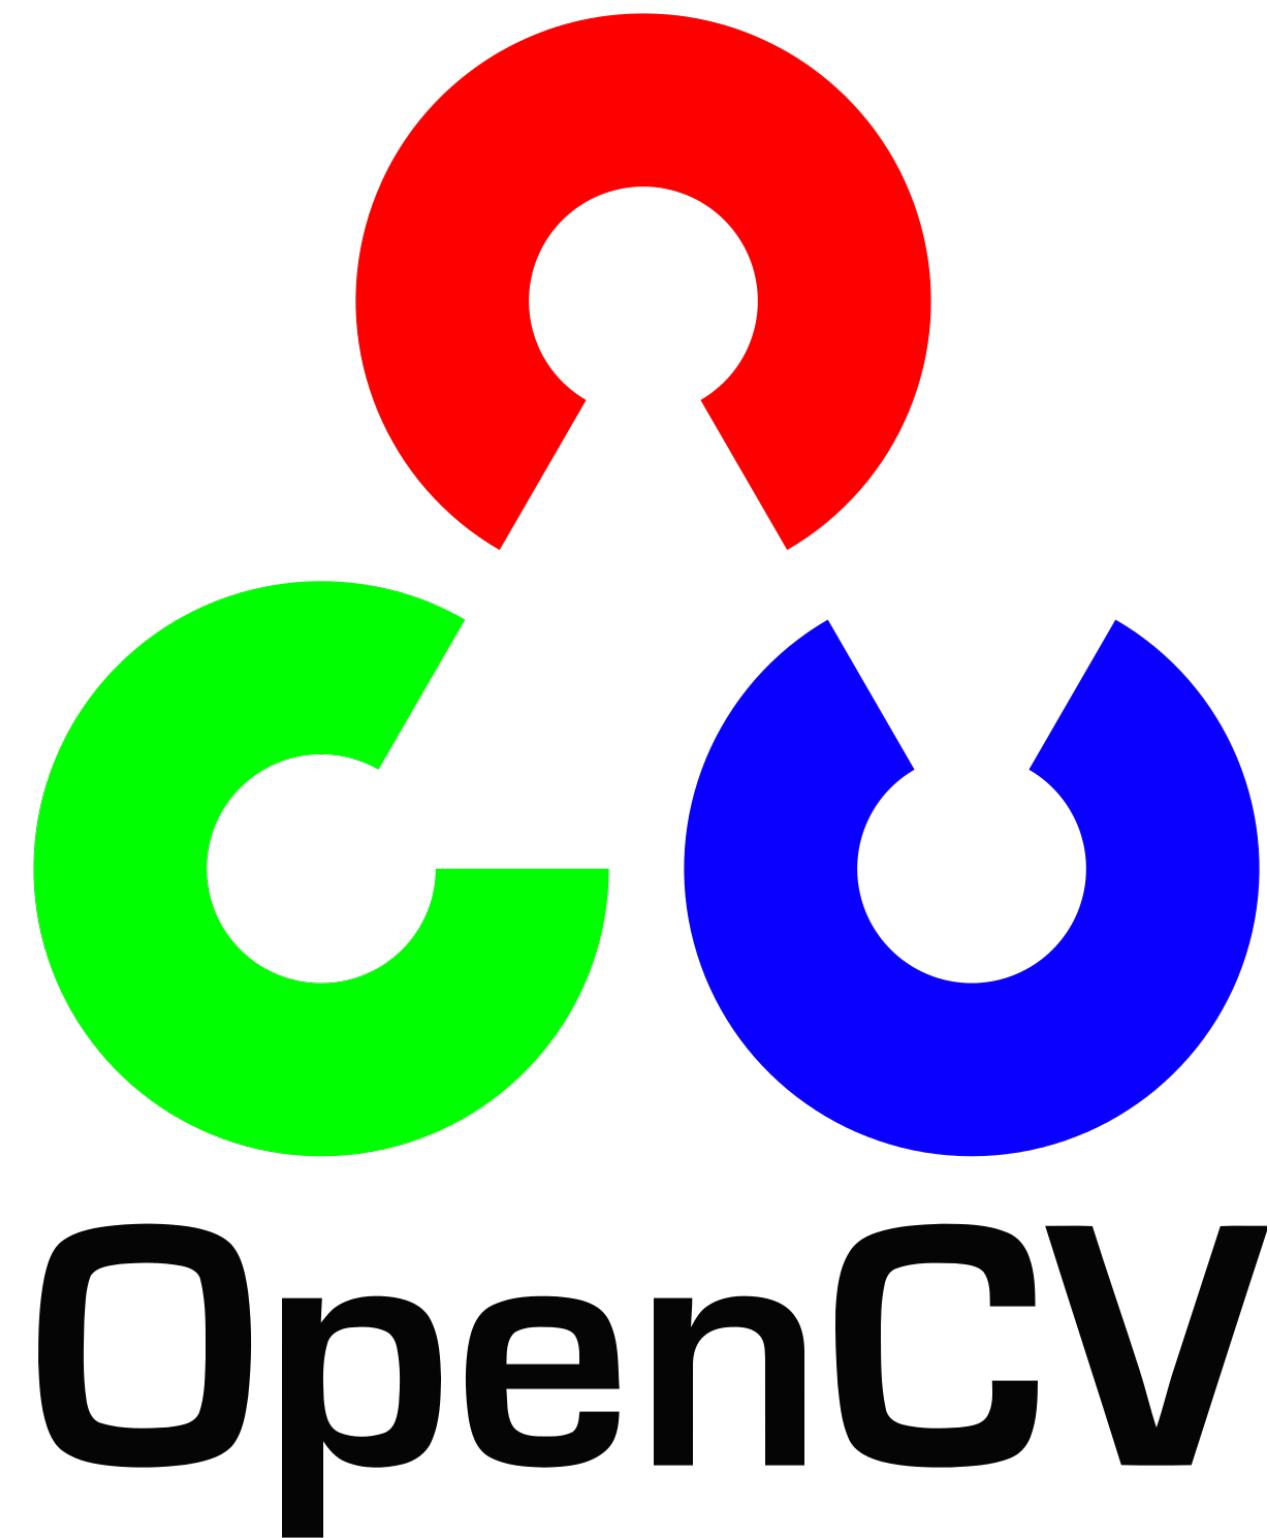
$$\text{Angle } (\theta) = \tan^{-1} \left(\frac{G_y}{G_x} \right)$$

$$\text{Edge Gradient } (G) = |G_x| + |G_y|$$

OpenCV

Installation

- OpenCV
 - https://docs.opencv.org/master/da/df6/tutorial_py_table_of_contents_setup.html
- Macs:
 - pip install opencv-python
- Imutils, scikit-image
 - Used for this lecture (for convenience)



OpenCV

Loading and Displaying an Image

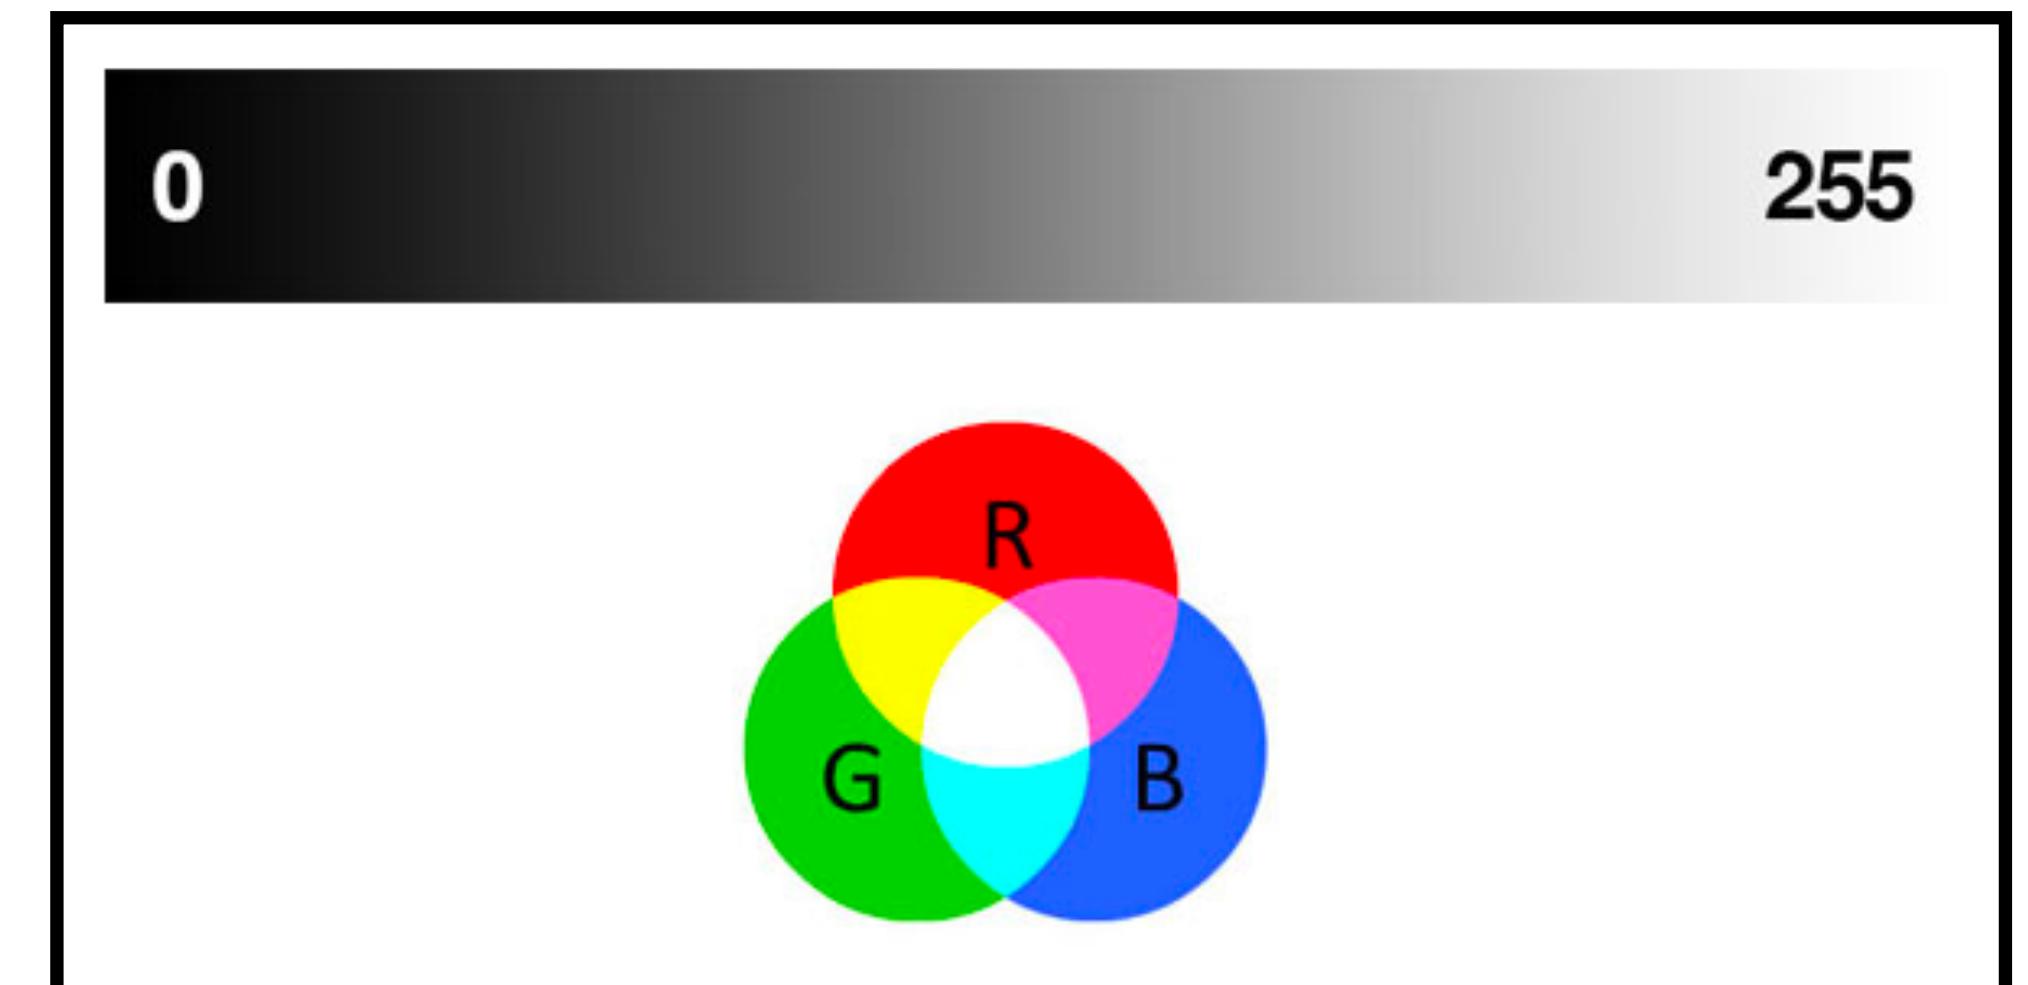


```
1 # import the necessary packages
2 import imutils
3 import cv2
4
5 # load the input image and show its dimensions, keeping in mind that
6 # images are represented as a multi-dimensional NumPy array with
7 # shape no. rows (height) x no. columns (width) x no. channels (depth)
8 image = cv2.imread("jurassicpark.jpg")
9 (h, w, d) = image.shape
10 print("width={}, height={}, depth={}".format(w, h, d))
11
12 # display the image to our screen -- we will need to click the window
13 # open by OpenCV and press a key on our keyboard to continue execution
14 cv2.imshow("Image", image)
15 cv2.waitKey(0)
```

OpenCV

Accessing Individual Pixels

- OpenCV color images in (B,G,R)
 - Each value has range [0,255]
 - $256 \times 256 \times 256 = 16777216$ colors



```
16  
17 # access the RGB pixel located at x=50, y=100, keepind in mind that  
18 # OpenCV stores images in BGR order rather than RGB  
19 (B, G, R) = image[100, 50]  
20 print("R={}, G={}, B={}".format(R, G, B))
```

OpenCV

Array Slicing and Cropping

```
22 # extract a 200x200 pixel square ROI (Region of Interest) from the
23 # input image
24 roi = image[150:350, 730:930]
25 cv2.imshow("ROI", roi)
26 cv2.waitKey(0)
```



OpenCV

Resizing Images

```
28 # resize the image to 200x200px, ignoring aspect ratio
29 resized = cv2.resize(image, (200, 200))
30 cv2.imshow("Fixed Resizing", resized)
31 cv2.waitKey(0)
```



OpenCV

Resizing Images

```
33 # fixed resizing and distort aspect ratio so let's resize the width
34 # to be 300px but compute the new height based on the aspect ratio
35 r = 300.0 / w
36 dim = (300, int(h * r))
37 resized = cv2.resize(image, dim)
38 cv2.imshow("Aspect Ratio Resize", resized)
39 cv2.waitKey(0)
```



OpenCV

Rotating Images

```
47 # let's rotate an image 45 degrees clockwise using OpenCV by first
48 # computing the image center, then constructing the rotation matrix,
49 # and then finally applying the affine warp
50 center = (w // 2, h // 2)
51 M = cv2.getRotationMatrix2D(center, -45, 1.0)
52 rotated = cv2.warpAffine(image, M, (w, h))
53 cv2.imshow("OpenCV Rotation", rotated)
54 cv2.waitKey(0)
```



OpenCV

Rotating Images

```
61 # OpenCV doesn't "care" if our rotated image is clipped after rotation
62 # so we can instead use another imutils convenience function to help
63 # us out
64 rotated = imutils.rotate_bound(image, 45)
65 cv2.imshow("Imutils Bound Rotation", rotated)
66 cv2.waitKey(0)
```



OpenCV

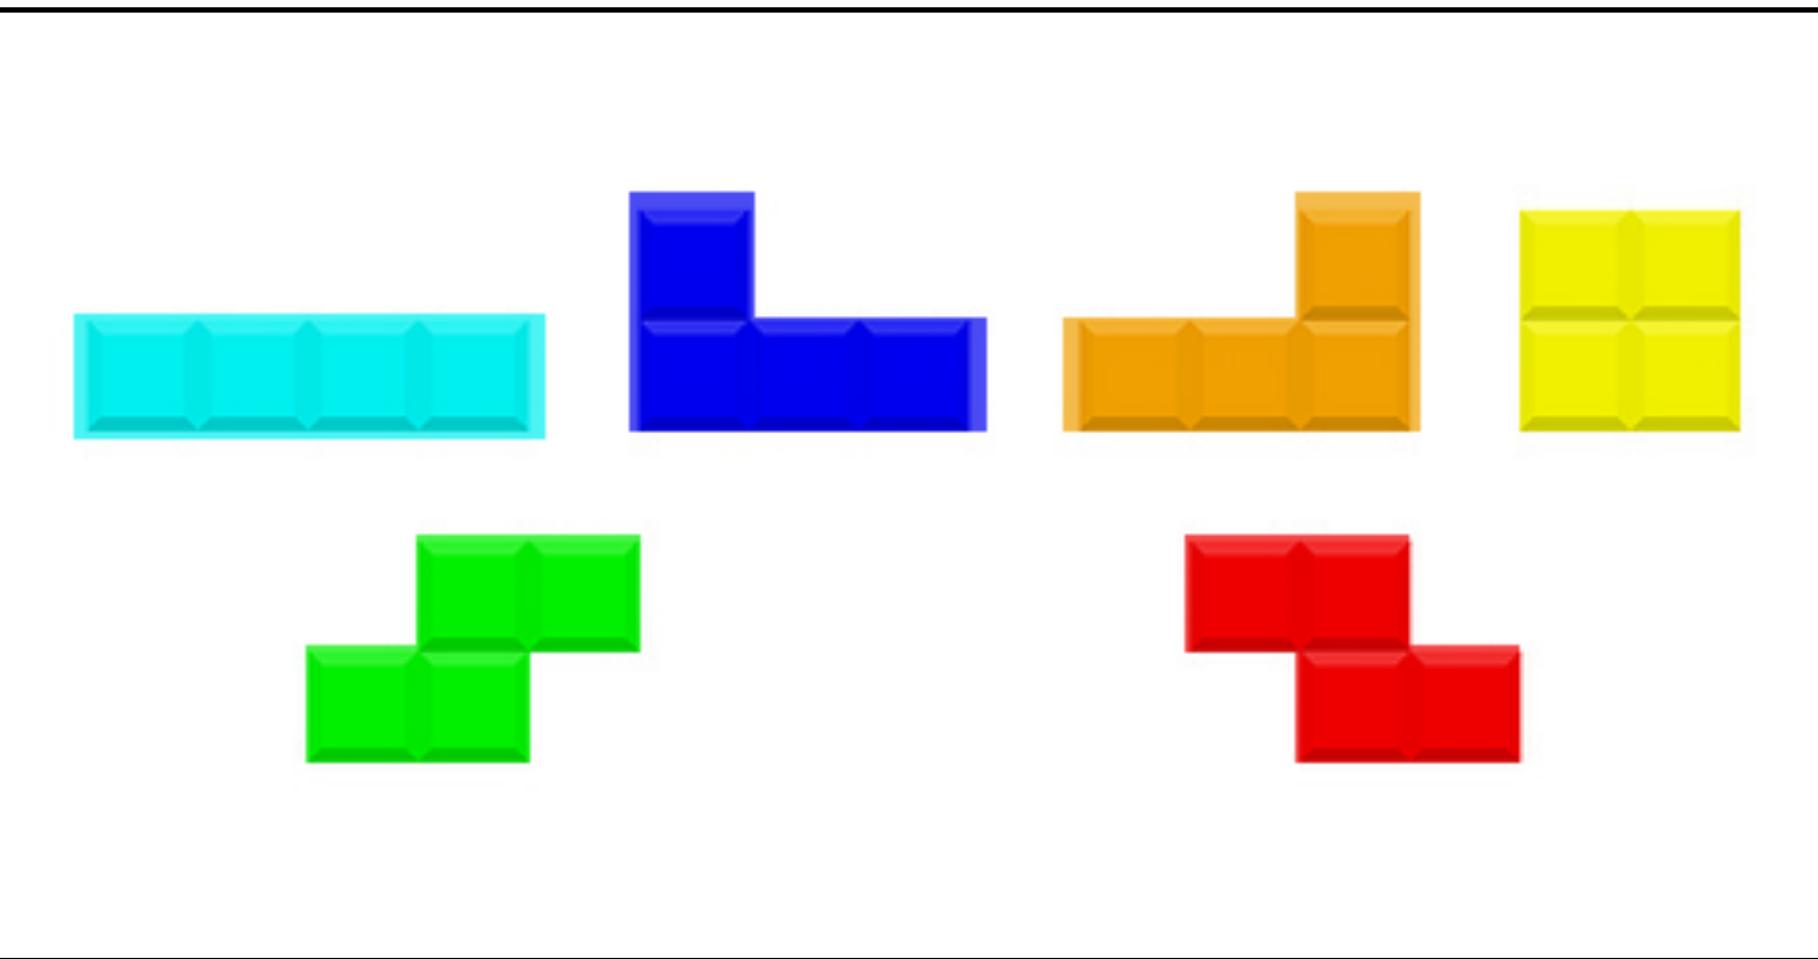
Smoothing Images

```
68 # apply a Gaussian blur with a 11x11 kernel to the image to smooth it,  
69 # useful when reducing high frequency noise  
70 blurred = cv2.GaussianBlur(image, (11, 11), 0)  
71 cv2.imshow("Blurred", blurred)  
72 cv2.waitKey(0)
```



OpenCV

Counting Objects

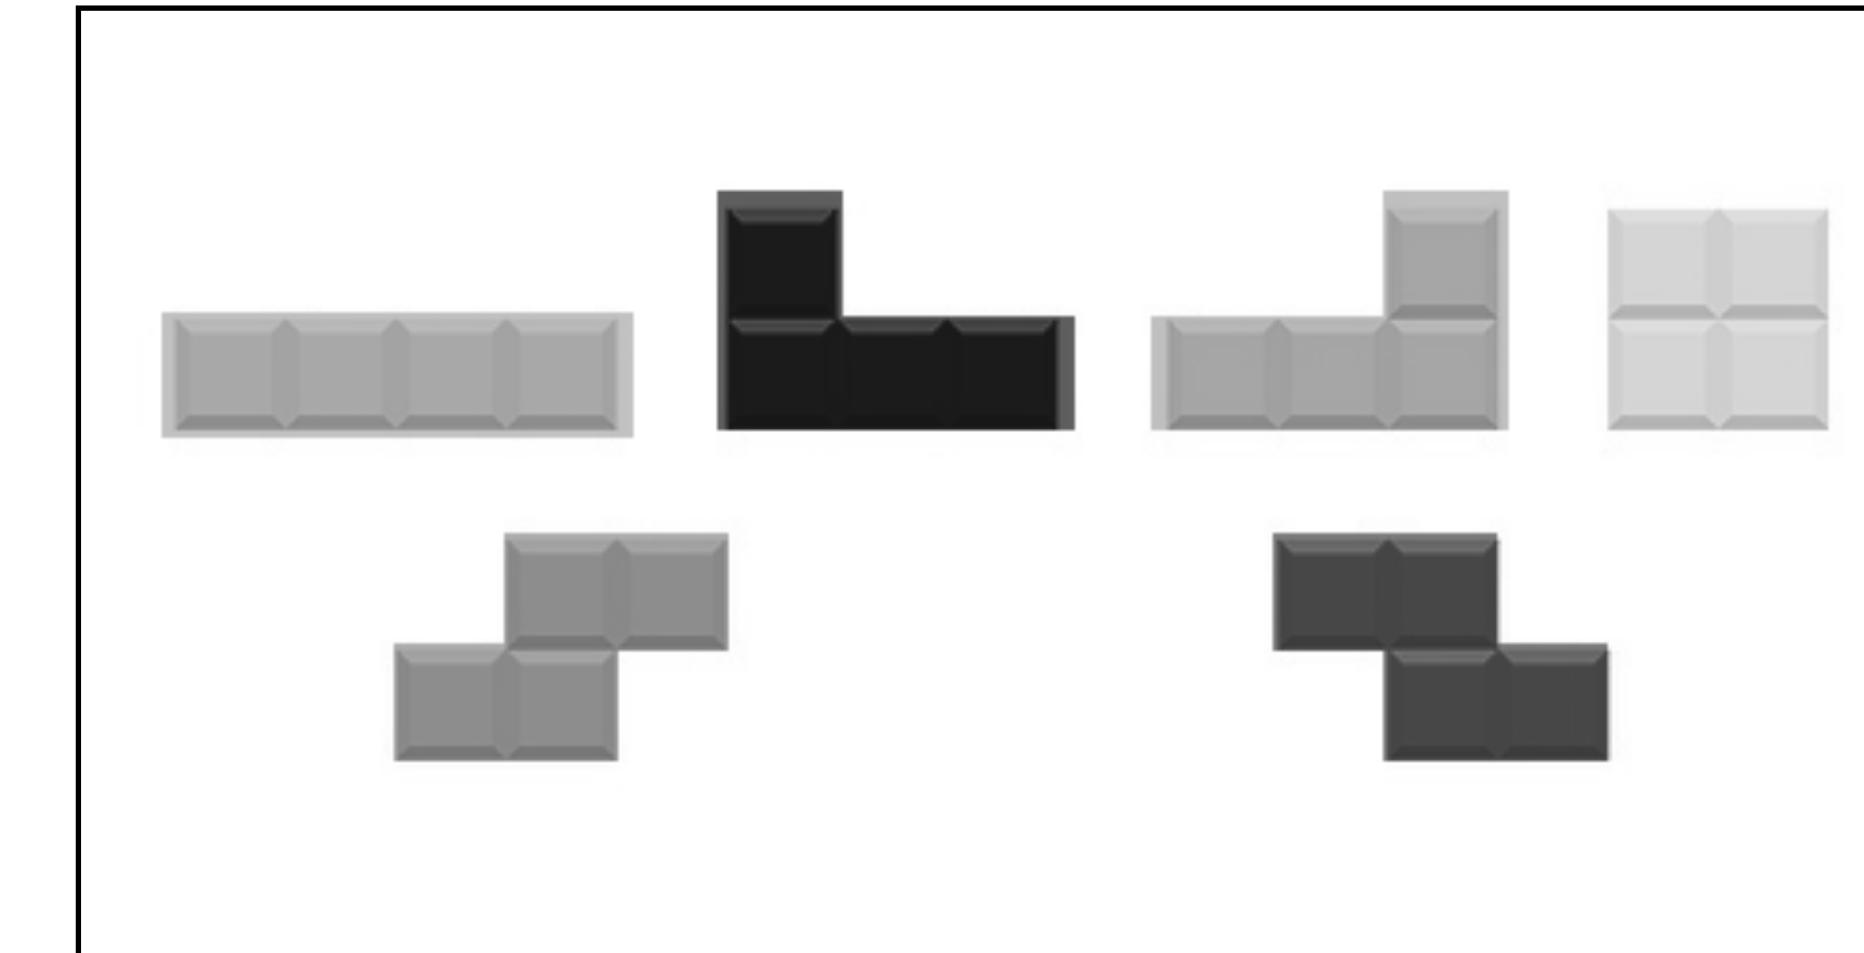
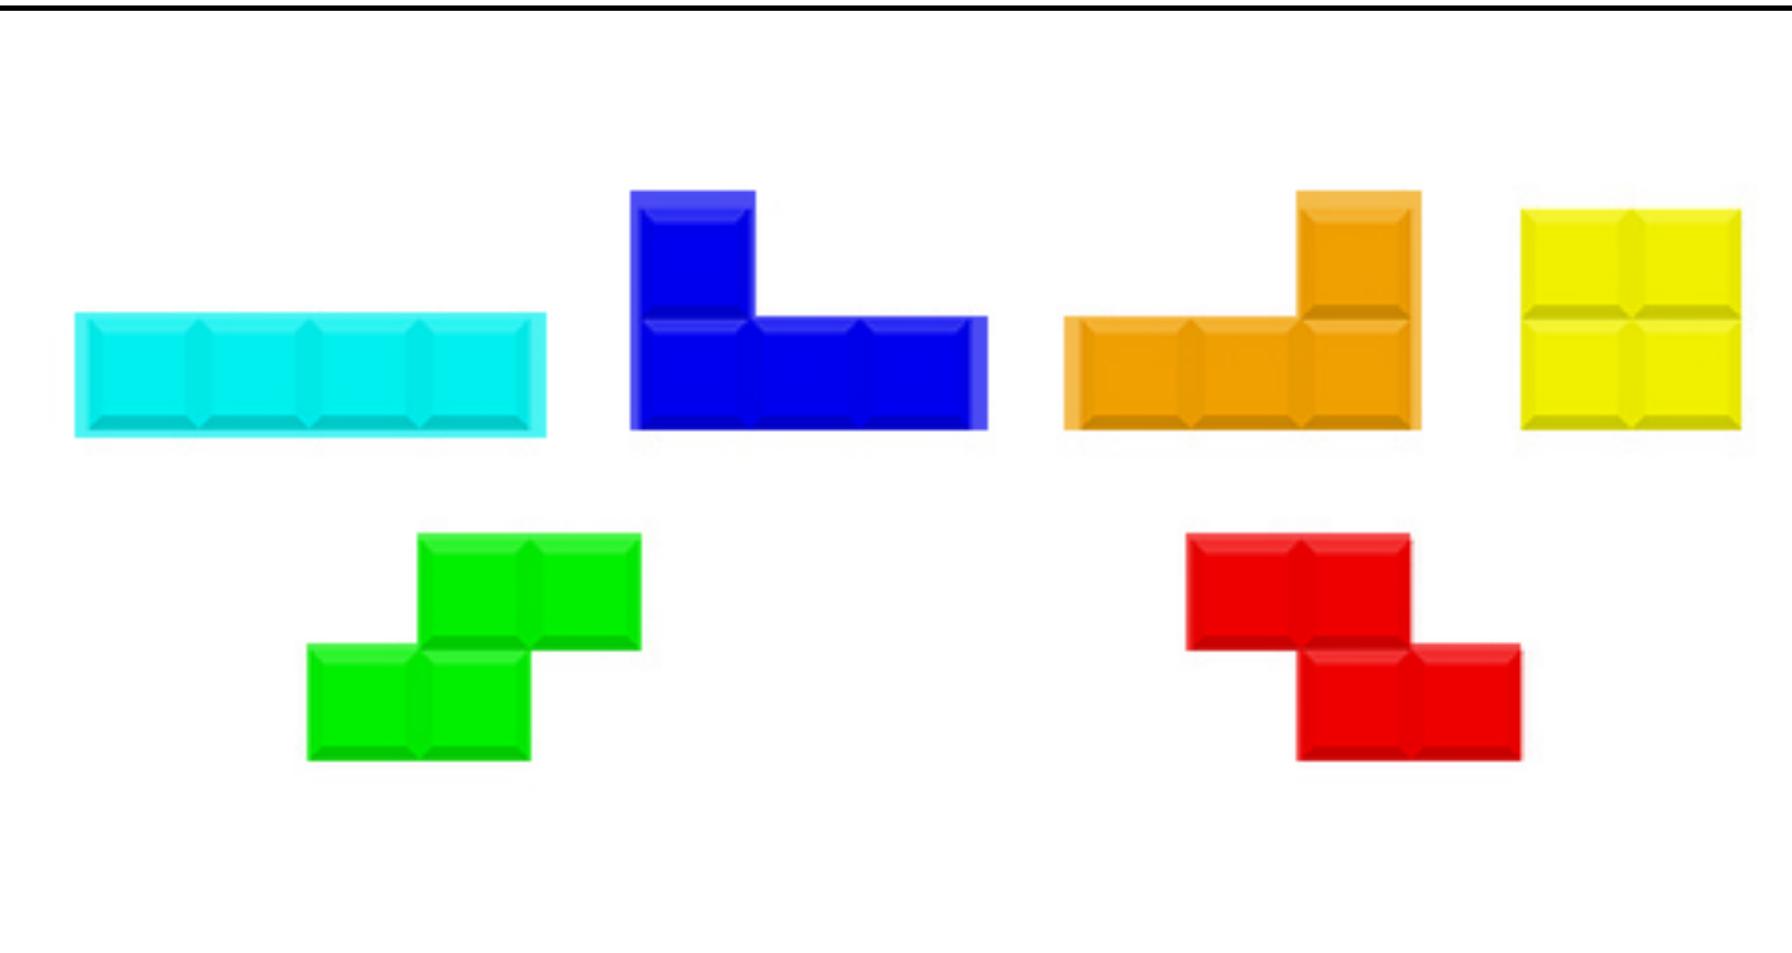


```
1 # import the necessary packages
2 import argparse
3 import imutils
4 import cv2
5
6 # construct the argument parser and parse the arguments
7 ap = argparse.ArgumentParser()
8 ap.add_argument("-i", "--image", required=True,
9                 help="path to input image")
10 args = vars(ap.parse_args())
```

OpenCV

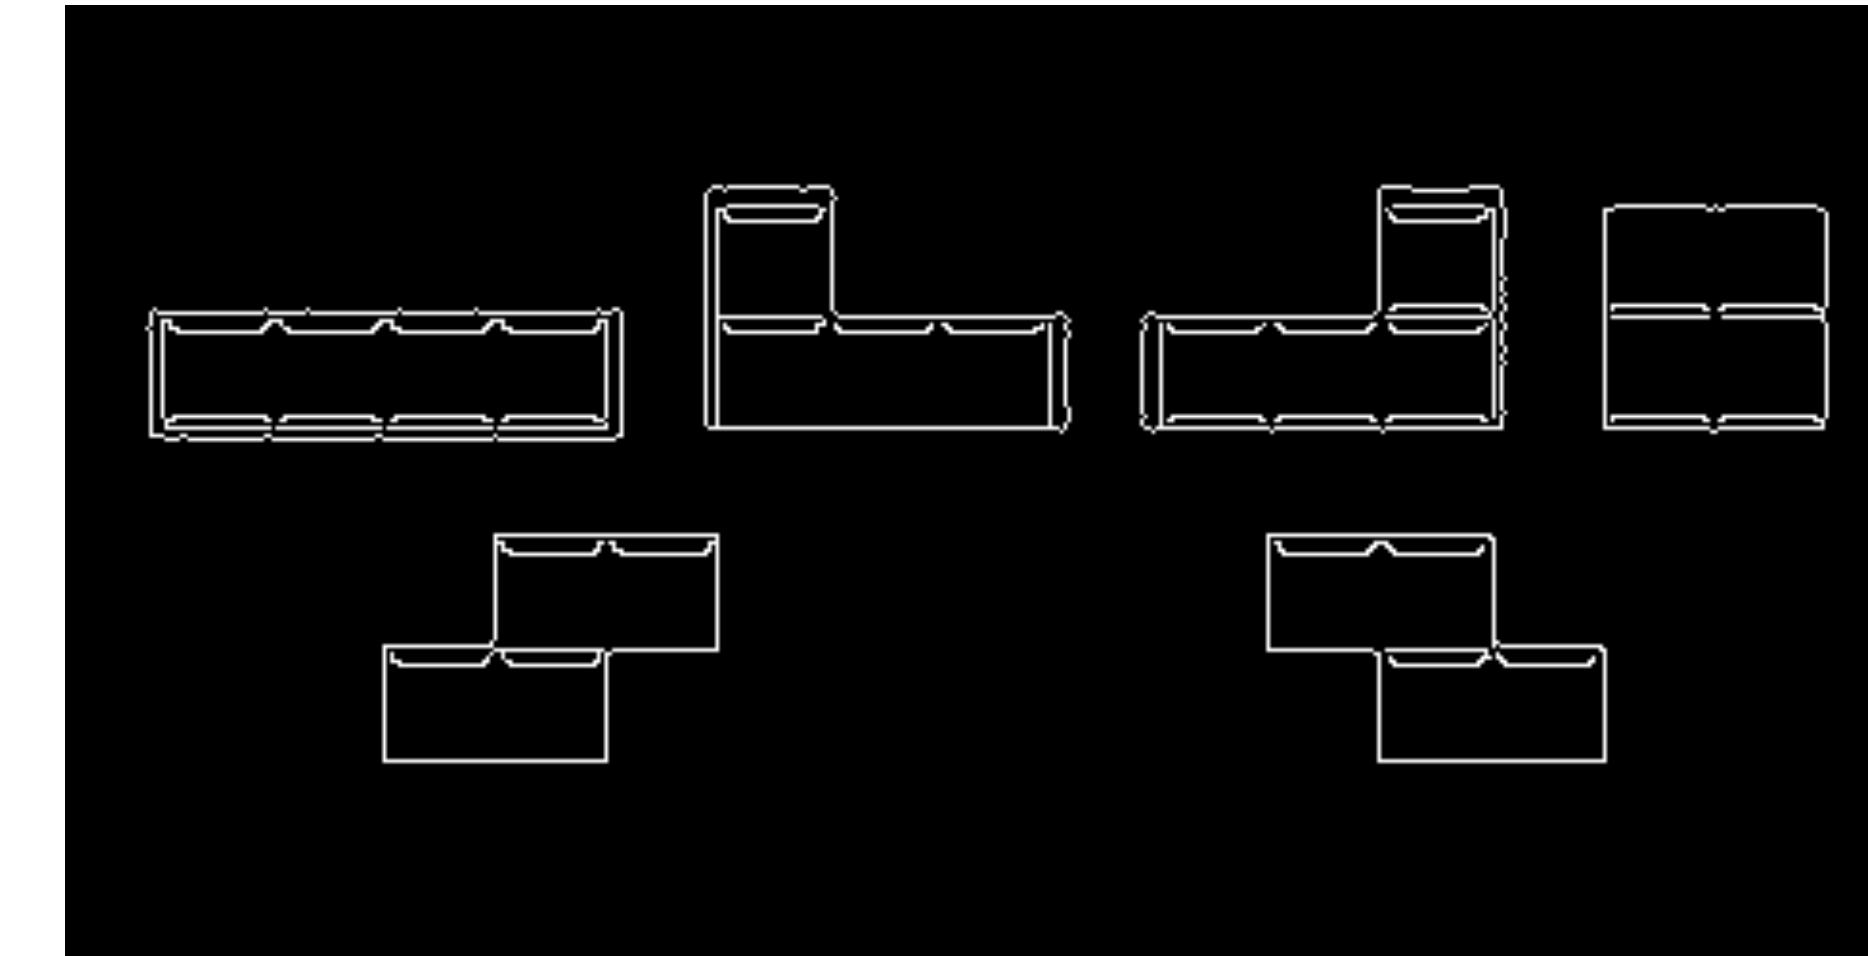
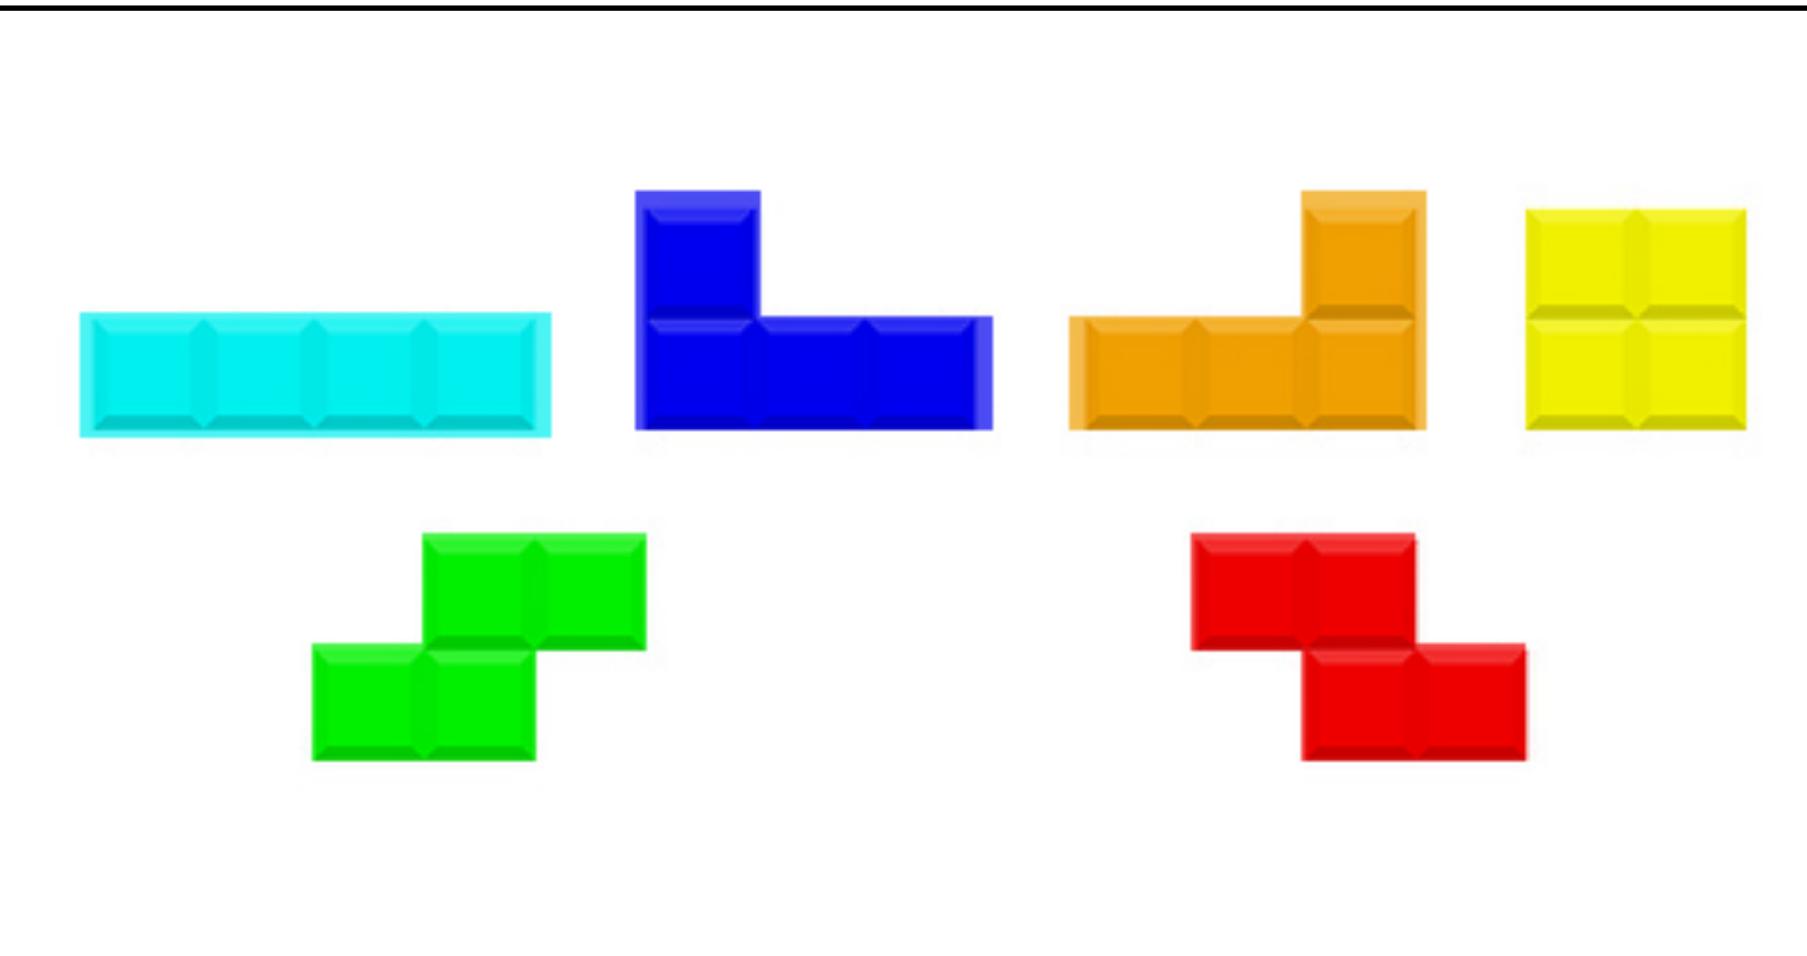
Converting to Grayscale

```
12 # load the input image (whose path was supplied via command line
13 # argument) and display the image to our screen
14 image = cv2.imread(args["image"])
15 cv2.imshow("Image", image)
16 cv2.waitKey(0)
17
18 # convert the image to grayscale
19 gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
20 cv2.imshow("Gray", gray)
21 cv2.waitKey(0)
```



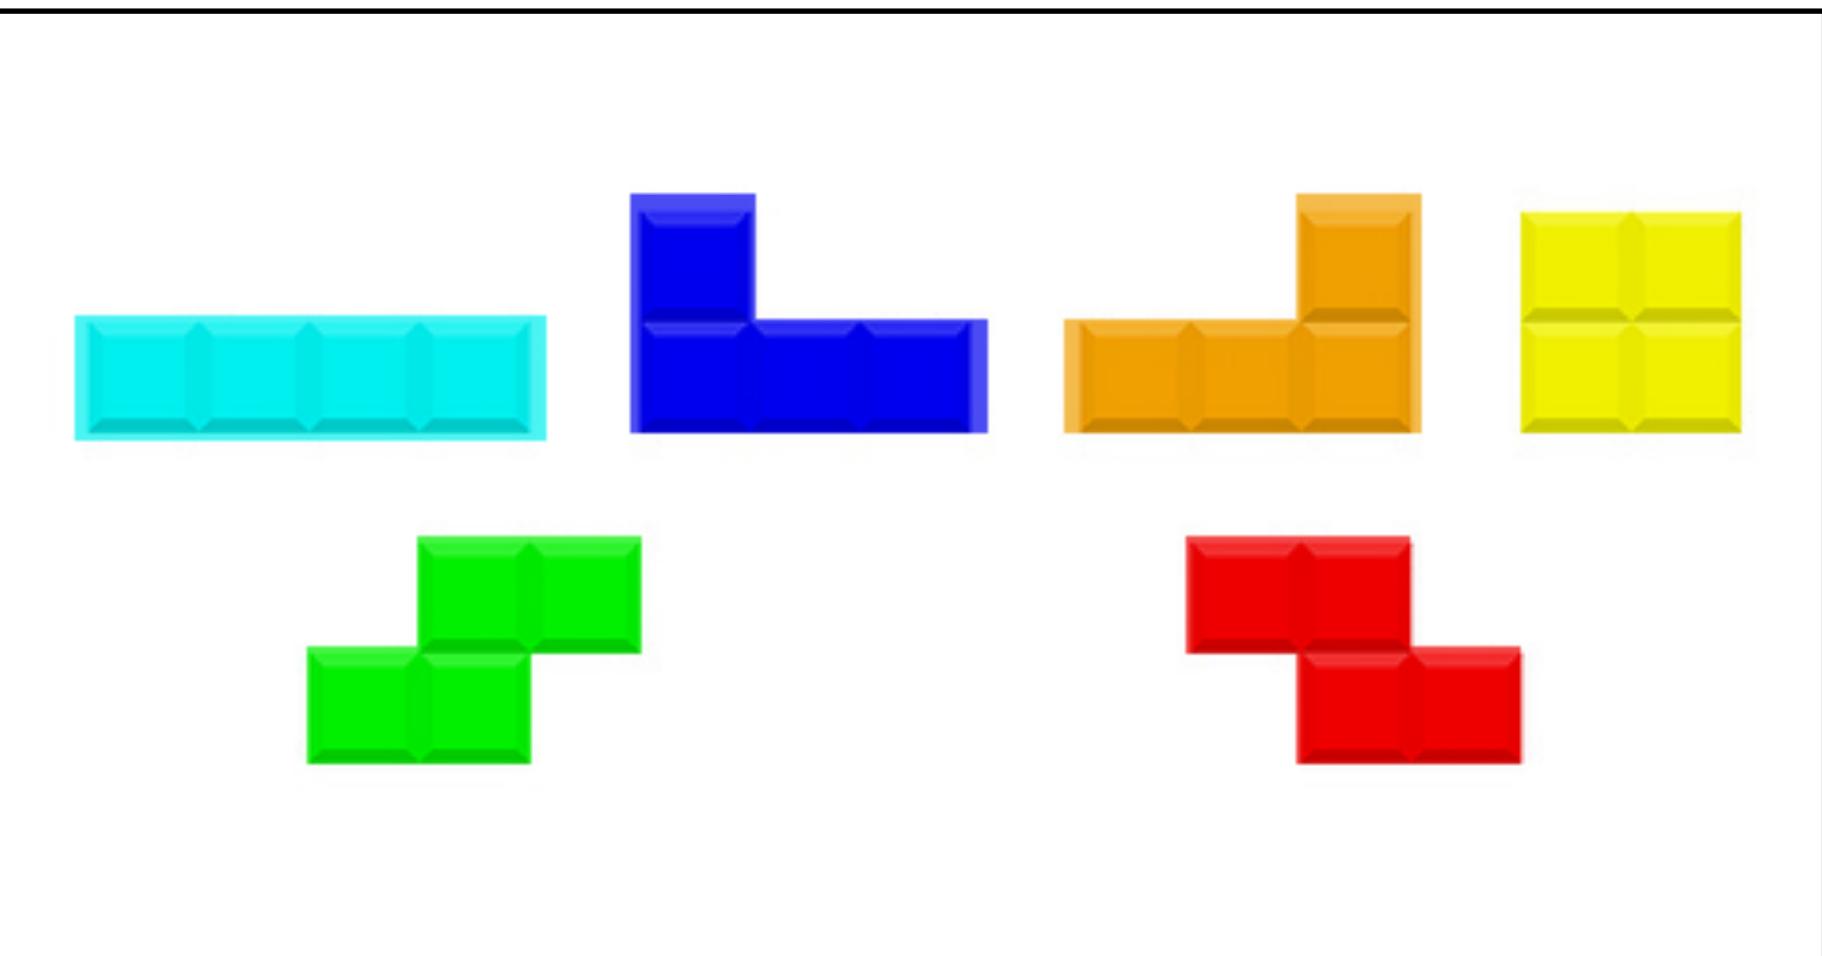
OpenCV Edge Detection

```
23 # applying edge detection we can find the outlines of objects in
24 # images
25 edged = cv2.Canny(gray, 30, 150)
26 cv2.imshow("Edged", edged)
27 cv2.waitKey(0)
```



OpenCV Thresholding

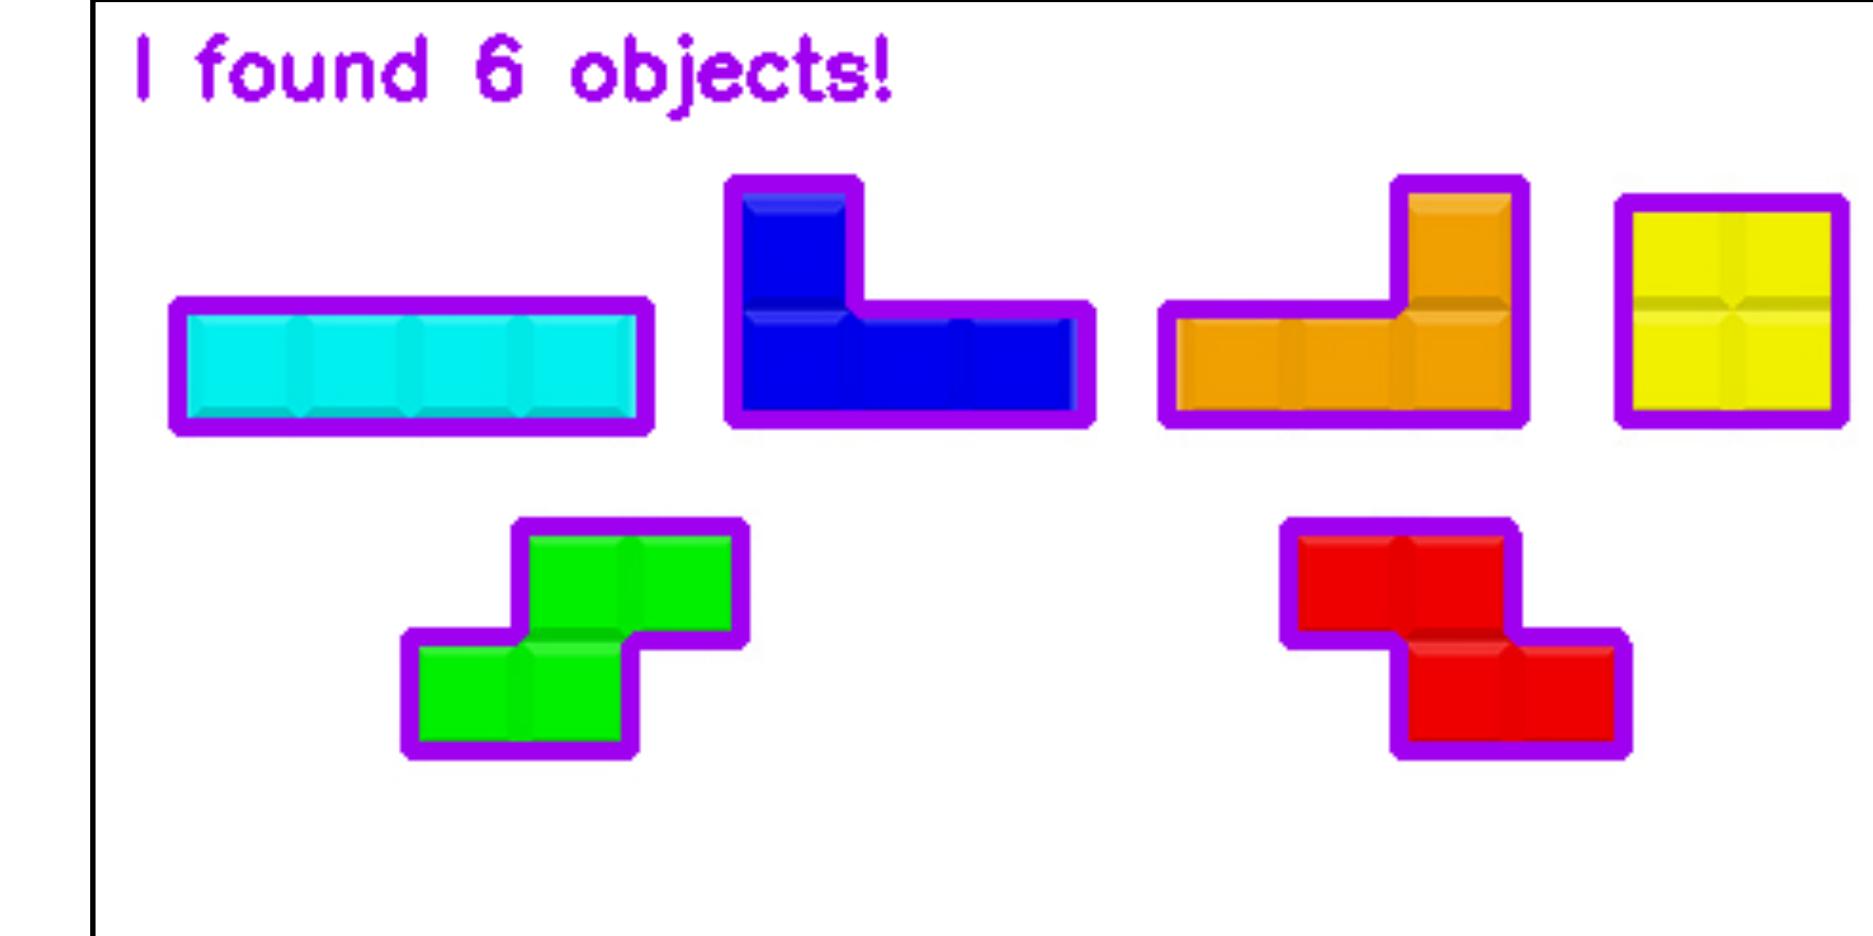
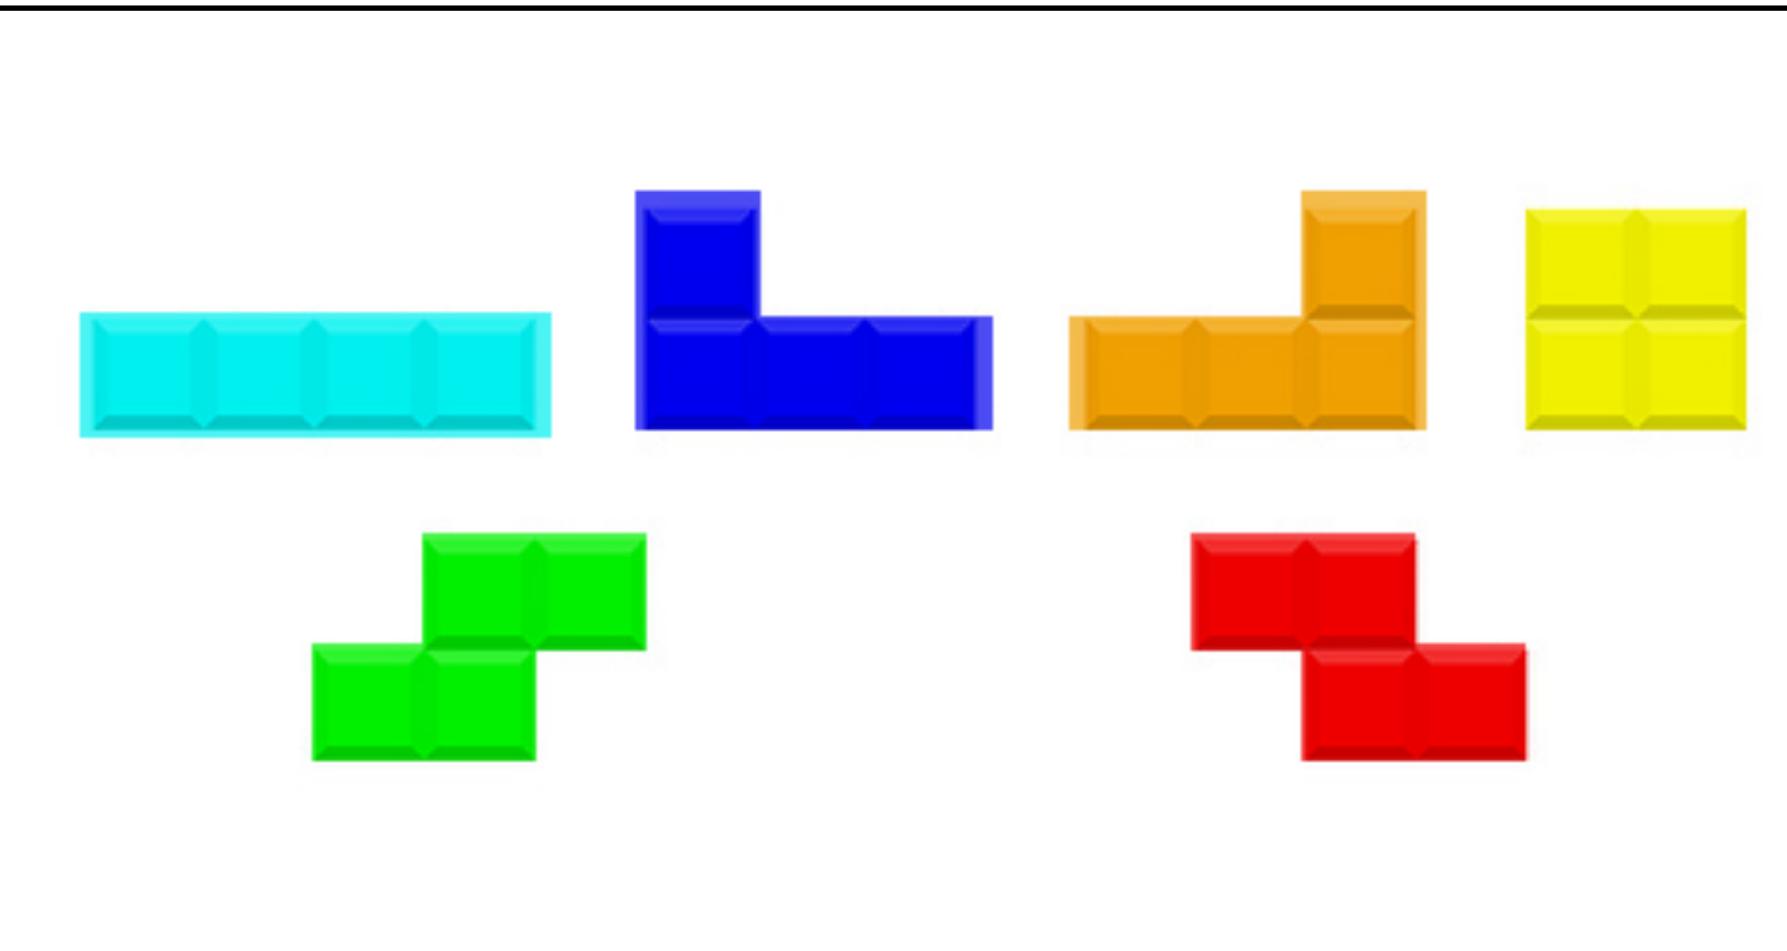
```
29 # threshold the image by setting all pixel values less than 225
30 # to 255 (white; foreground) and all pixel values >= 225 to 255
31 # (black; background), thereby segmenting the image
32 thresh = cv2.threshold(gray, 225, 255, cv2.THRESH_BINARY_INV)[1]
33 cv2.imshow("Thresh", thresh)
34 cv2.waitKey(0)
```



OpenCV

Detecting/Drawing Contours

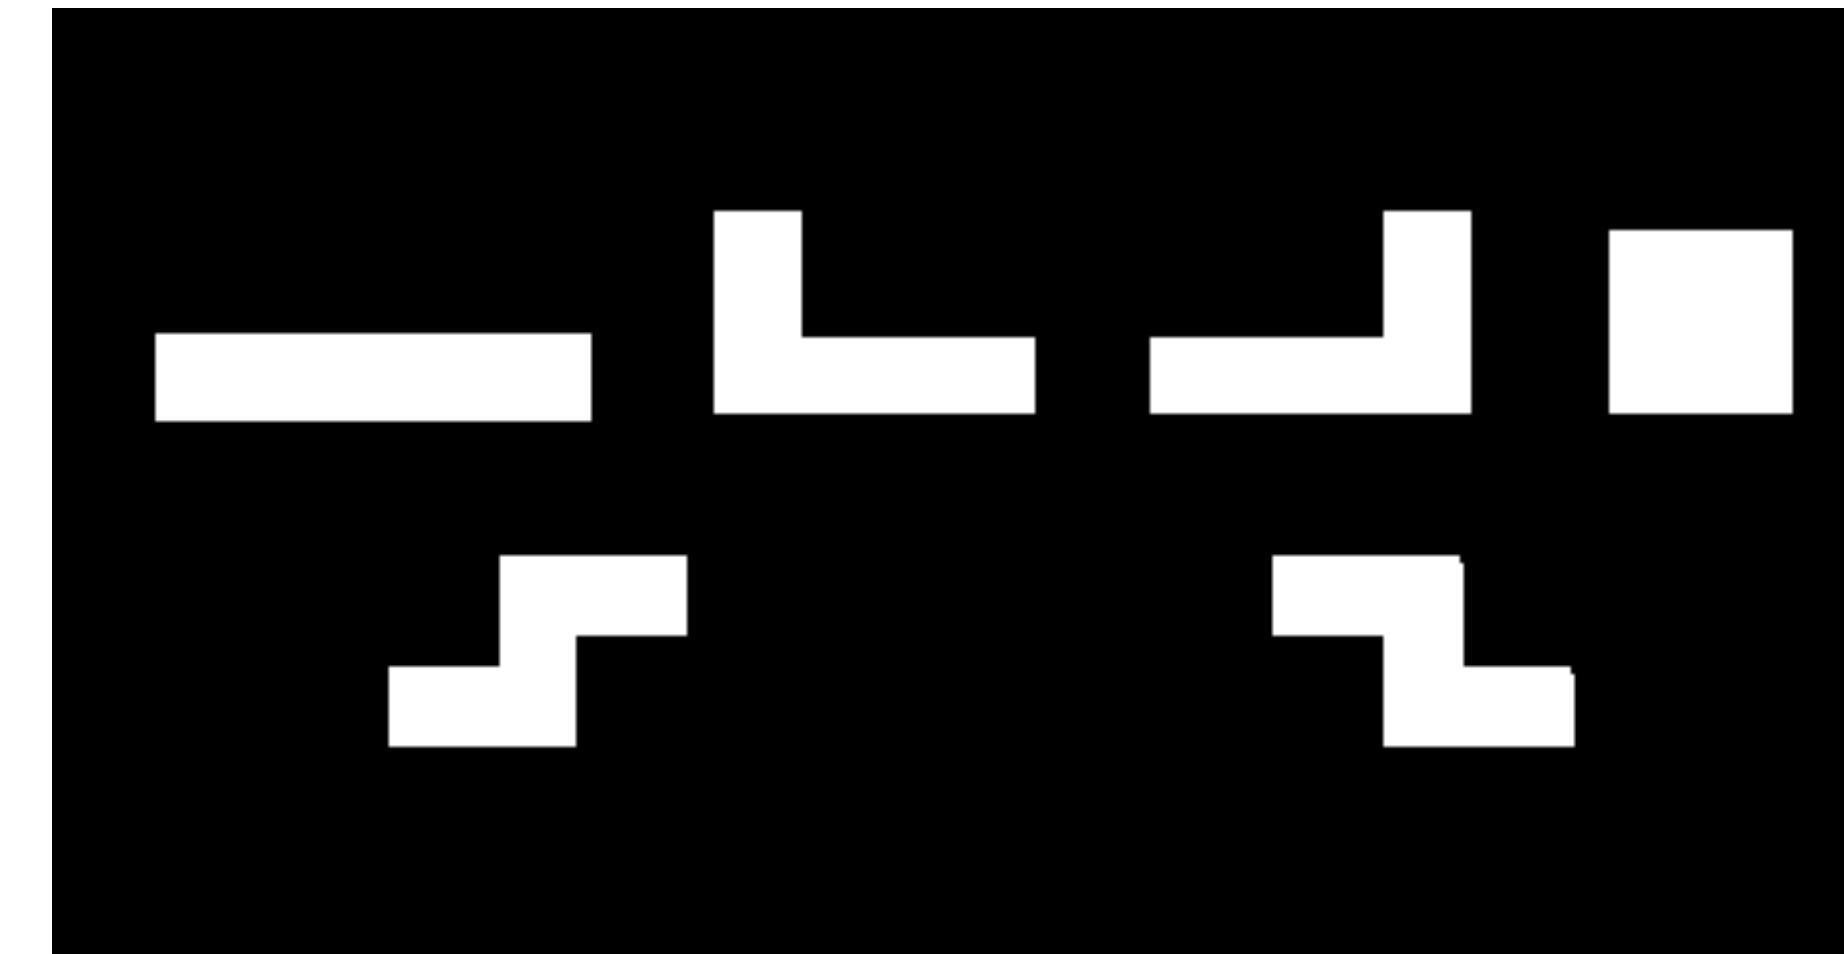
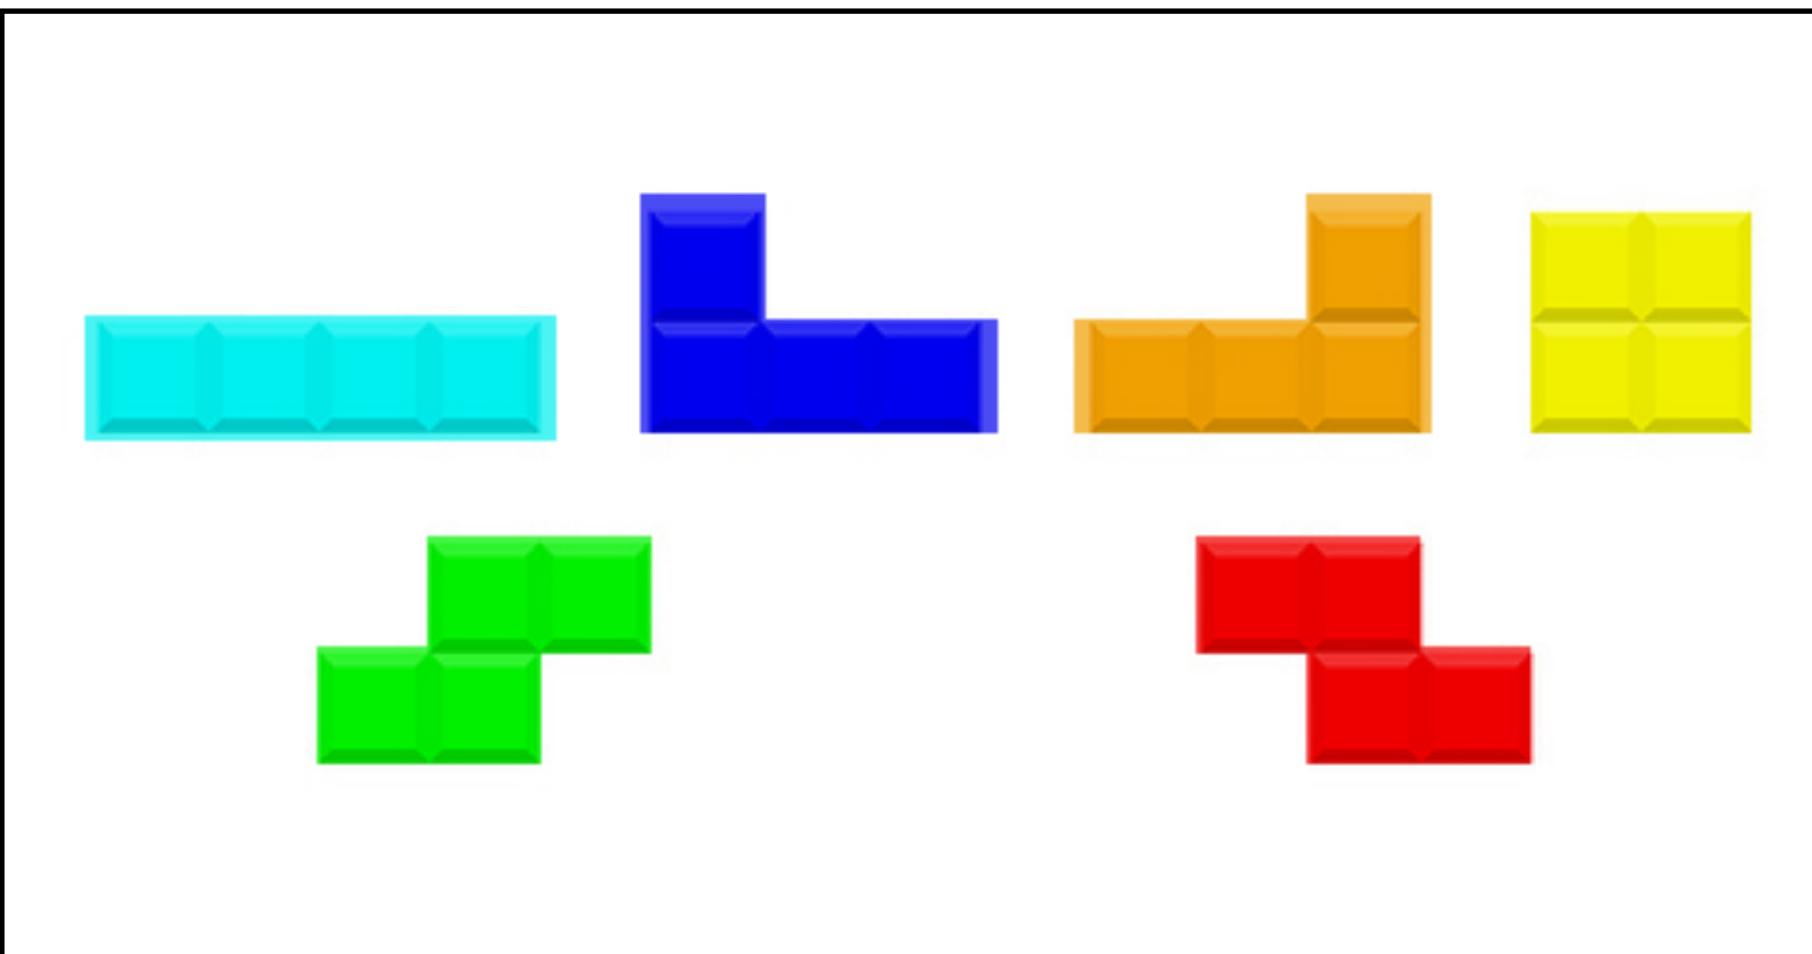
```
36 # find contours (i.e., outlines) of the foreground objects in the
37 # thresholded image
38 cnts = cv2.findContours(thresh.copy(), cv2.RETR_EXTERNAL,
39 |   cv2.CHAIN_APPROX_SIMPLE)
40 cnts = imutils.grab_contours(cnts)
41 output = image.copy()
42 # loop over the contours
43 for c in cnts:
44     # draw each contour on the output image with a 3px thick purple
45     # outline, then display the output contours one at a time
46     cv2.drawContours(output, [c], -1, (240, 0, 159), 3)
47     cv2.imshow("Contours", output)
48     cv2.waitKey(0)
```



OpenCV

Erosions and Dilations

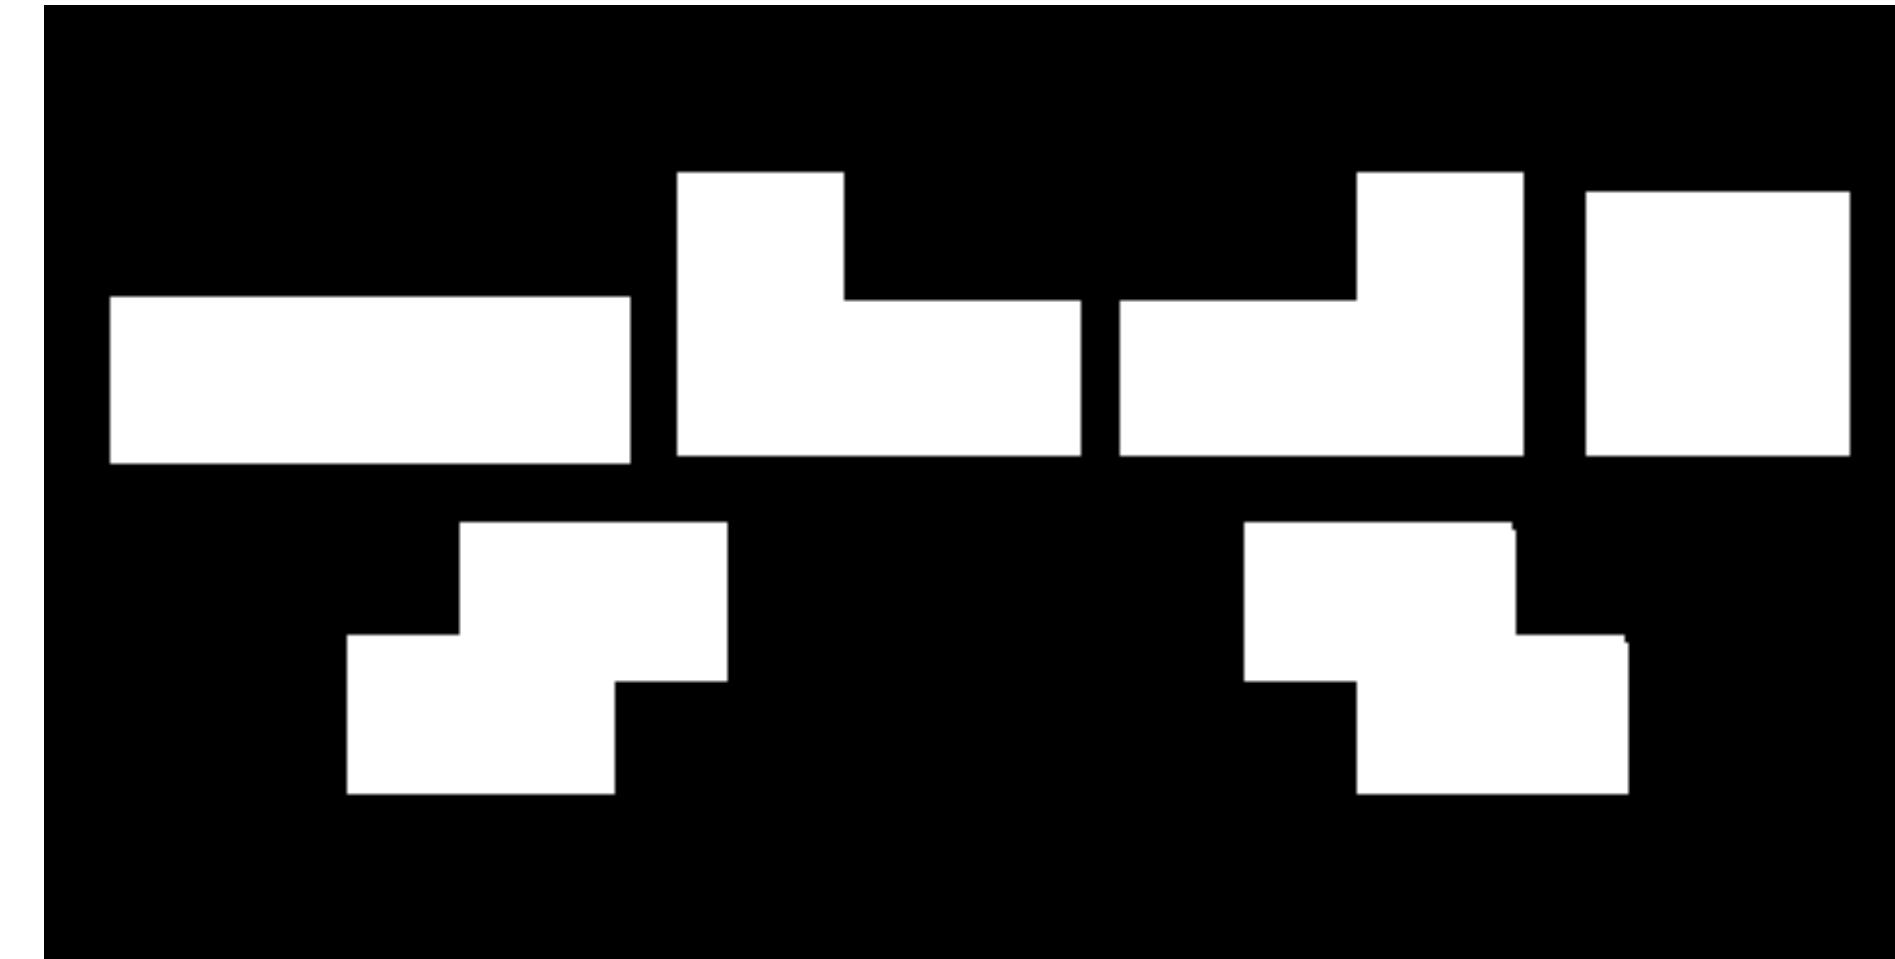
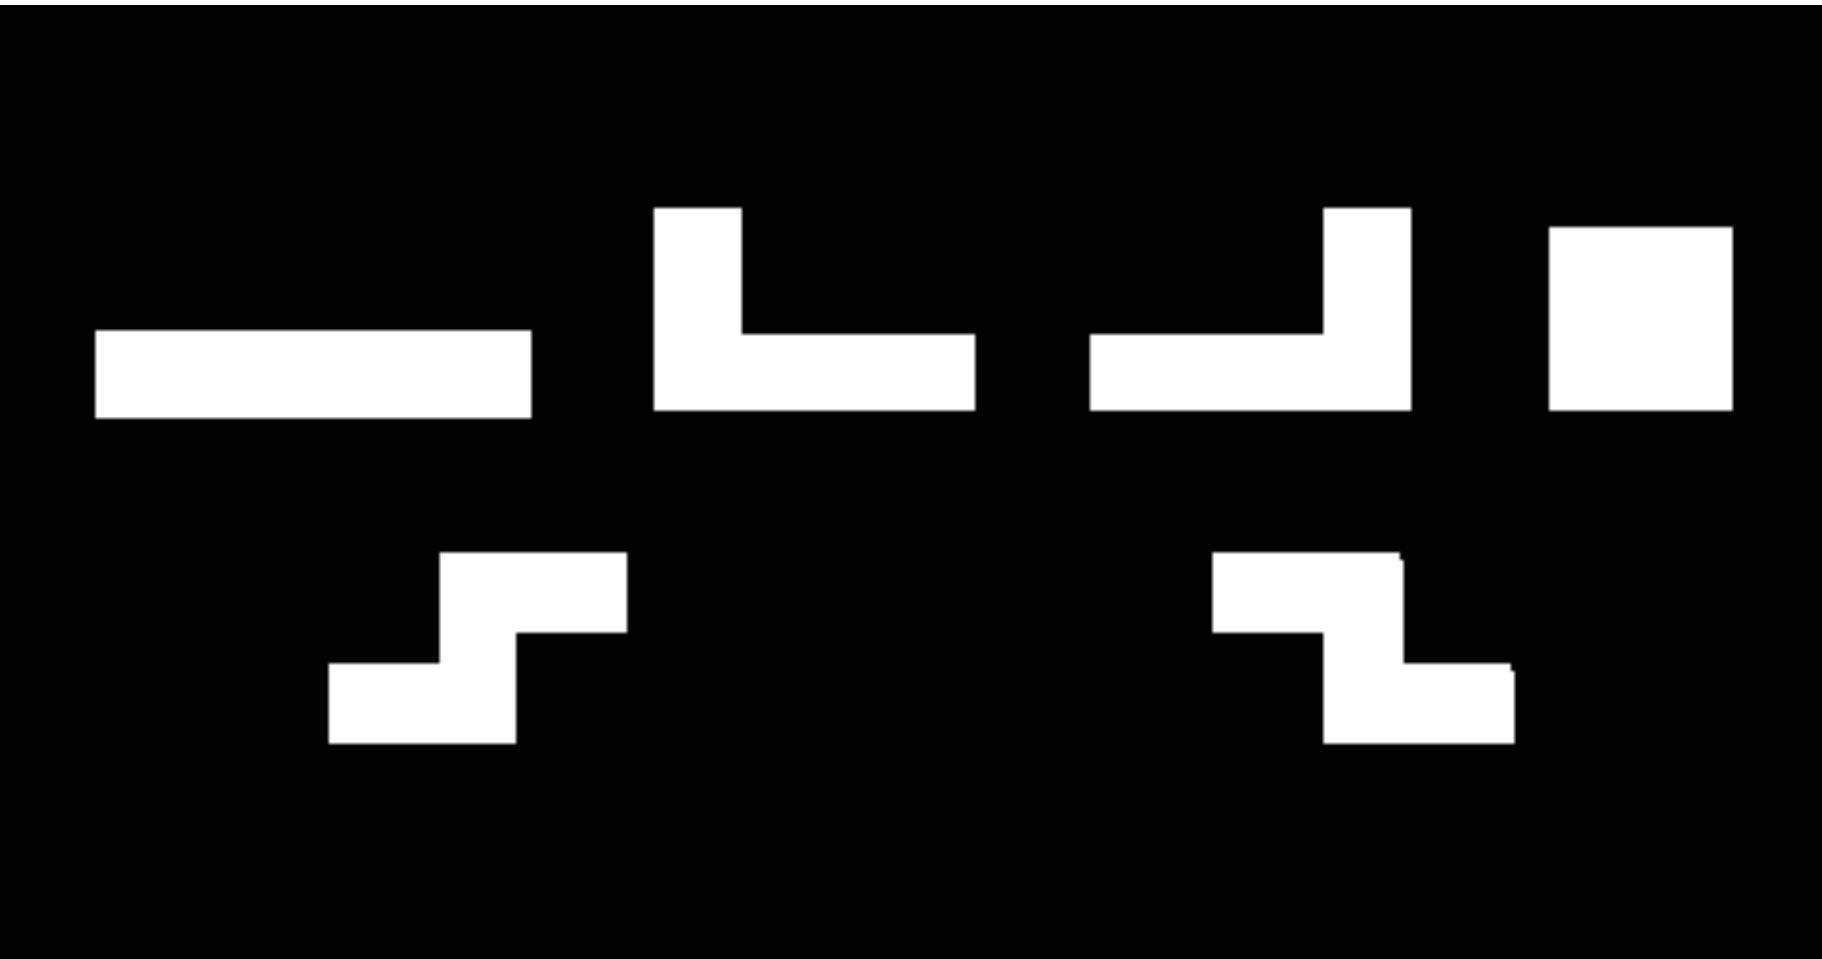
```
58 # we apply erosions to reduce the size of foreground objects
59 mask = thresh.copy()
60 mask = cv2.erode(mask, None, iterations=5)
61 cv2.imshow("Eroded", mask)
62 cv2.waitKey(0)
```



OpenCV

Erosions and Dilations

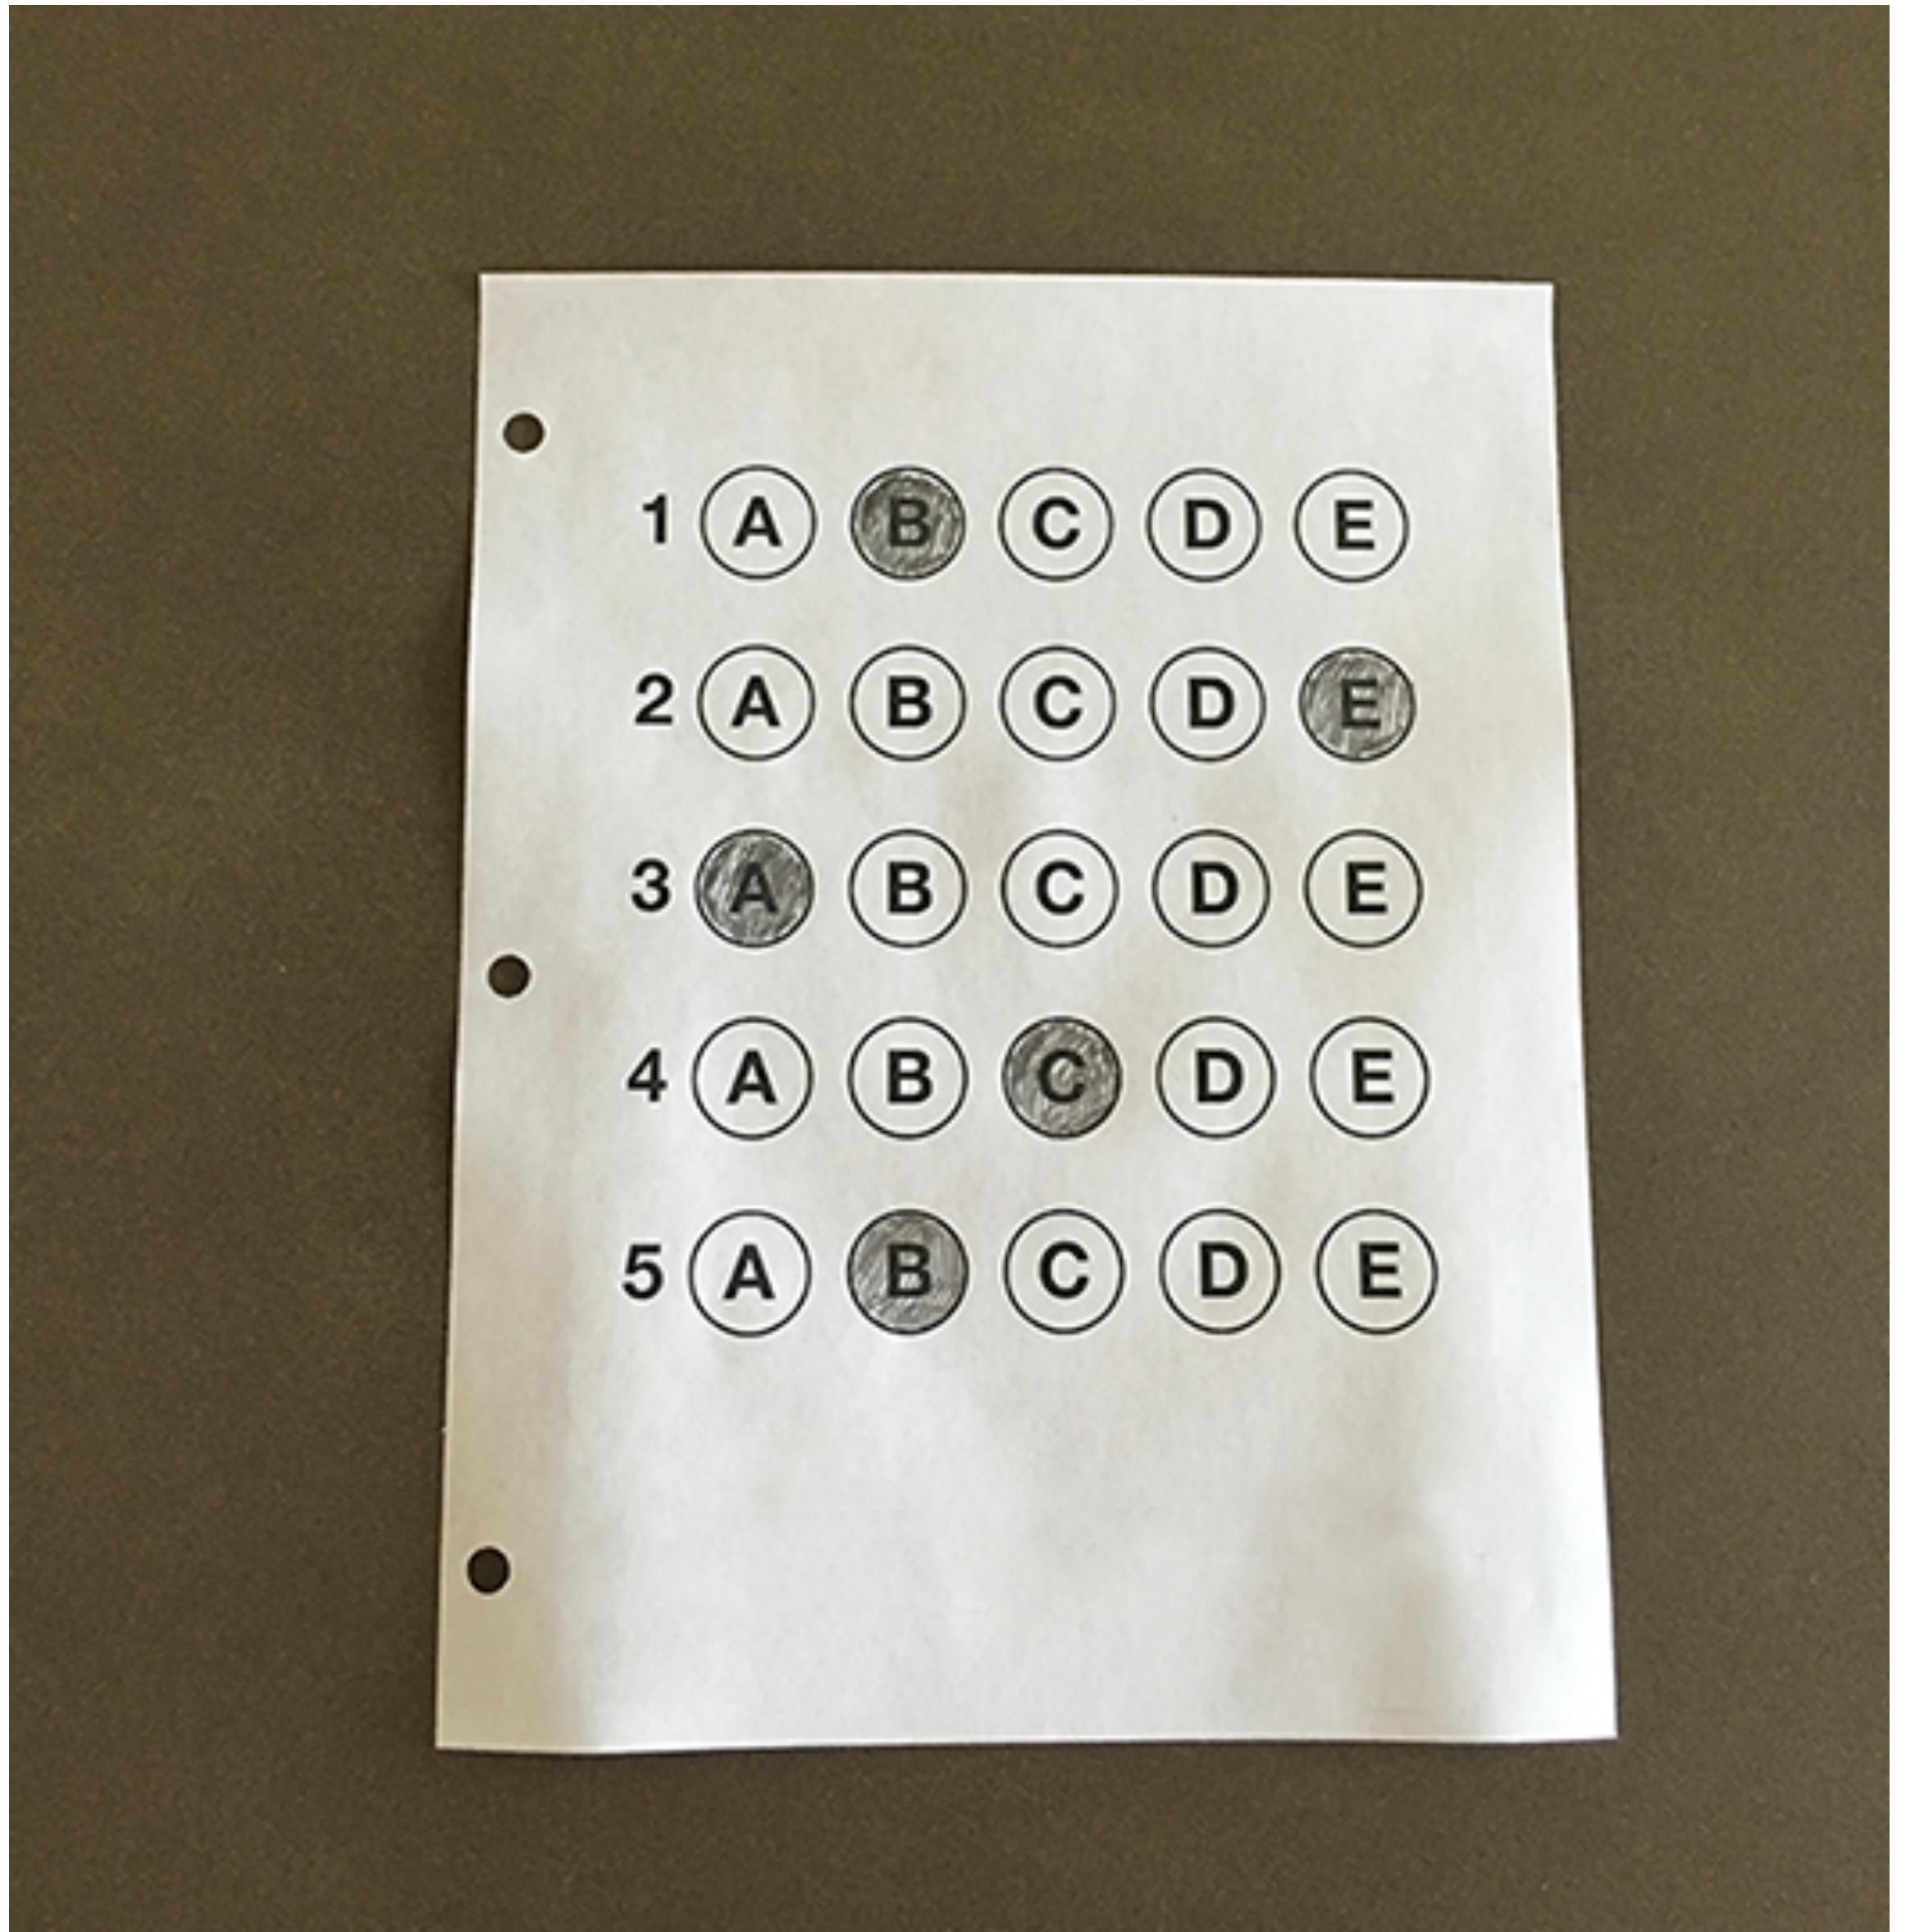
```
64 # similarly, dilations can increase the size of the ground objects
65 mask = thresh.copy()
66 mask = cv2.dilate(mask, None, iterations=5)
67 cv2.imshow("Dilated", mask)
68 cv2.waitKey(0)
```



OMR

Bubble Sheet Scanner

- Detect exam in an image
- Apply transform to extract top-down view of exam
- Extract set of bubbles
- Sort questions/bubbles
- Determine the marked bubble in each row
- Lookup correct answer
- Repeat for all questions in exam



OMR

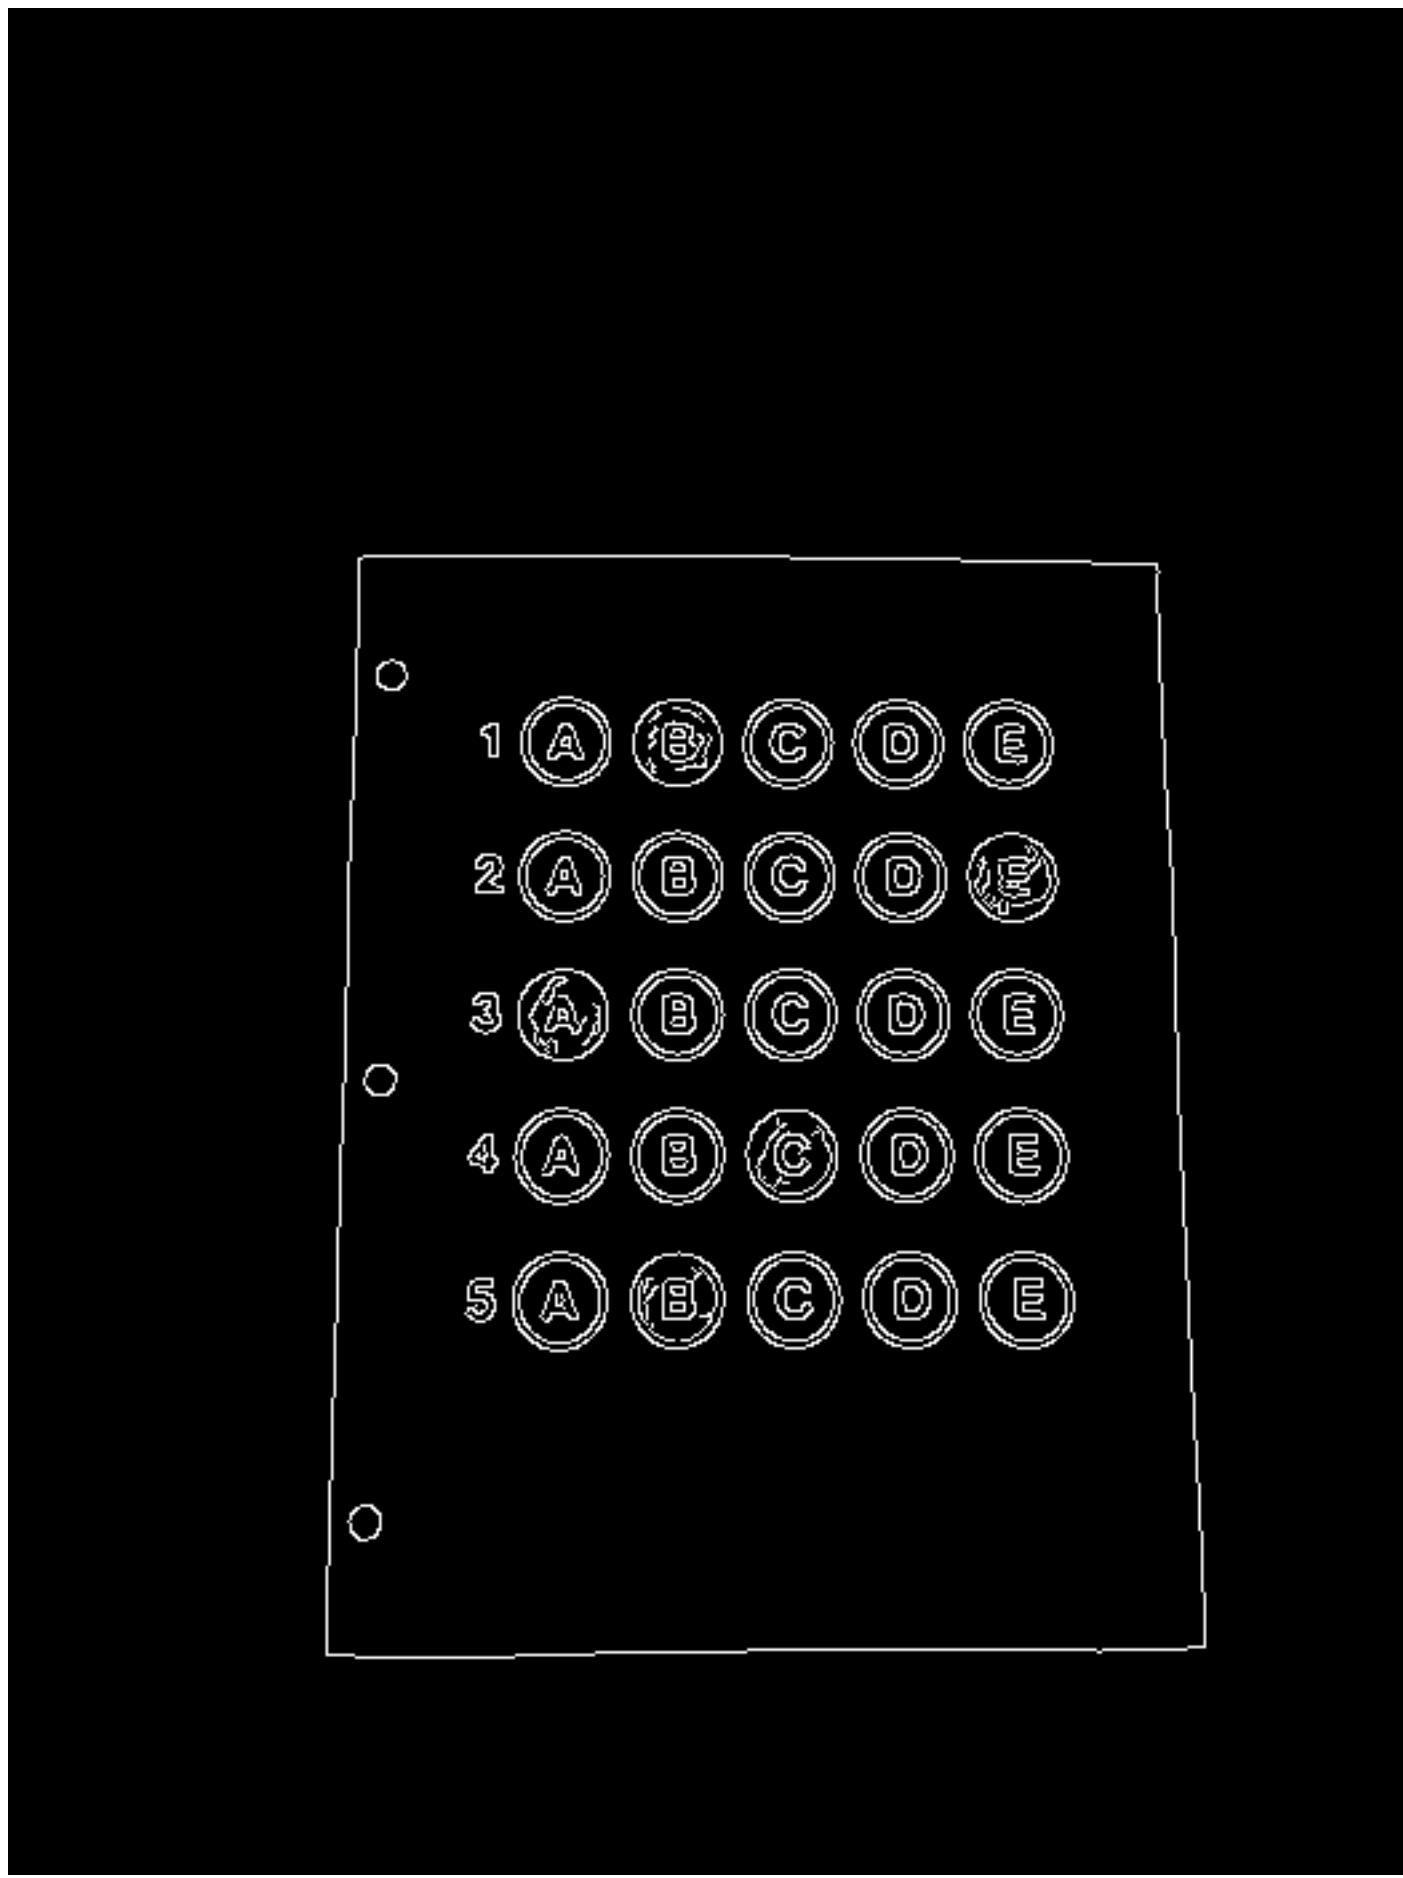
Test Grader

```
1 # import the necessary packages
2 from imutils.perspective import four_point_transform
3 from imutils import contours
4 import numpy as np
5 import argparse
6 import imutils
7 import cv2
8
9 # construct the argument parse and parse the arguments
10 ap = argparse.ArgumentParser()
11 ap.add_argument("-i", "--image", required=True,
12                 help="path to the input image")
13 args = vars(ap.parse_args())
14
15 # define the answer key which maps the question number
16 # to the correct answer
17 ANSWER_KEY = {0: 1, 1: 4, 2: 0, 3: 3, 4: 1}
18
```

Question #1: B
Question #2: E
Question #3: A
Question #4: D
Question #5: B

OMR

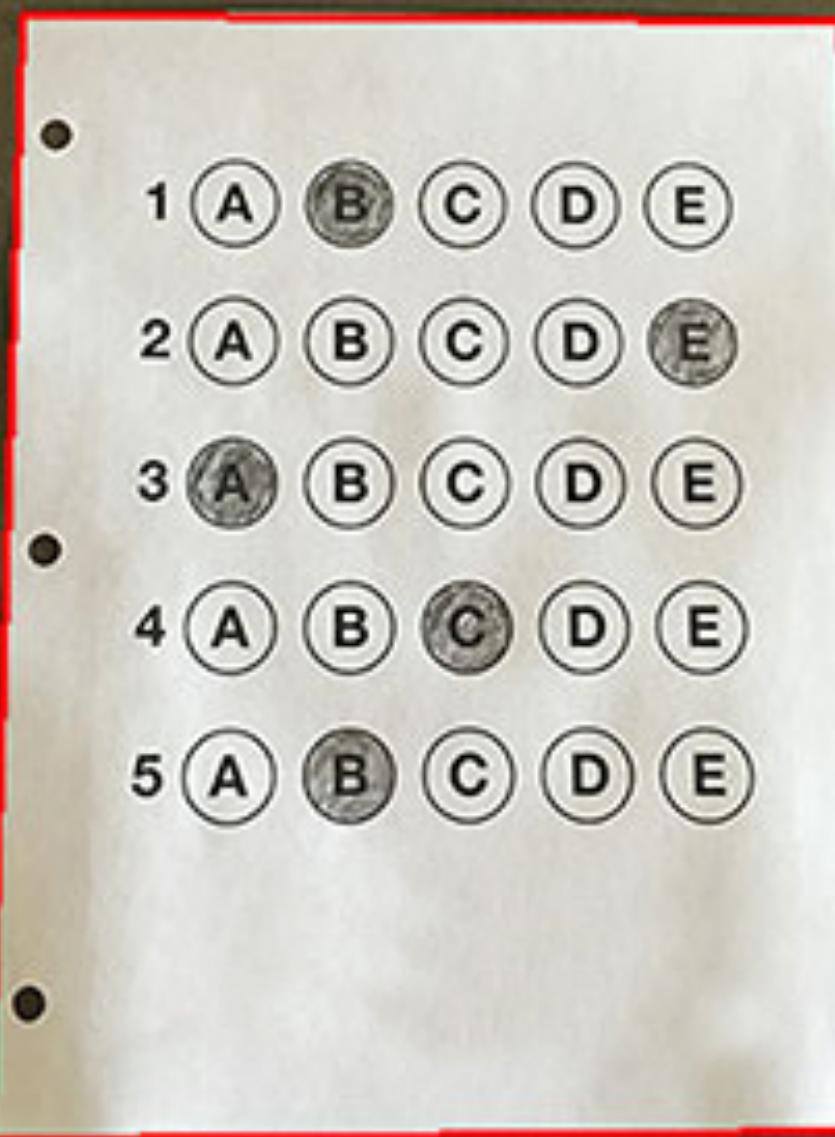
Pre-Process Image



```
19 # load the image, convert it to grayscale, blur it
20 # slightly, then find edges
21 image = cv2.imread(args["image"])
22 gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
23 blurred = cv2.GaussianBlur(gray, (5, 5), 0)
24 edged = cv2.Canny(blurred, 75, 200)
25
```

OMR

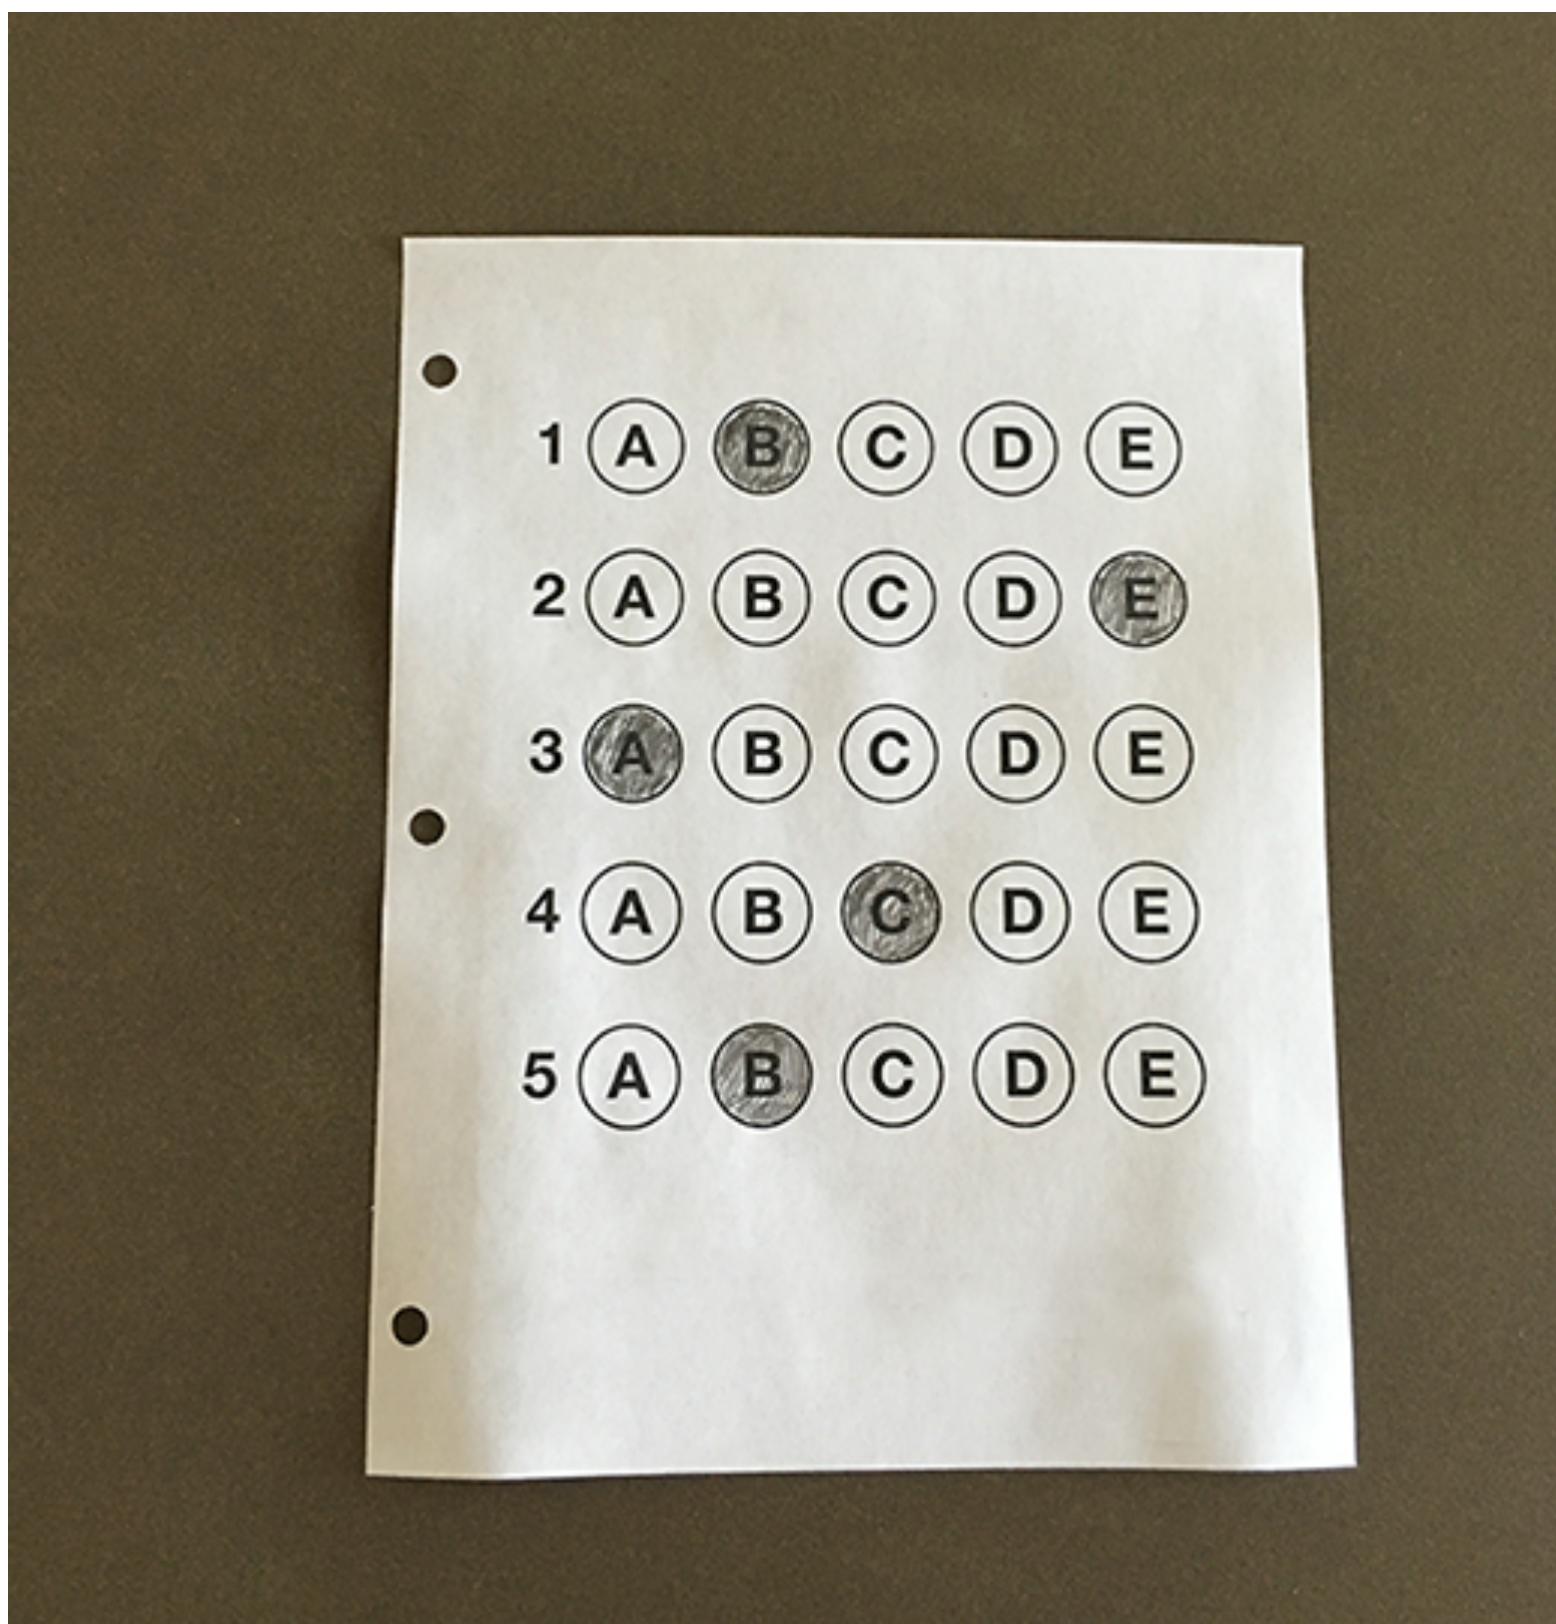
Find the exam



```
26 # find contours in the edge map, then initialize
27 # the contour that corresponds to the document
28 cnts = cv2.findContours(edged.copy(), cv2.RETR_EXTERNAL,
29 |   cv2.CHAIN_APPROX_SIMPLE)
30 cnts = imutils.grab_contours(cnts)
31 docCnt = None
32
33 # ensure that at least one contour was found
34 if len(cnts) > 0:
35     # sort the contours according to their size in
36     # descending order
37     cnts = sorted(cnts, key=cv2.contourArea, reverse=True)
38
39     # loop over the sorted contours
40     for c in cnts:
41         # approximate the contour
42         peri = cv2.arcLength(c, True)
43         approx = cv2.approxPolyDP(c, 0.02 * peri, True)
44
45         # if our approximated contour has four points,
46         # then we can assume we have found the paper
47         if len(approx) == 4:
48             docCnt = approx
49             break
50
```

OMR

Top-down view of document



```
51 # apply a four point perspective transform to both the
52 # original image and grayscale image to obtain a top-down
53 # birds eye view of the paper
54 paper = four_point_transform(image, docCnt.reshape(4, 2))
55 warped = four_point_transform(gray, docCnt.reshape(4, 2))
56
```

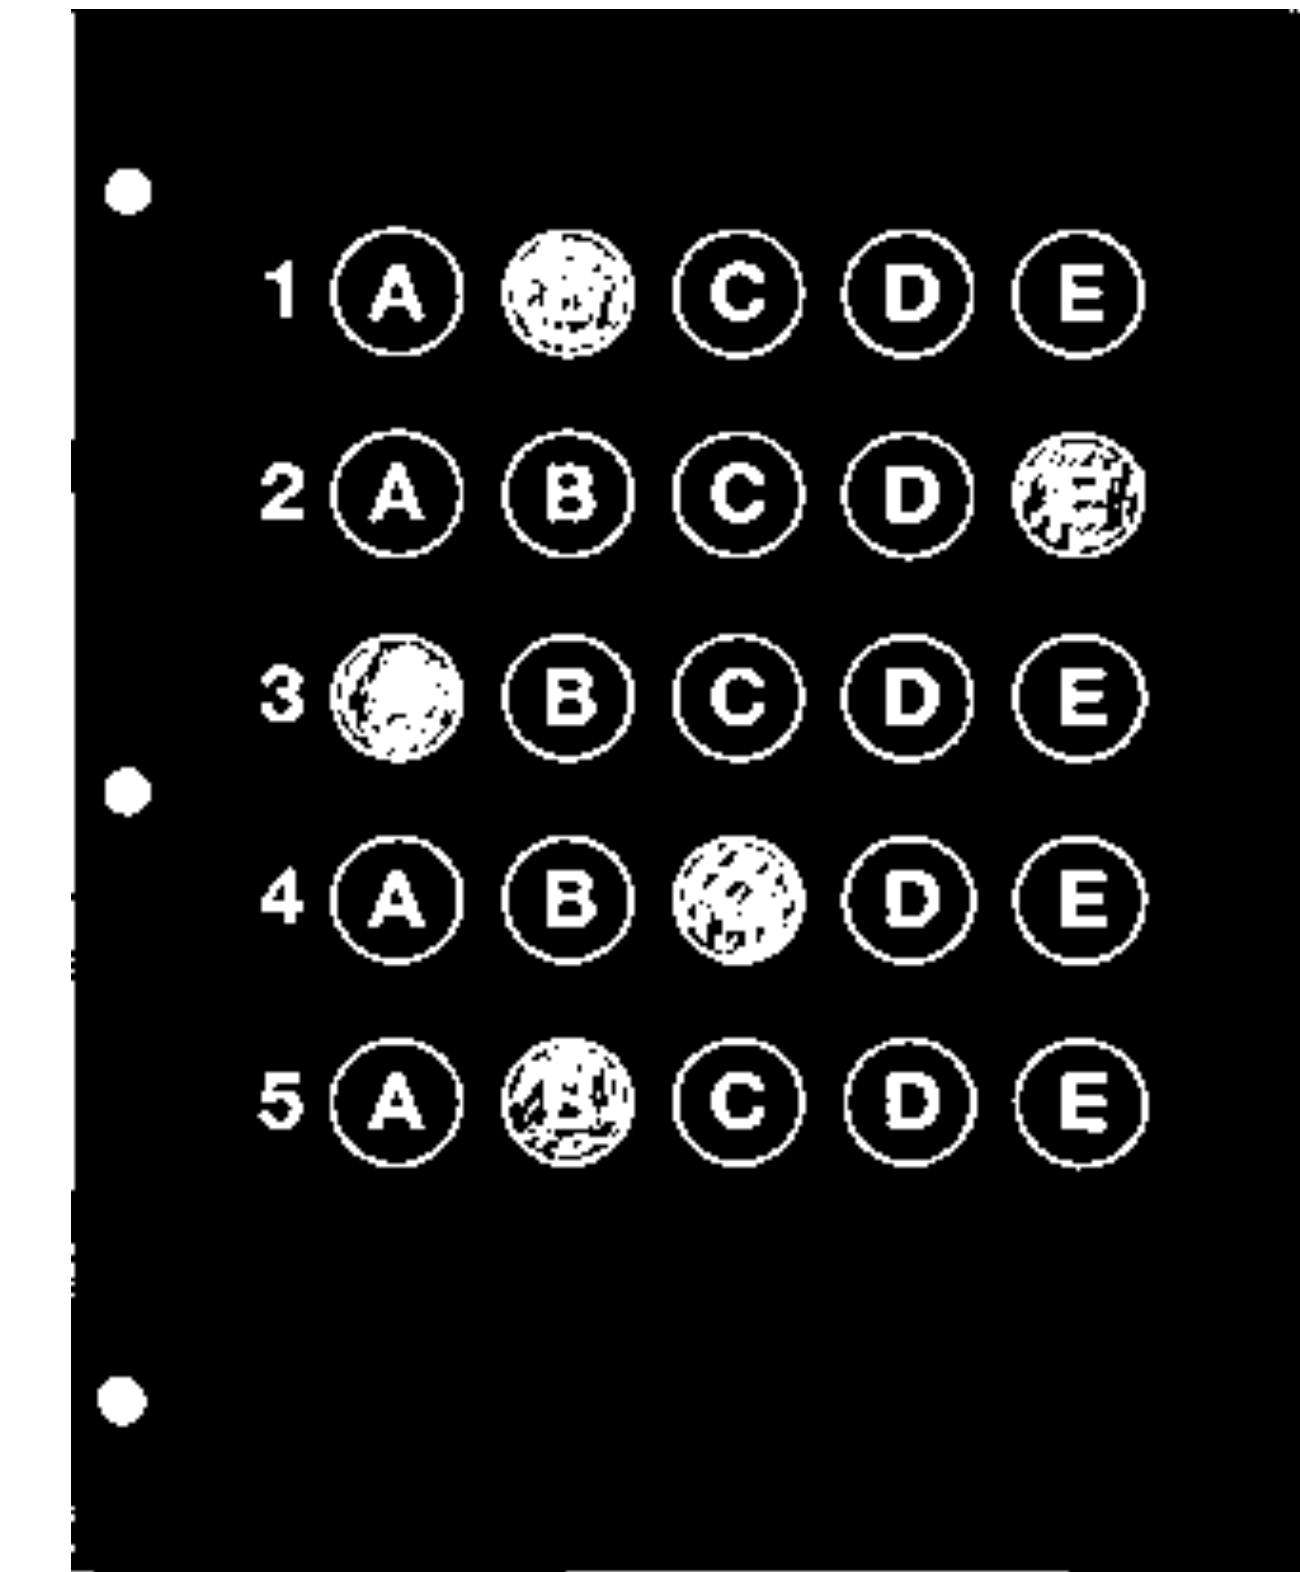


OMR

Top-down view of document

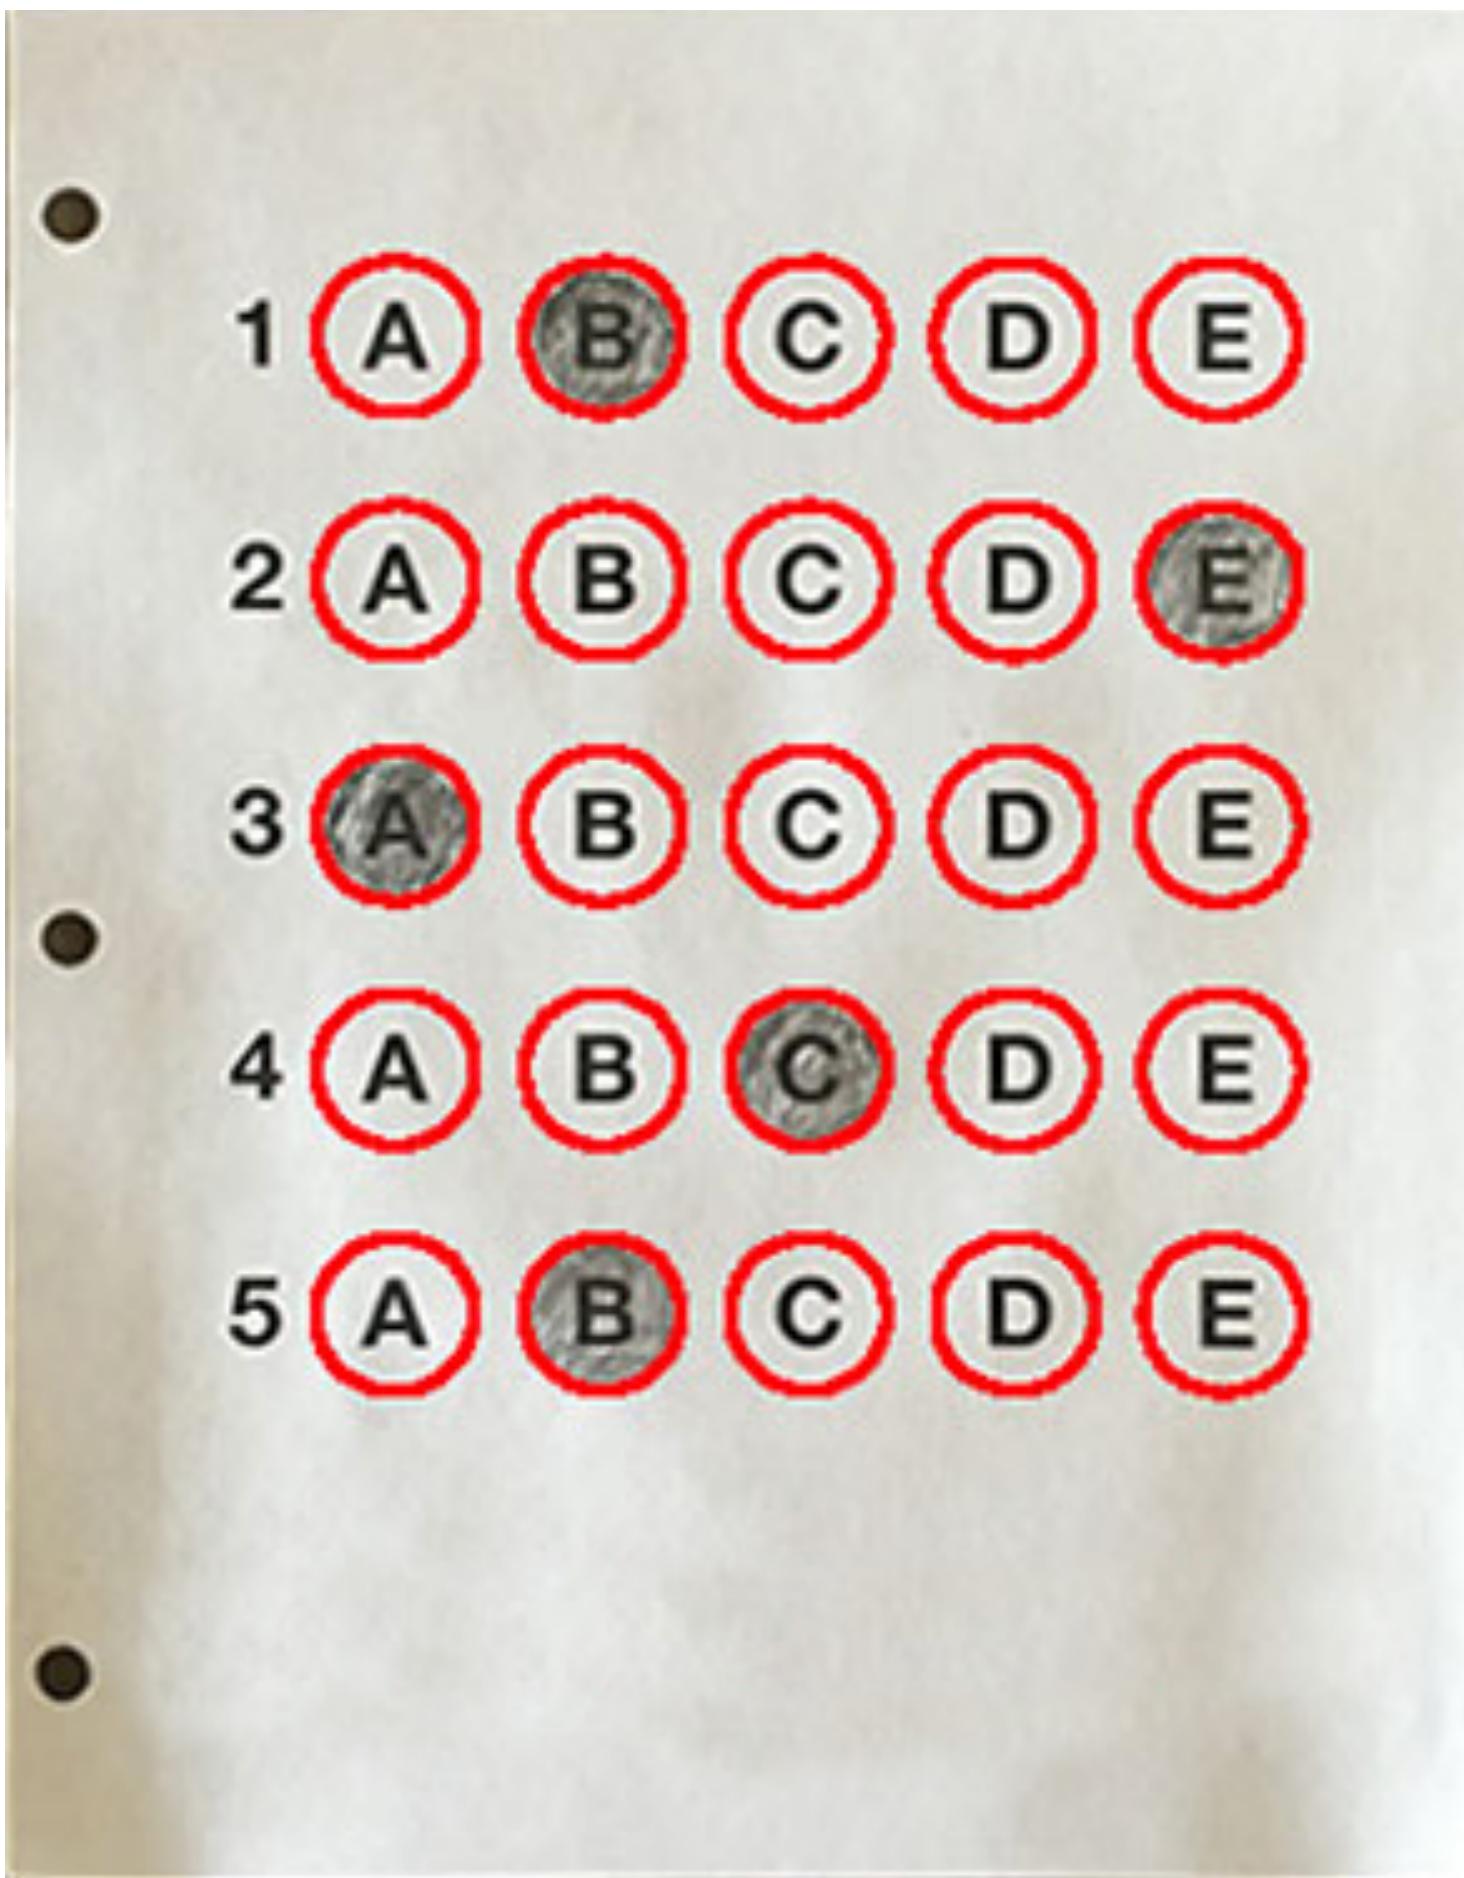


```
57 # apply Otsu's thresholding method to binarize the warped  
58 # piece of paper  
59 thresh = cv2.threshold(warped, 0, 255,  
60 |   cv2.THRESH_BINARY_INV | cv2.THRESH_OTSU)[1]  
61
```



OMR

Finding Bubbles



```
62 # find contours in the thresholded image, then initialize
63 # the list of contours that correspond to questions
64 cnts = cv2.findContours(thresh.copy(), cv2.RETR_EXTERNAL,
65     cv2.CHAIN_APPROX_SIMPLE)
66 cnts = imutils.grab_contours(cnts)
67 questionCnts = []
68 # loop over the contours
69 for c in cnts:
70     # compute the bounding box of the contour, then use the
71     # bounding box to derive the aspect ratio
72     (x, y, w, h) = cv2.boundingRect(c)
73     ar = w / float(h)
74     # in order to label the contour as a question, region
75     # should be sufficiently wide, sufficiently tall, and
76     # have an aspect ratio approximately equal to 1
77     if w >= 20 and h >= 20 and ar >= 0.9 and ar <= 1.1:
78         questionCnts.append(c)
79 |
```

OMR Grading



```
82 # sort the question contours top-to-bottom, then initialize
83 # the total number of correct answers
84 questionCnts = contours.sort_contours(questionCnts,
85 |   method="top-to-bottom")[0]
86 correct = 0
87 # each question has 5 possible answers, to loop over the
88 # question in batches of 5
89 for (q, i) in enumerate(np.arange(0, len(questionCnts), 5)):
90     # sort the contours for the current question from
91     # left to right, then initialize the index of the
92     # bubbled answer
93     cnts = contours.sort_contours(questionCnts[i:i + 5])[0]
94     bubbled = None
```

OMR

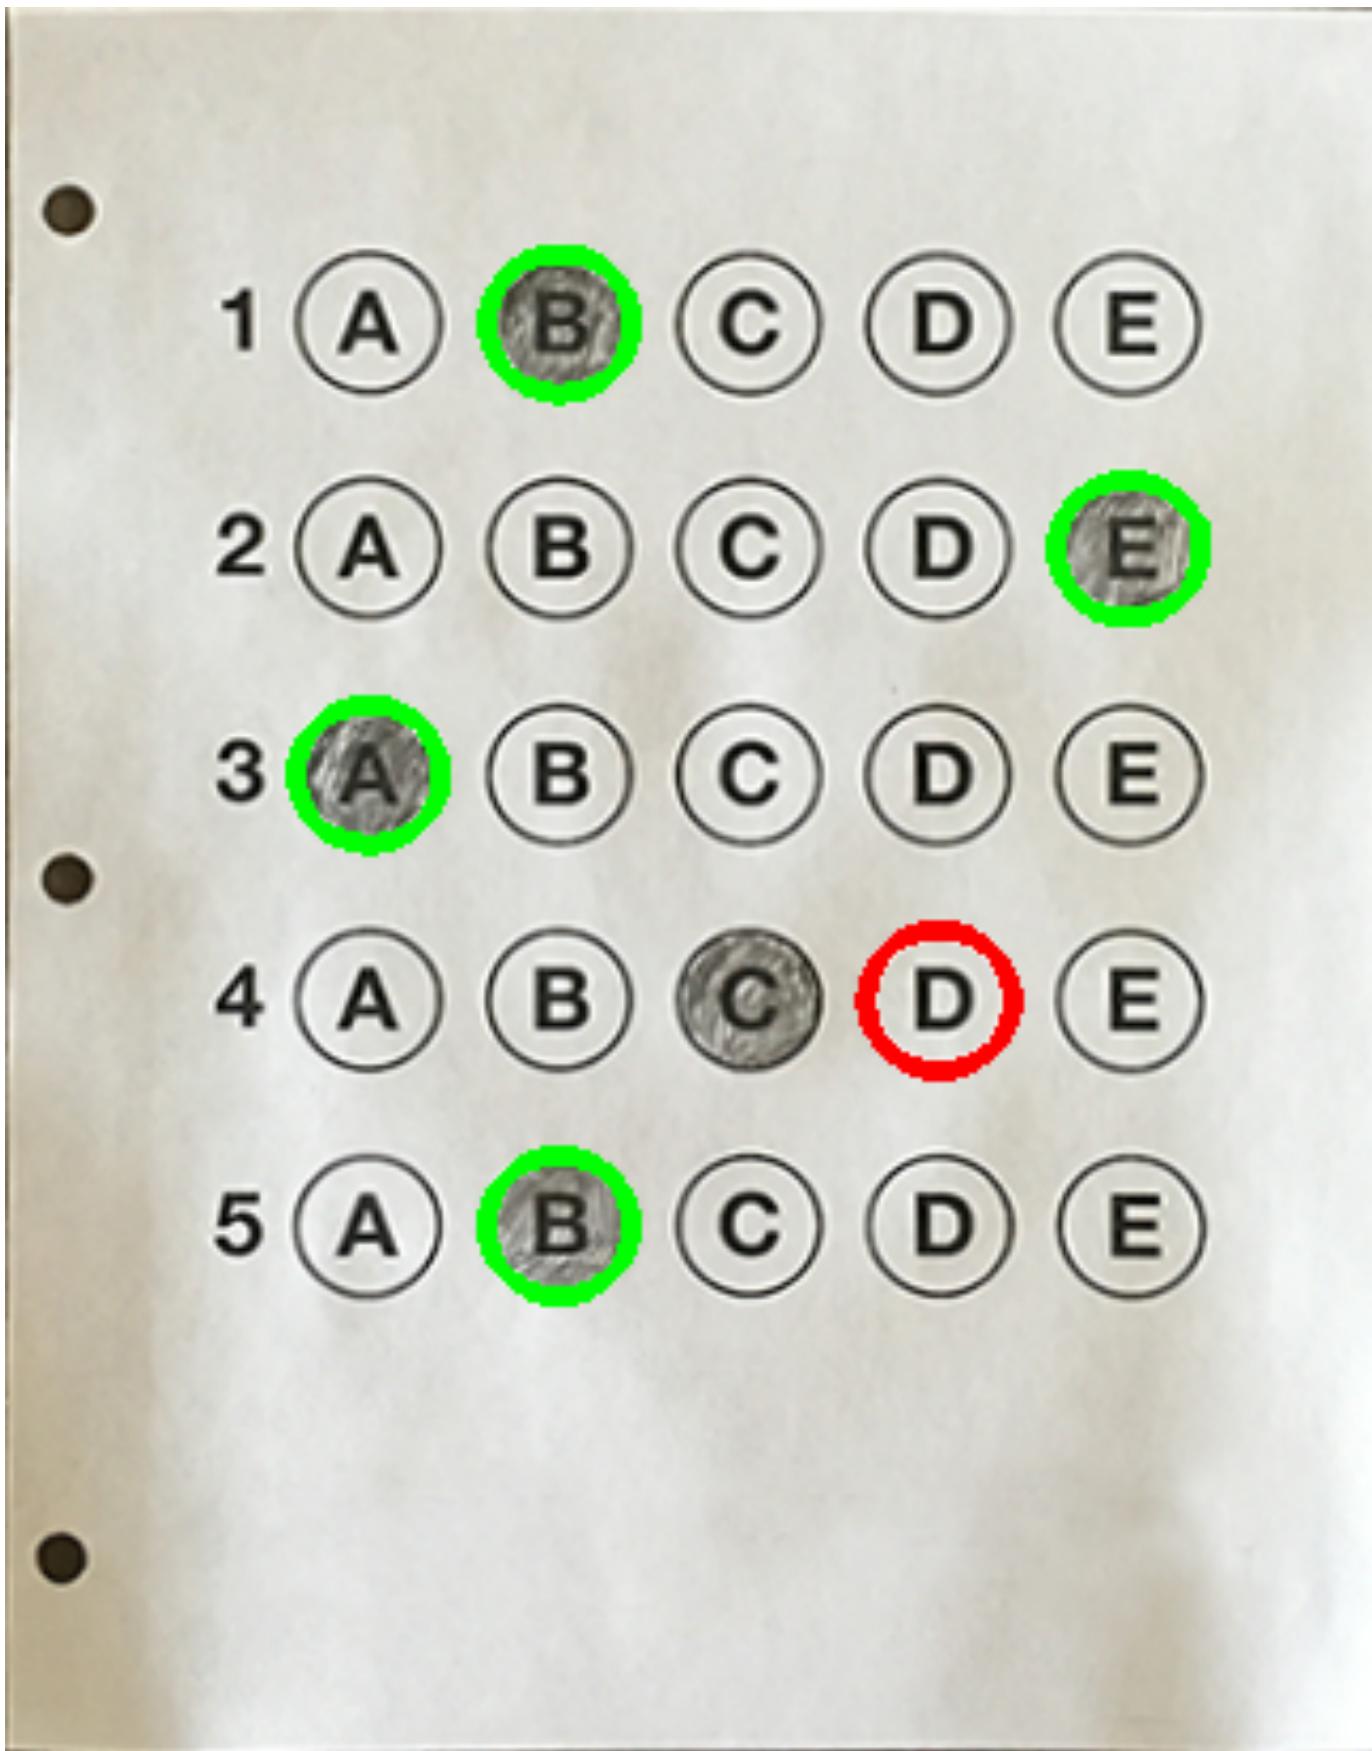
What's Filled In

A

```
97 # loop over the sorted contours
98 for (j, c) in enumerate(cnts):
99     # construct a mask that reveals only the current
100    # "bubble" for the question
101    mask = np.zeros(thresh.shape, dtype="uint8")
102    cv2.drawContours(mask, [c], -1, 255, -1)
103
104    # apply the mask to the thresholded image, then
105    # count the number of non-zero pixels in the
106    # bubble area
107    mask = cv2.bitwise_and(thresh, thresh, mask=mask)
108    total = cv2.countNonZero(mask)
109
110    # if the current total has a larger number of total
111    # non-zero pixels, then we are examining the currently
112    # bubbled-in answer
113    if bubbled is None or total > bubbled[0]:
114        bubbled = (total, j)
115
```

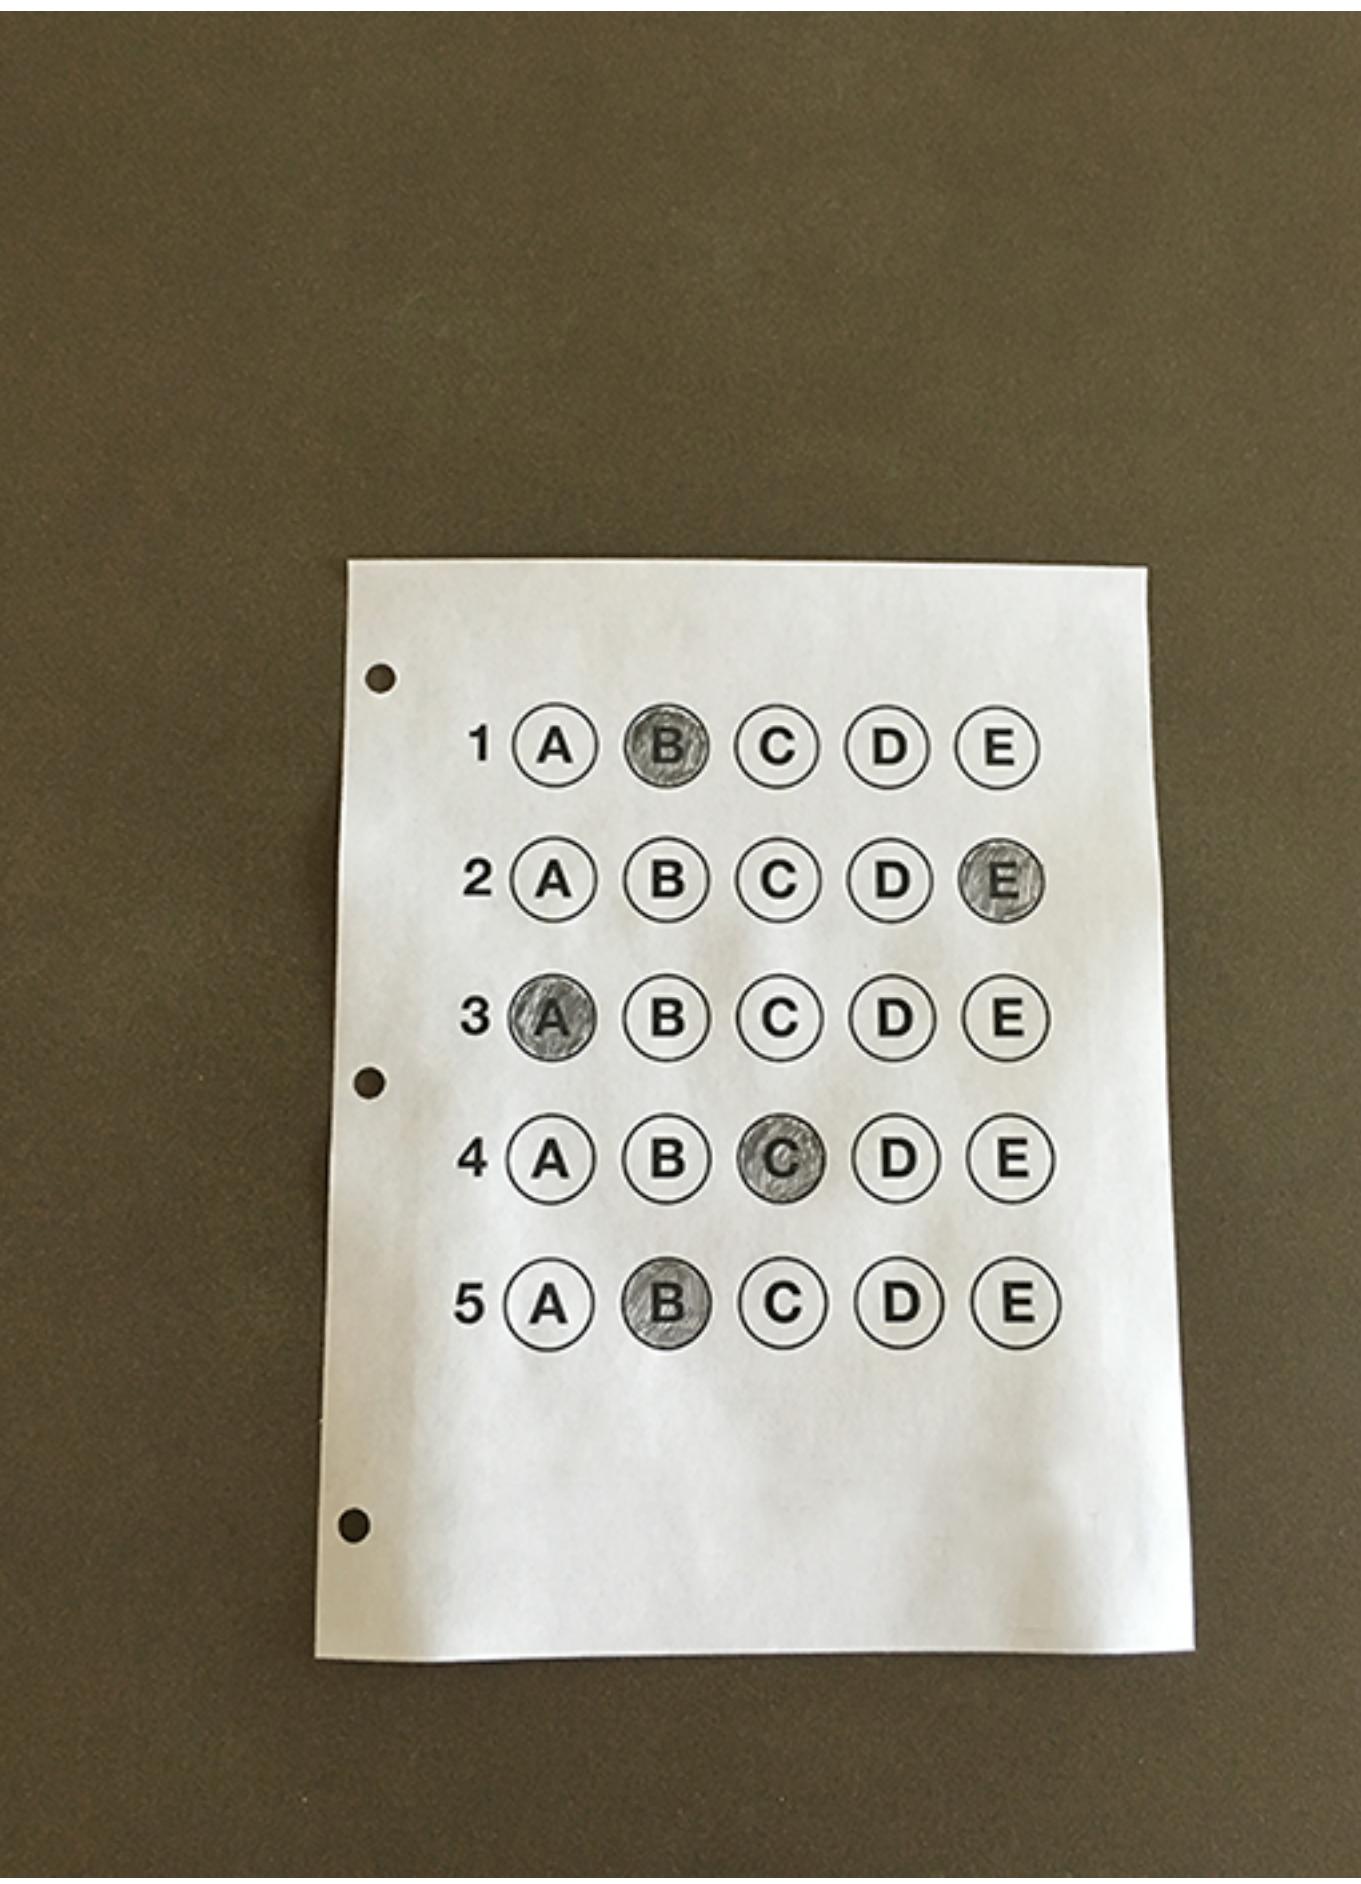
OMR

Answer Key

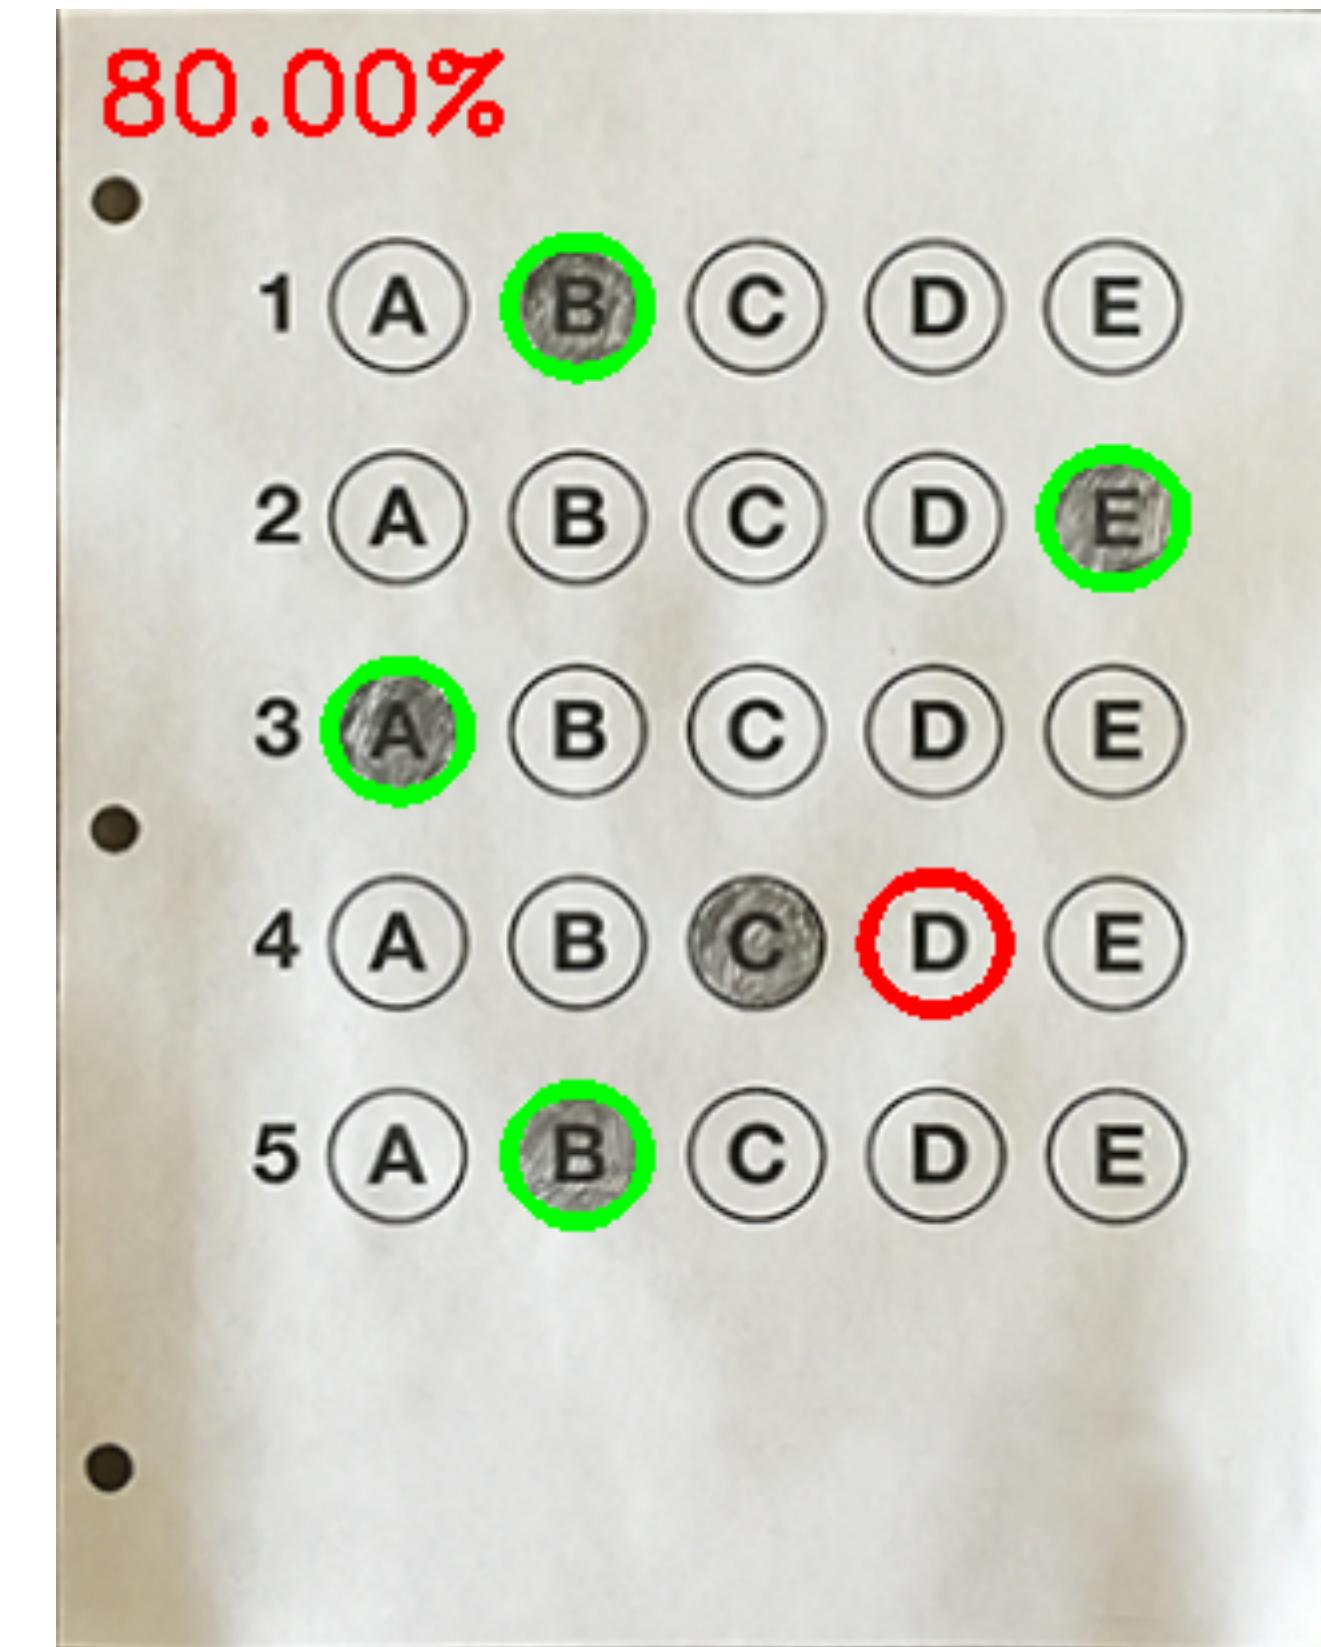


```
116 # initialize the contour color and the index of the
117 # *correct* answer
118 color = (0, 0, 255)
119 k = ANSWER_KEY[q]
120
121 # check to see if the bubbled answer is correct
122 if k == bubbled[1]:
123     color = (0, 255, 0)
124     correct += 1
125
126 # draw the outline of the correct answer on the test
127 cv2.drawContours(paper, [cnts[k]], -1, color, 3)
```

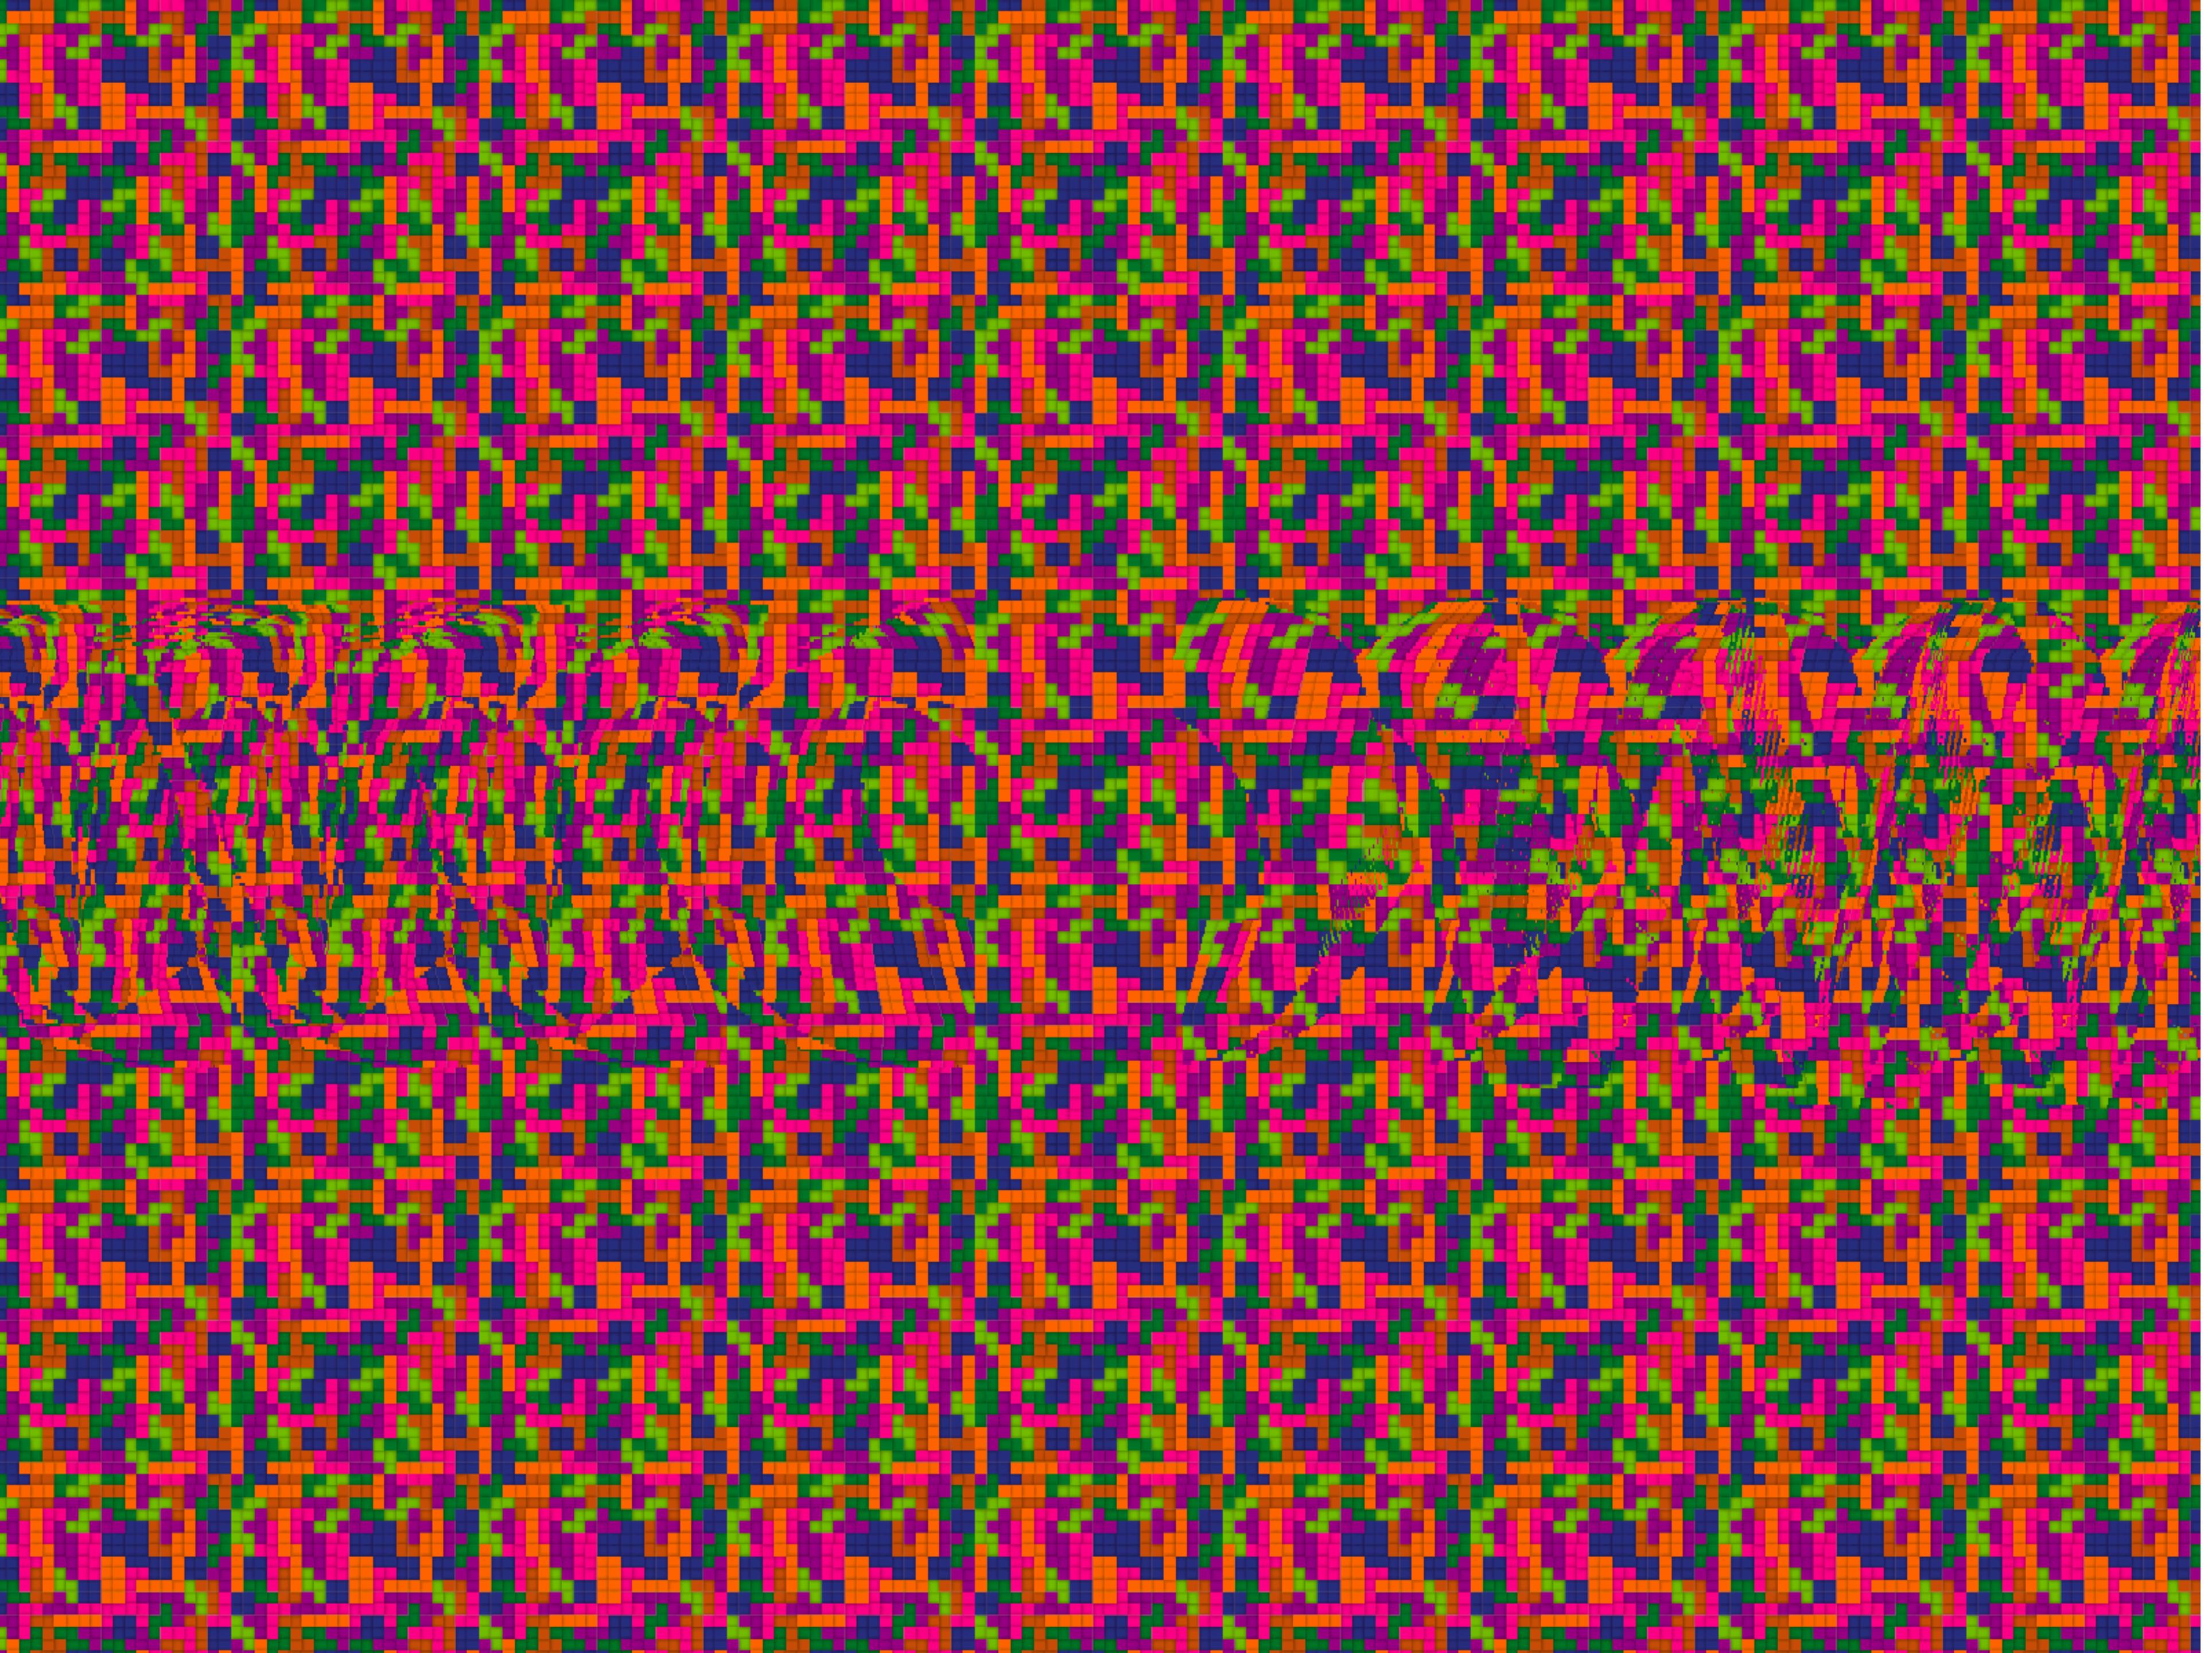
OMR Scoring



```
129 # grab the test taker
130 score = (correct / 5.0) * 100
131 print("[INFO] score: {:.2f}%".format(score))
132 cv2.putText(paper, "{}.2f%".format(score), (10, 30),
133 | cv2.FONT_HERSHEY_SIMPLEX, 0.9, (0, 0, 255), 2)
134 cv2.imshow("Original", image)
135 cv2.imshow("Exam", paper)
136 cv2.waitKey(0)
```



Autostereograms



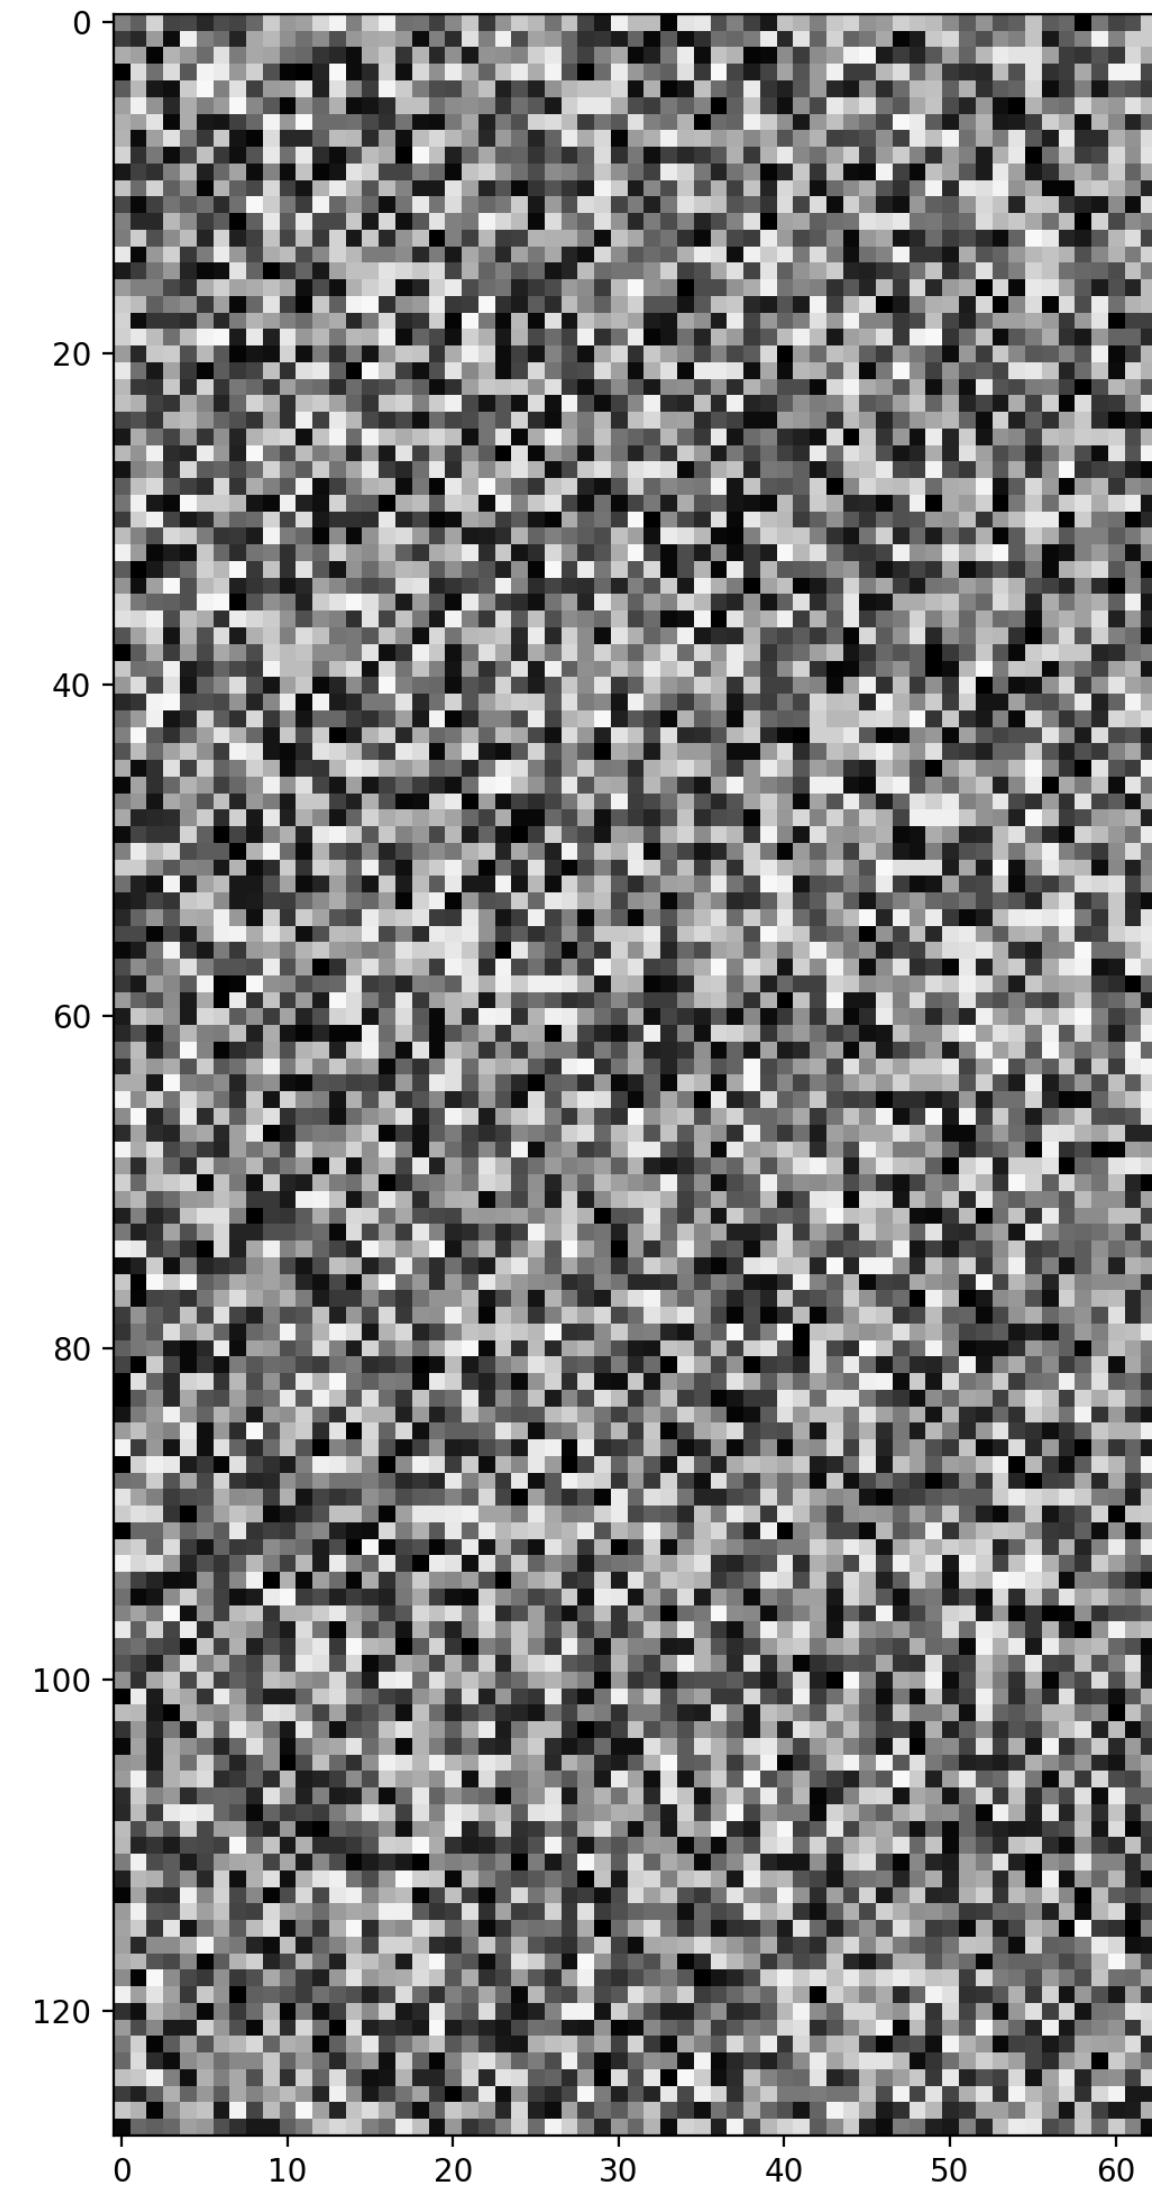
Depth Map Based Stereograms

1. Pattern

```
def make_pattern(shape=(16, 16), levels=64):
    "Creates a pattern from gray values."
    return np.random.randint(0, levels - 1, shape) / levels
```

```
def display(img, colorbar=False):
    "Displays an image."
    plt.figure(figsize=(10, 10))
    if len(img.shape) == 2:
        i = skimage.io.imshow(img, cmap='gray')
    else:
        i = skimage.io.imshow(img)
    if colorbar:
        plt.colorbar(i, shrink=0.5, label='depth')
    plt.tight_layout()
    plt.show()
```

```
58
59     pattern = make_pattern(shape=(128, 64))
60     display(pattern)
```

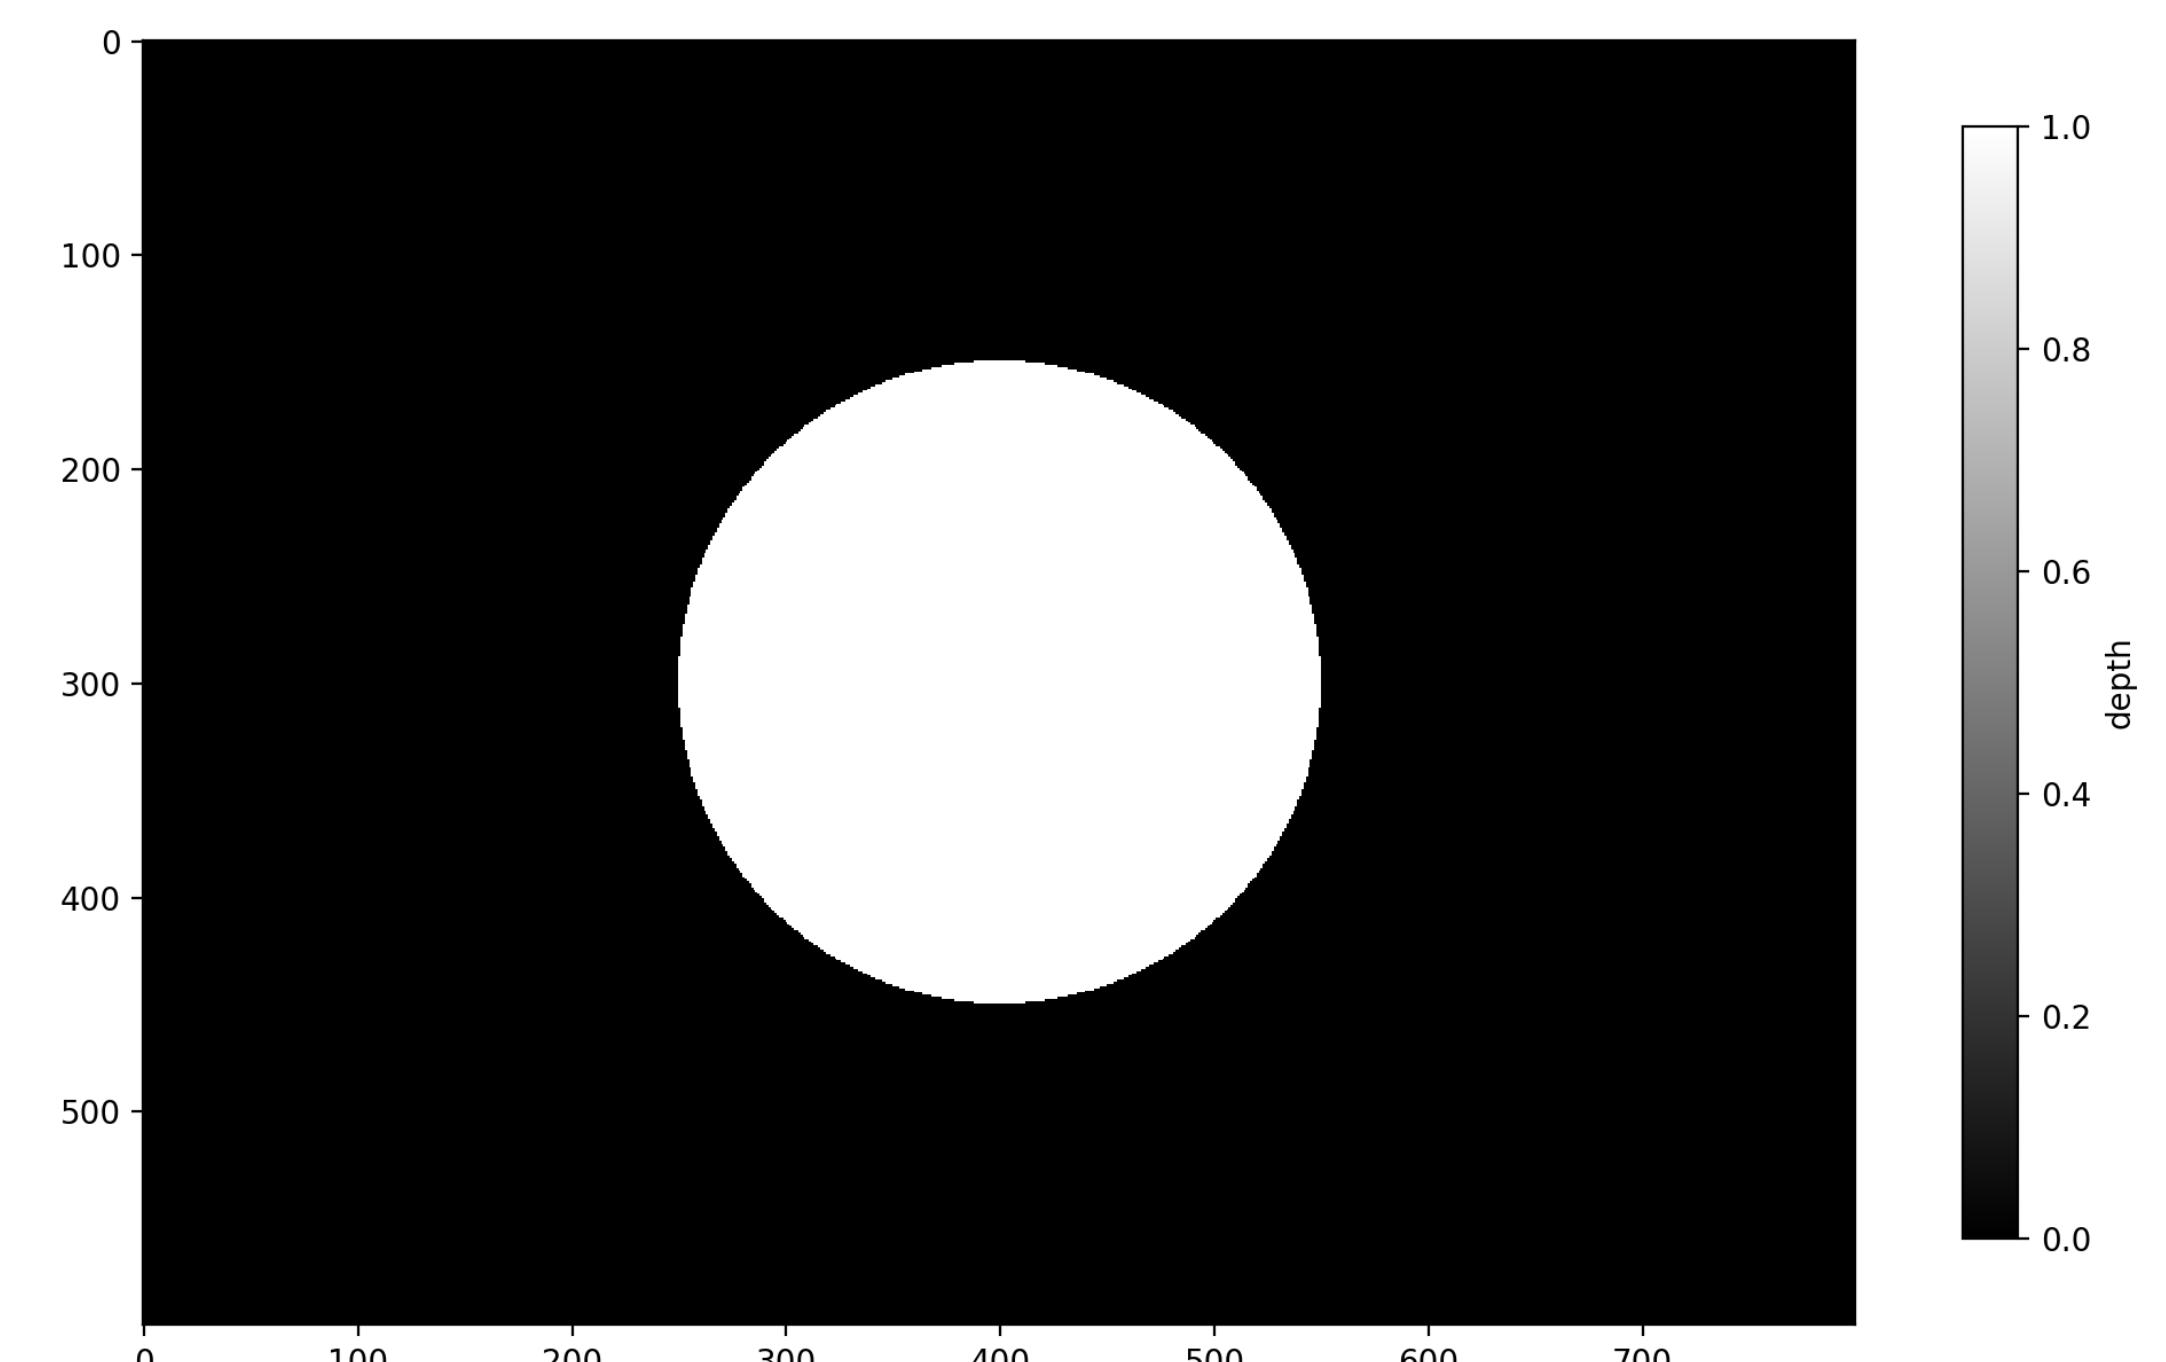


Depth Map Based Stereograms

2. Depth Map

```
def create_circular_depthmap(shape=(600, 800), center=None, radius=100):
    "Creates a circular depthmap, centered on the image."
    depthmap = np.zeros(shape, dtype=float)
    r = np.arange(depthmap.shape[0])
    c = np.arange(depthmap.shape[1])
    R, C = np.meshgrid(r, c, indexing='ij')
    if center is None:
        center = np.array([r.max() / 2, c.max() / 2])
    d = np.sqrt((R - center[0])**2 + (C - center[1])**2)
    depthmap += (d < radius)
    return depthmap
```

```
depthmap = create_circular_depthmap(radius=150)
display(depthmap, colorbar=True)
```



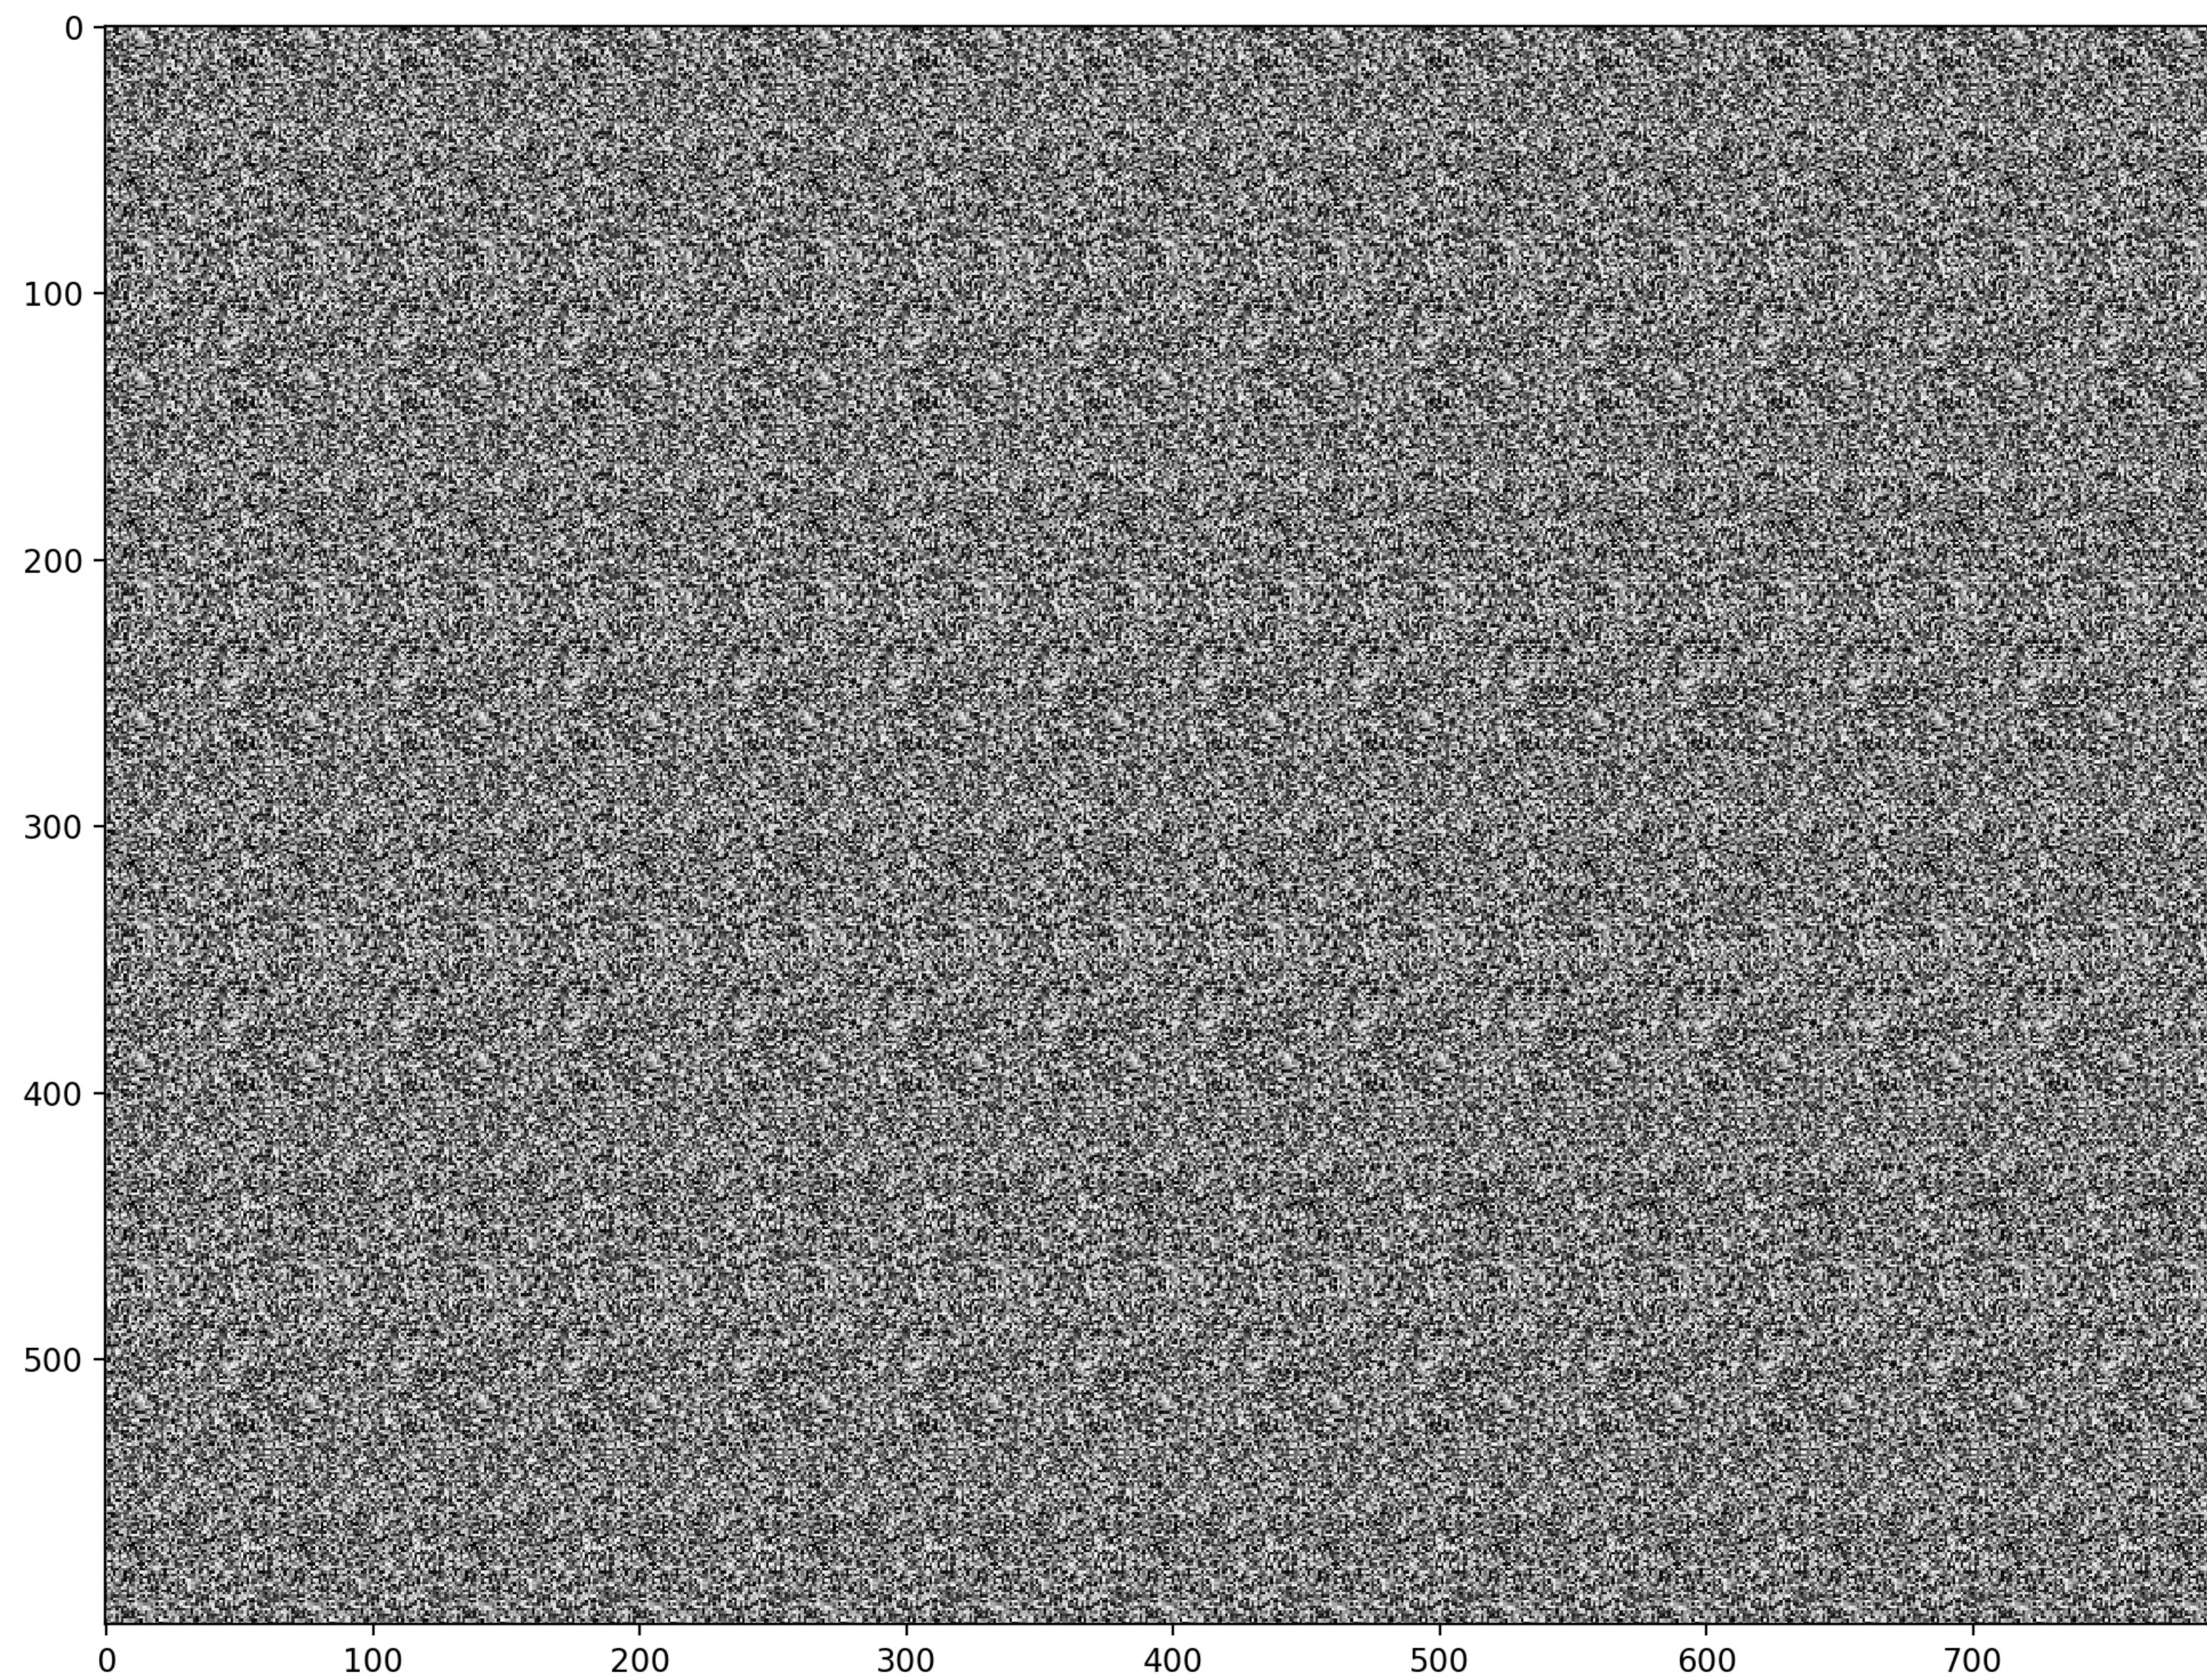
Depth Map Based Stereograms

3. Autostereogram

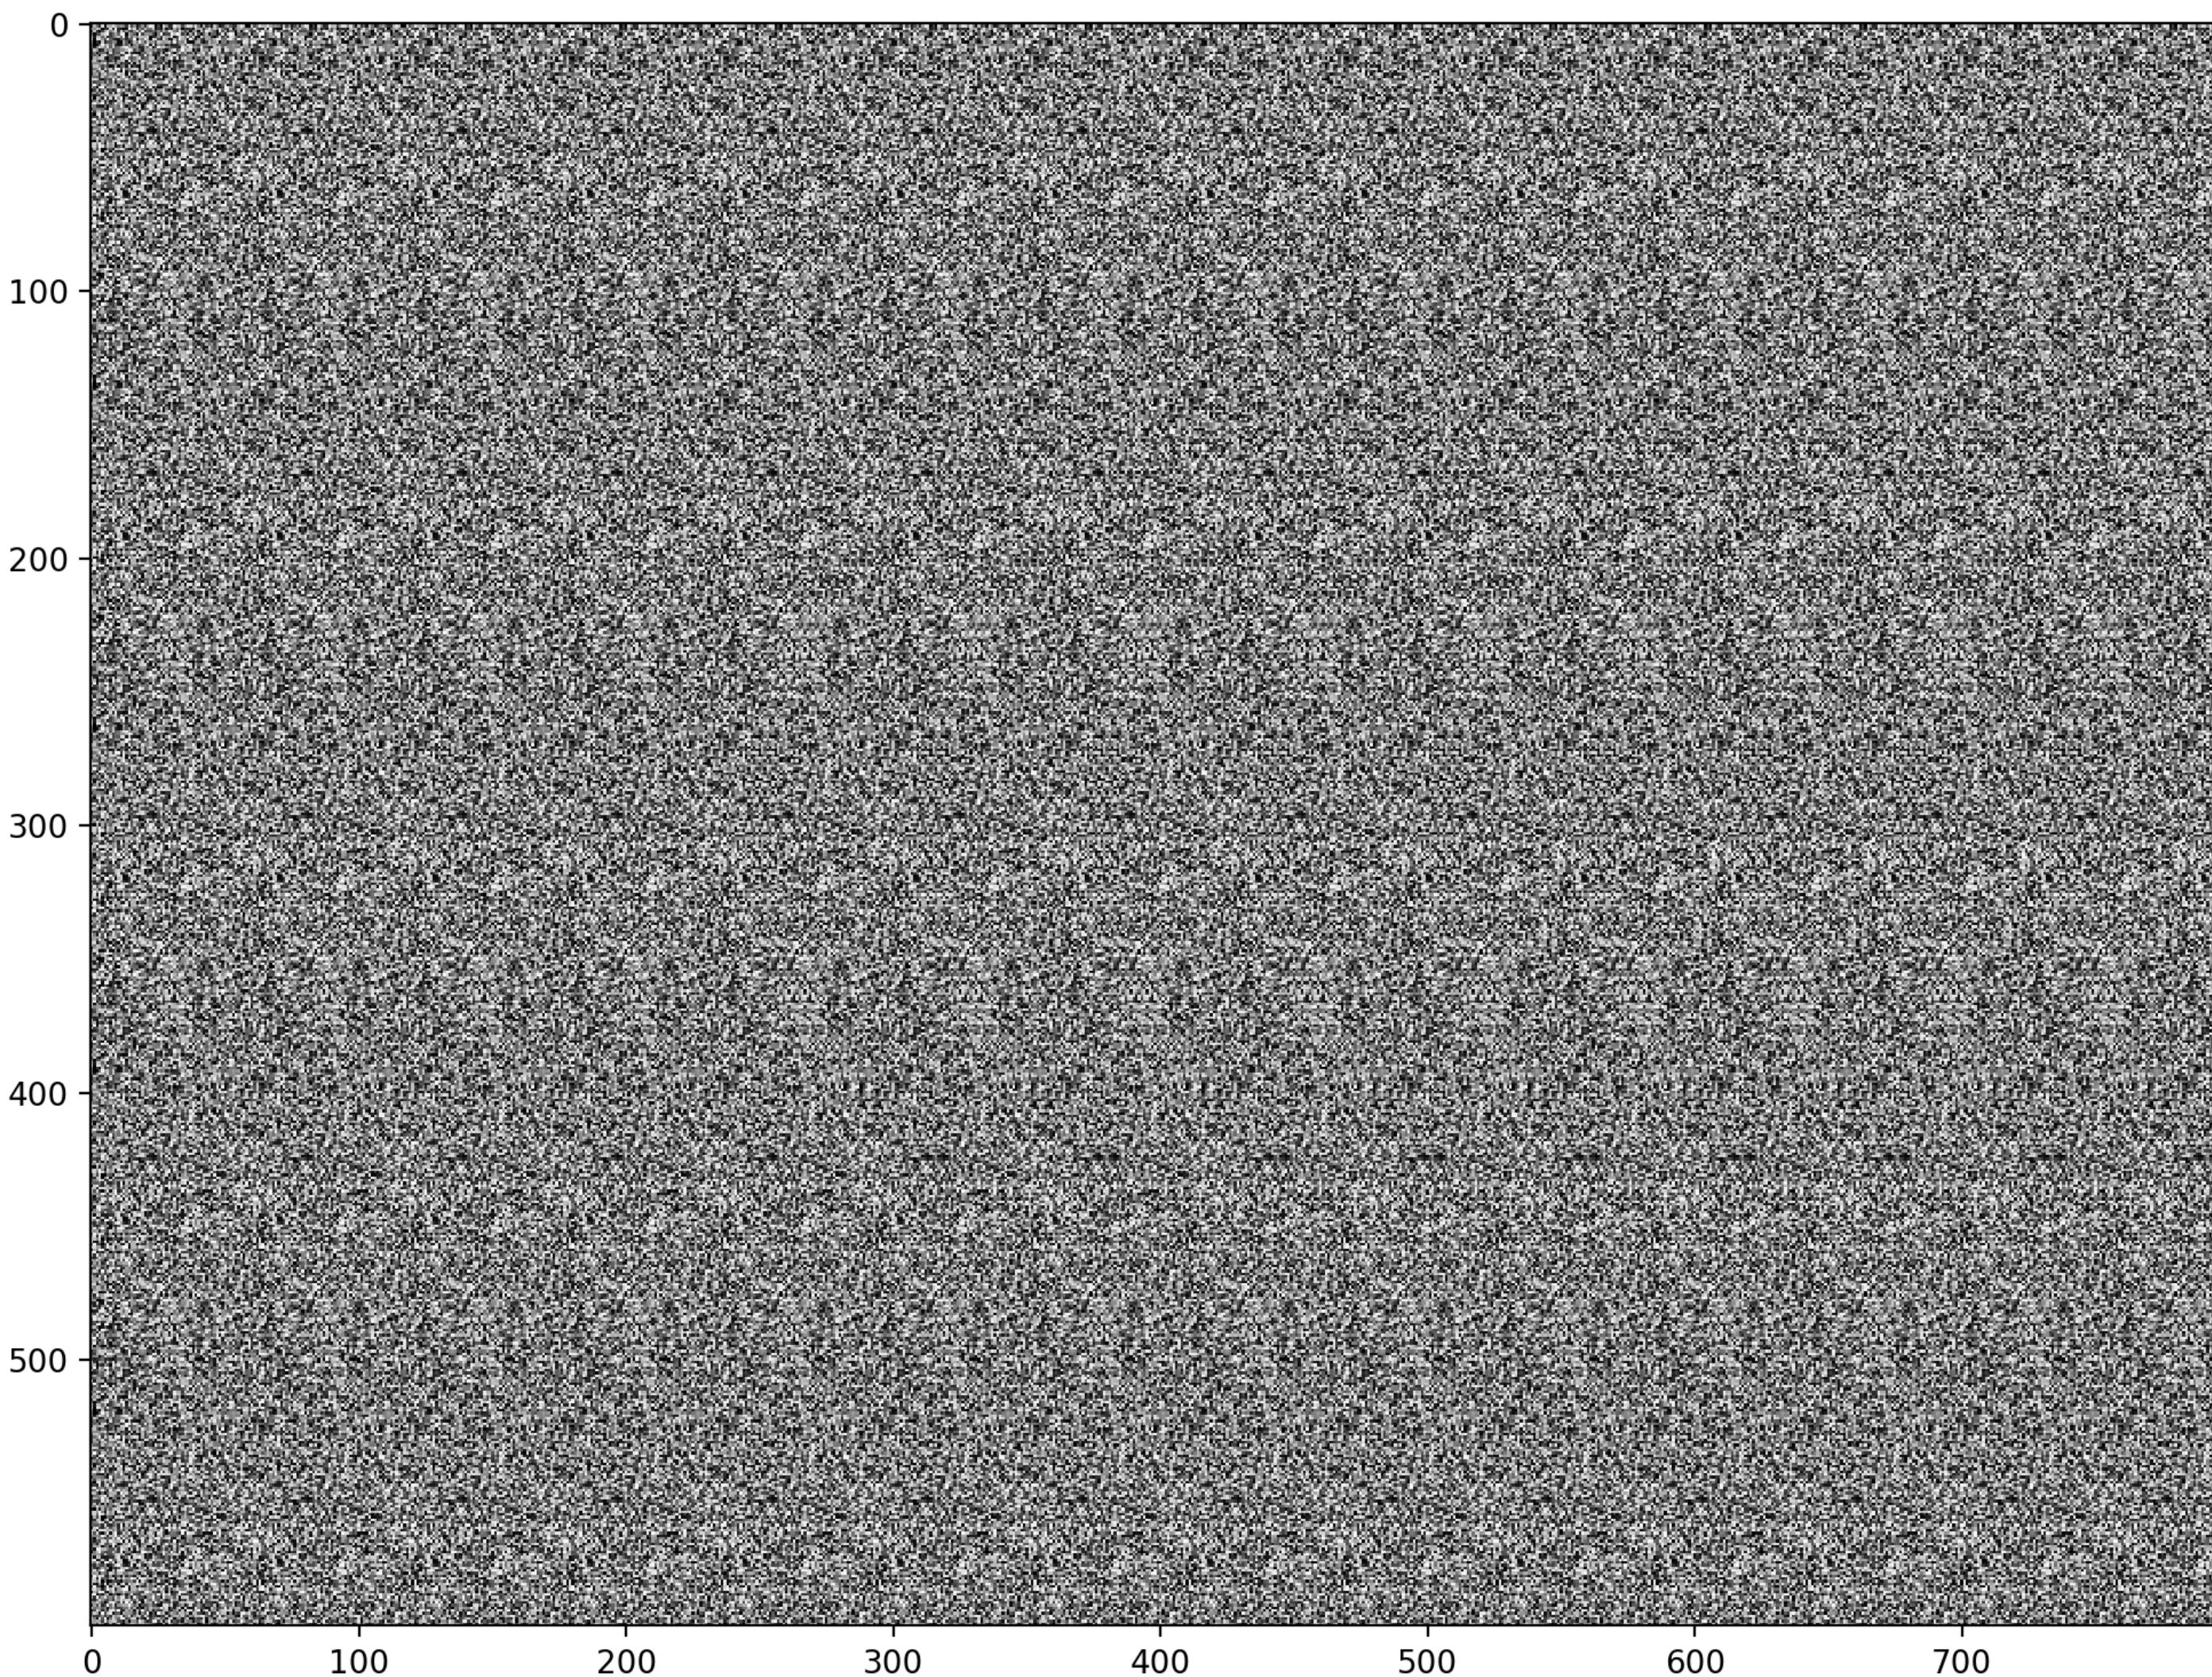
```
def normalize(depthmap):
    "Normalizes values of depthmap to [0, 1] range."
    if depthmap.max() > depthmap.min():
        return (depthmap - depthmap.min()) / (depthmap.max() - depthmap.min())
    else:
        return depthmap

def make_autostereogram(depthmap, pattern, shift_amplitude=0.1, invert=False):
    "Creates an autostereogram from depthmap and pattern."
    depthmap = normalize(depthmap)
    if invert:
        depthmap = 1 - depthmap
    autostereogram = np.zeros_like(depthmap, dtype=pattern.dtype)
    for r in np.arange(autostereogram.shape[0]):
        for c in np.arange(autostereogram.shape[1]):
            if c < pattern.shape[1]:
                autostereogram[r, c] = pattern[r % pattern.shape[0], c]
            else:
                shift = int(depthmap[r, c] * shift_amplitude * pattern.shape[1])
                autostereogram[r, c] = autostereogram[r, c - pattern.shape[1] + shift]
    return autostereogram
```

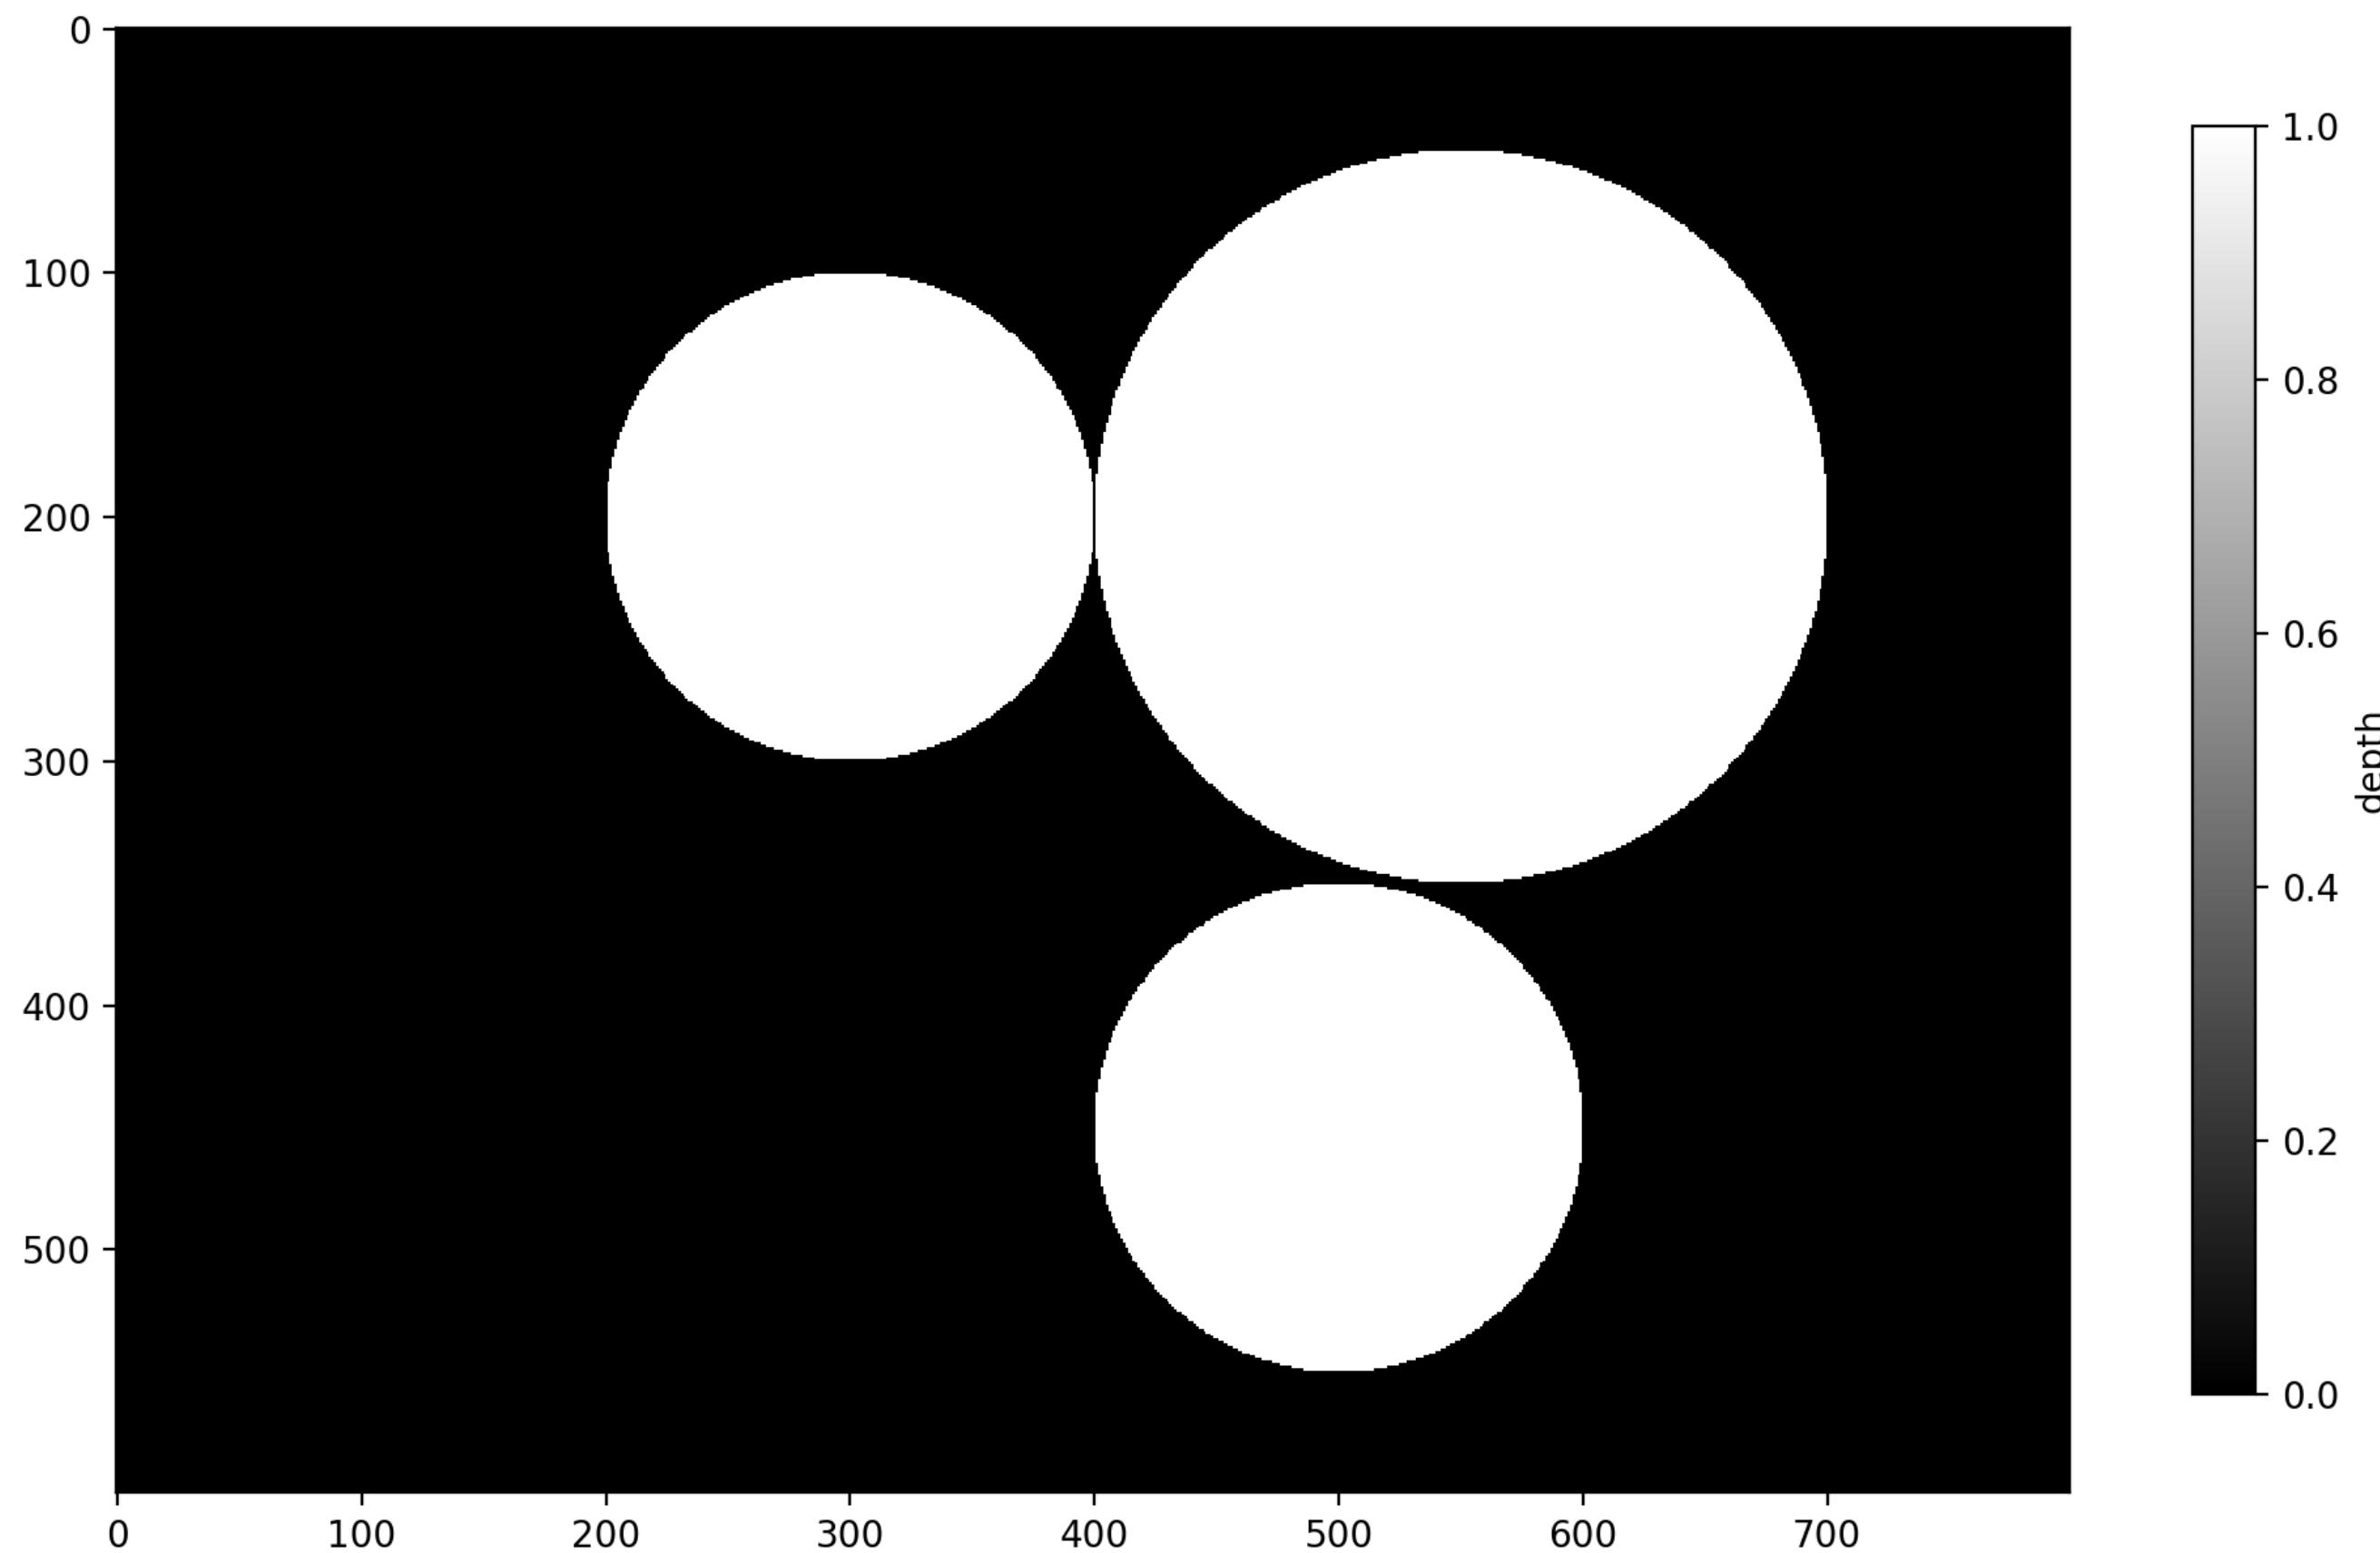
```
autostereogram = make_autostereogram(depthmap, pattern)
display(autostereogram)
```



```
autostereogram = make_autostereogram(depthmap, pattern, invert=True)
display(autostereogram)
```



```
depthmap = create_circular_depthmap(center=(200, 300), radius=100) + \
           create_circular_depthmap(center=(450, 500), radius=100) + \
           create_circular_depthmap(center=(200, 550), radius=150)
depthmap = normalize(depthmap)
display(depthmap, colorbar=True)
autostereogram = make_autostereogram(depthmap, pattern)
display(autostereogram)
```



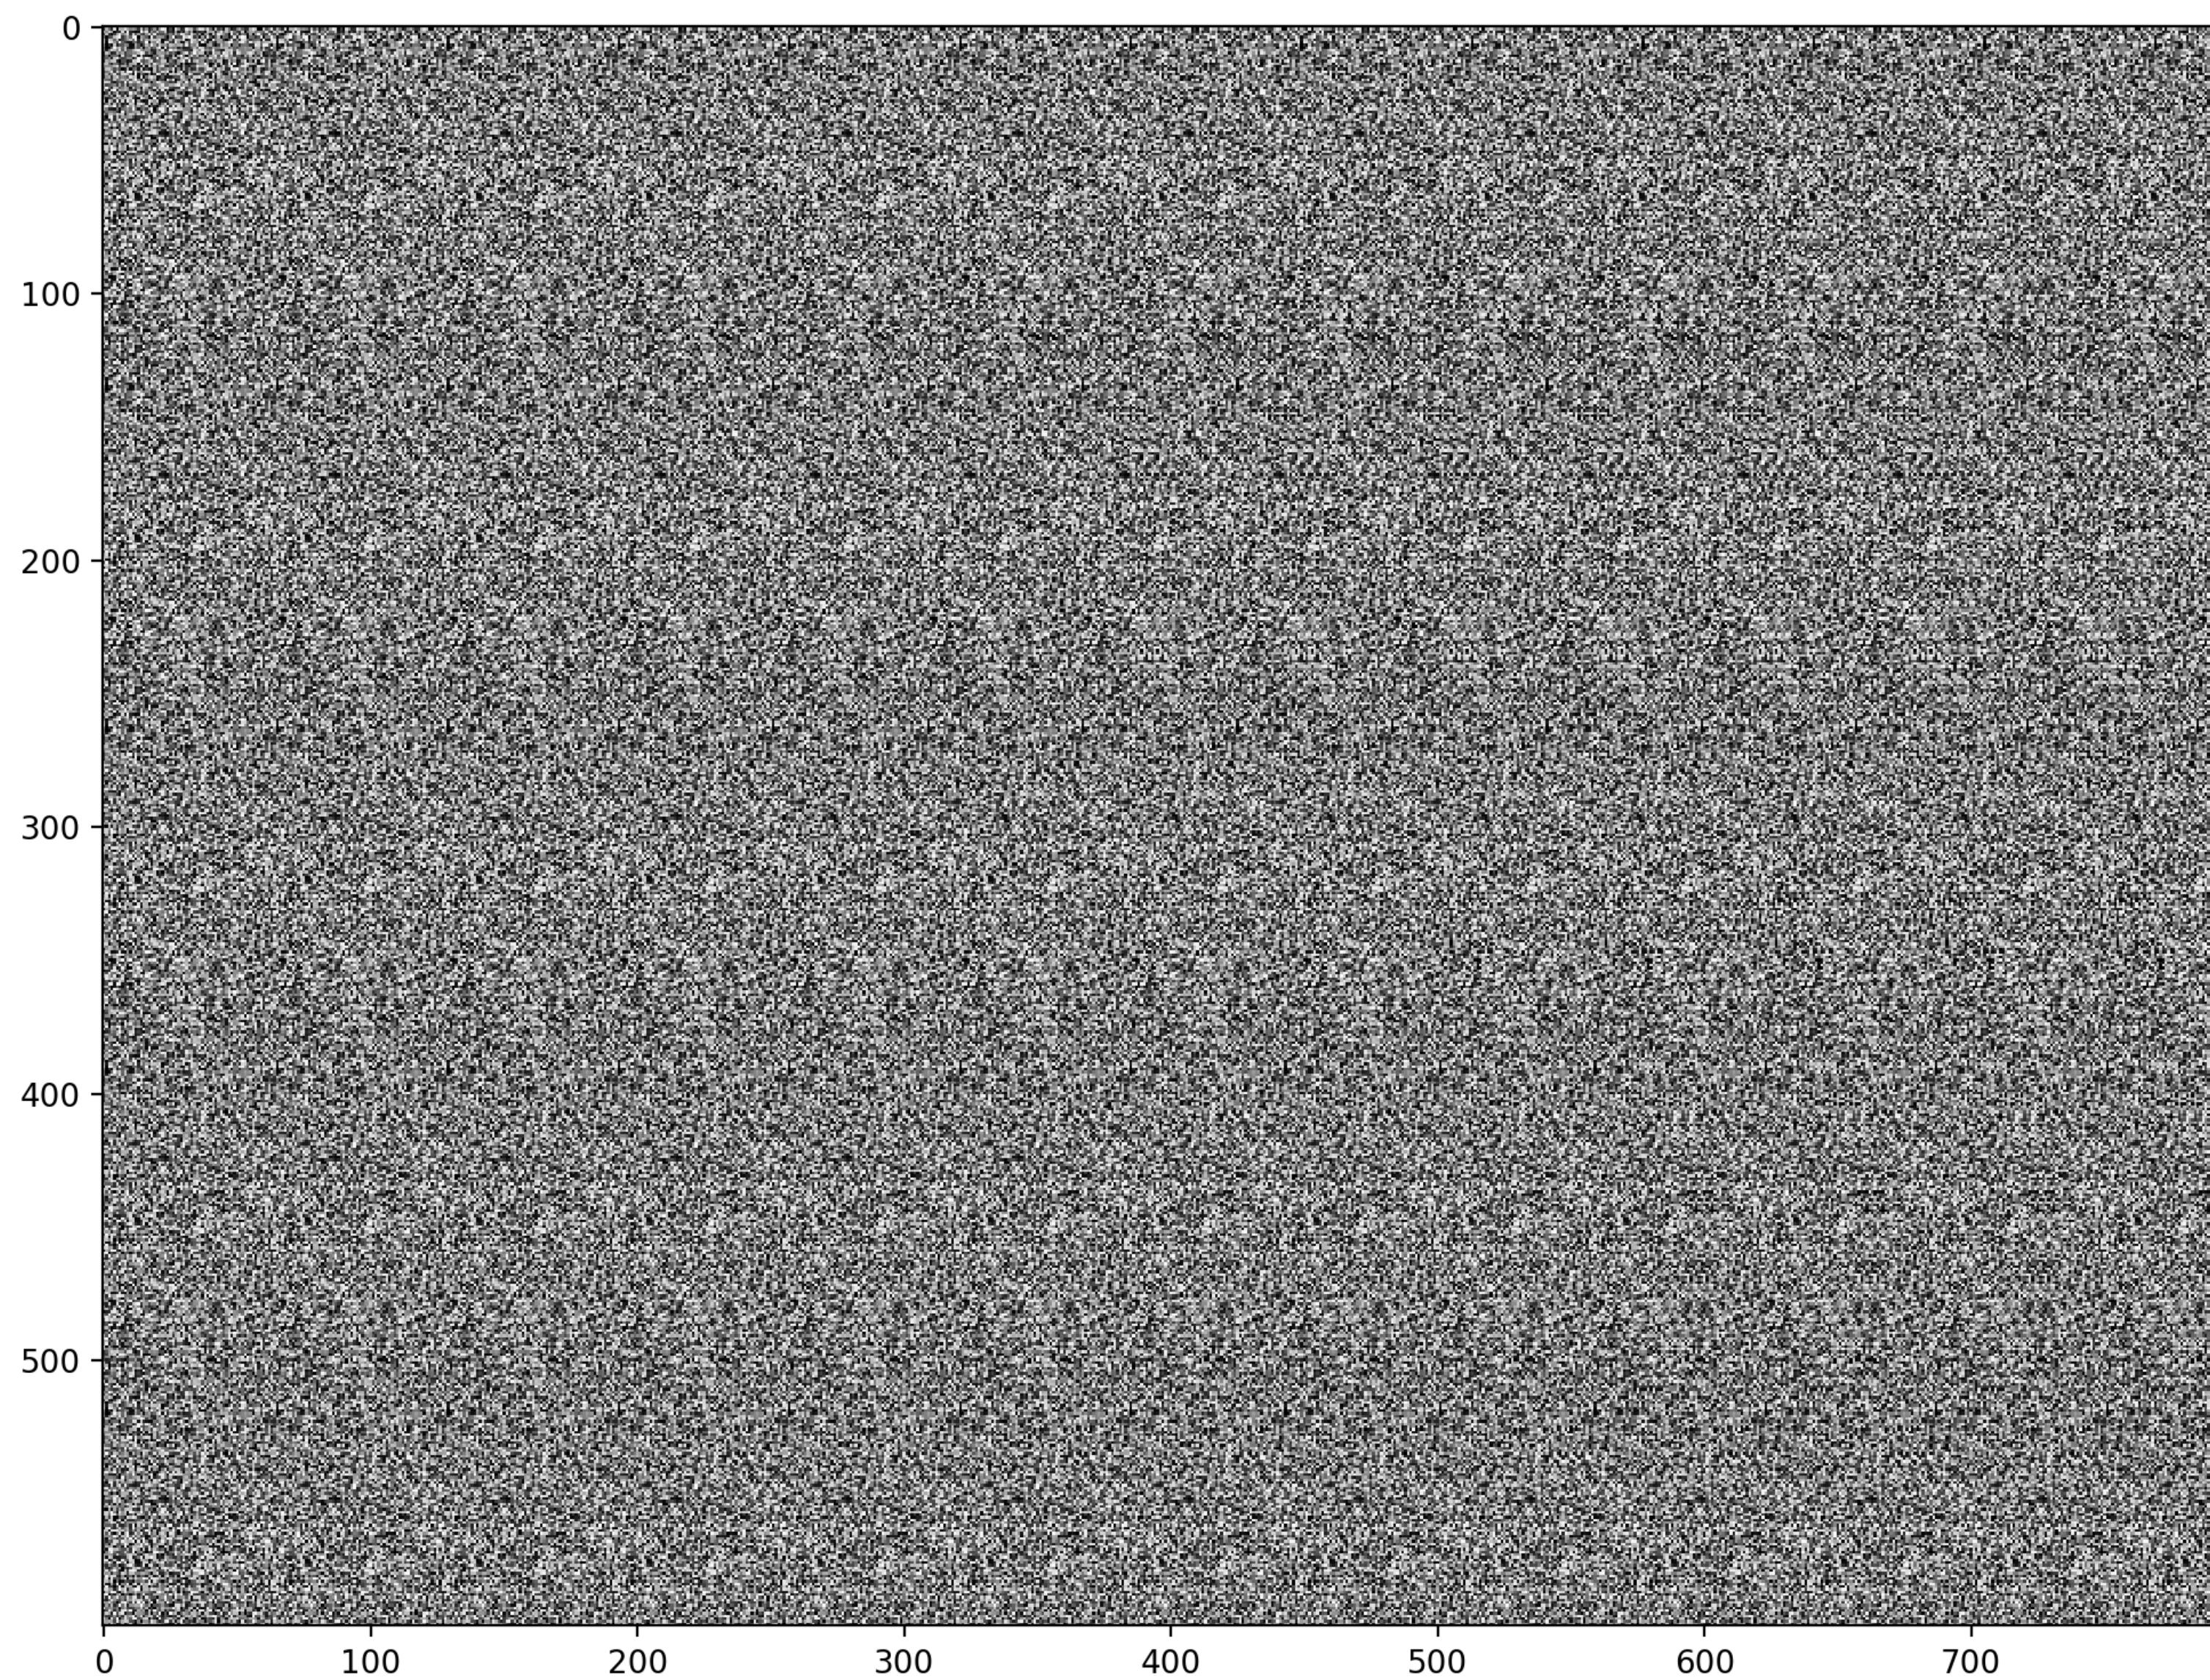
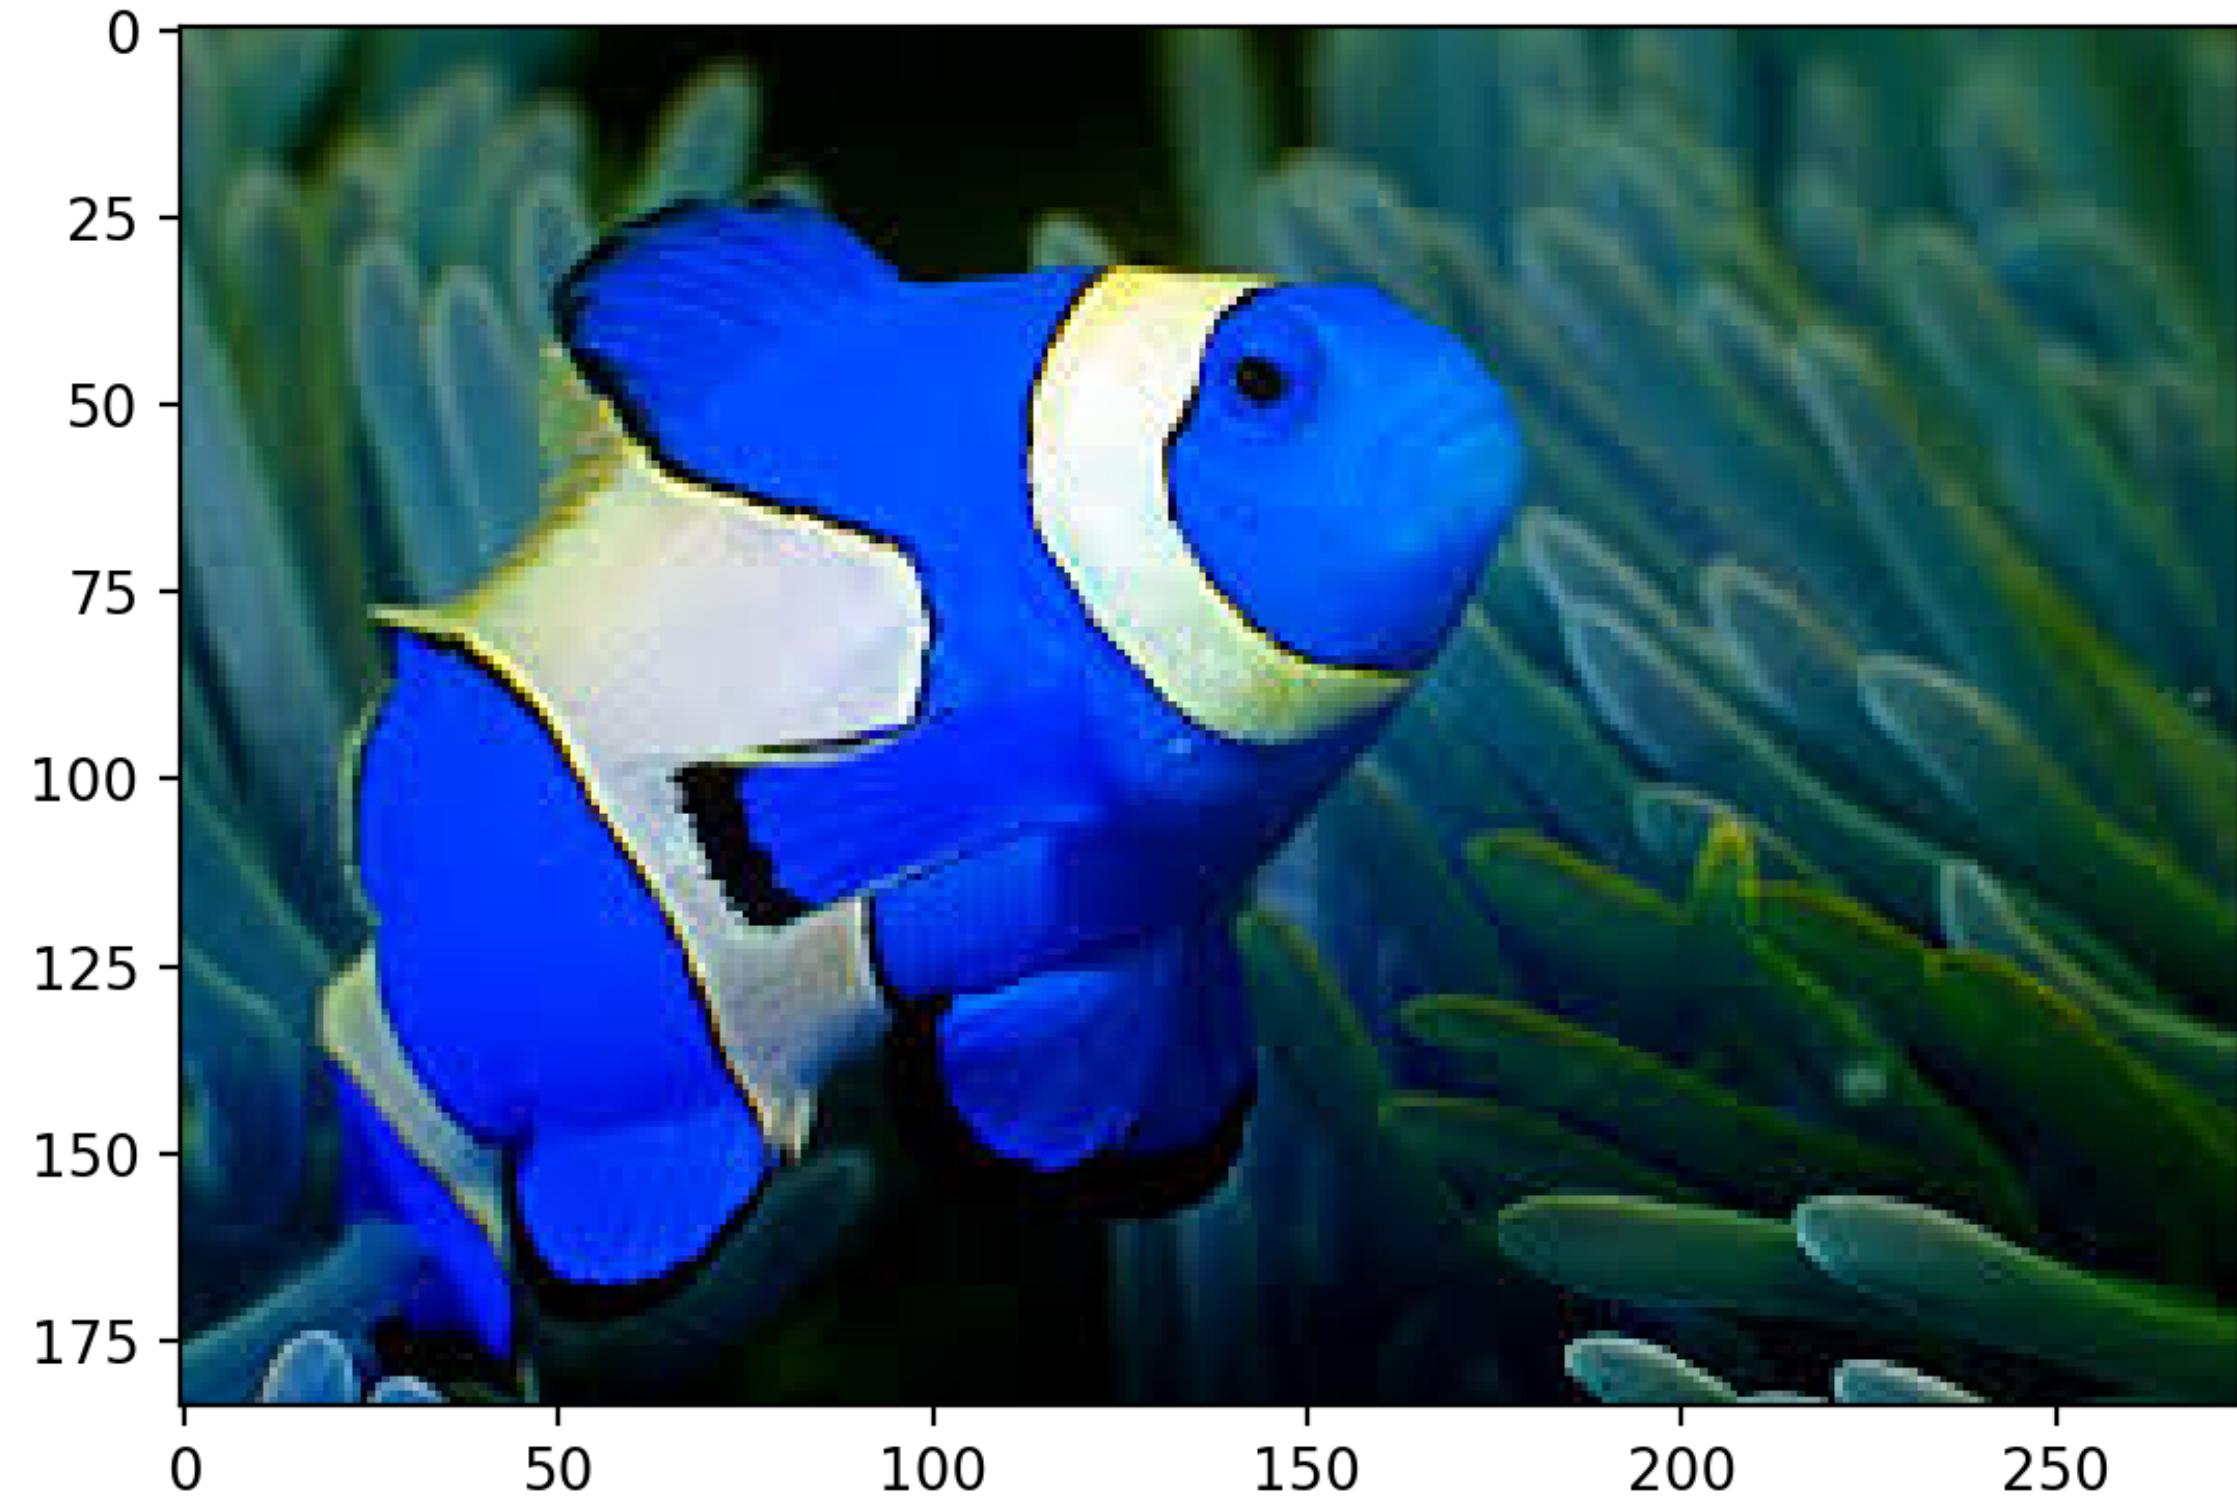


Image Segmentation

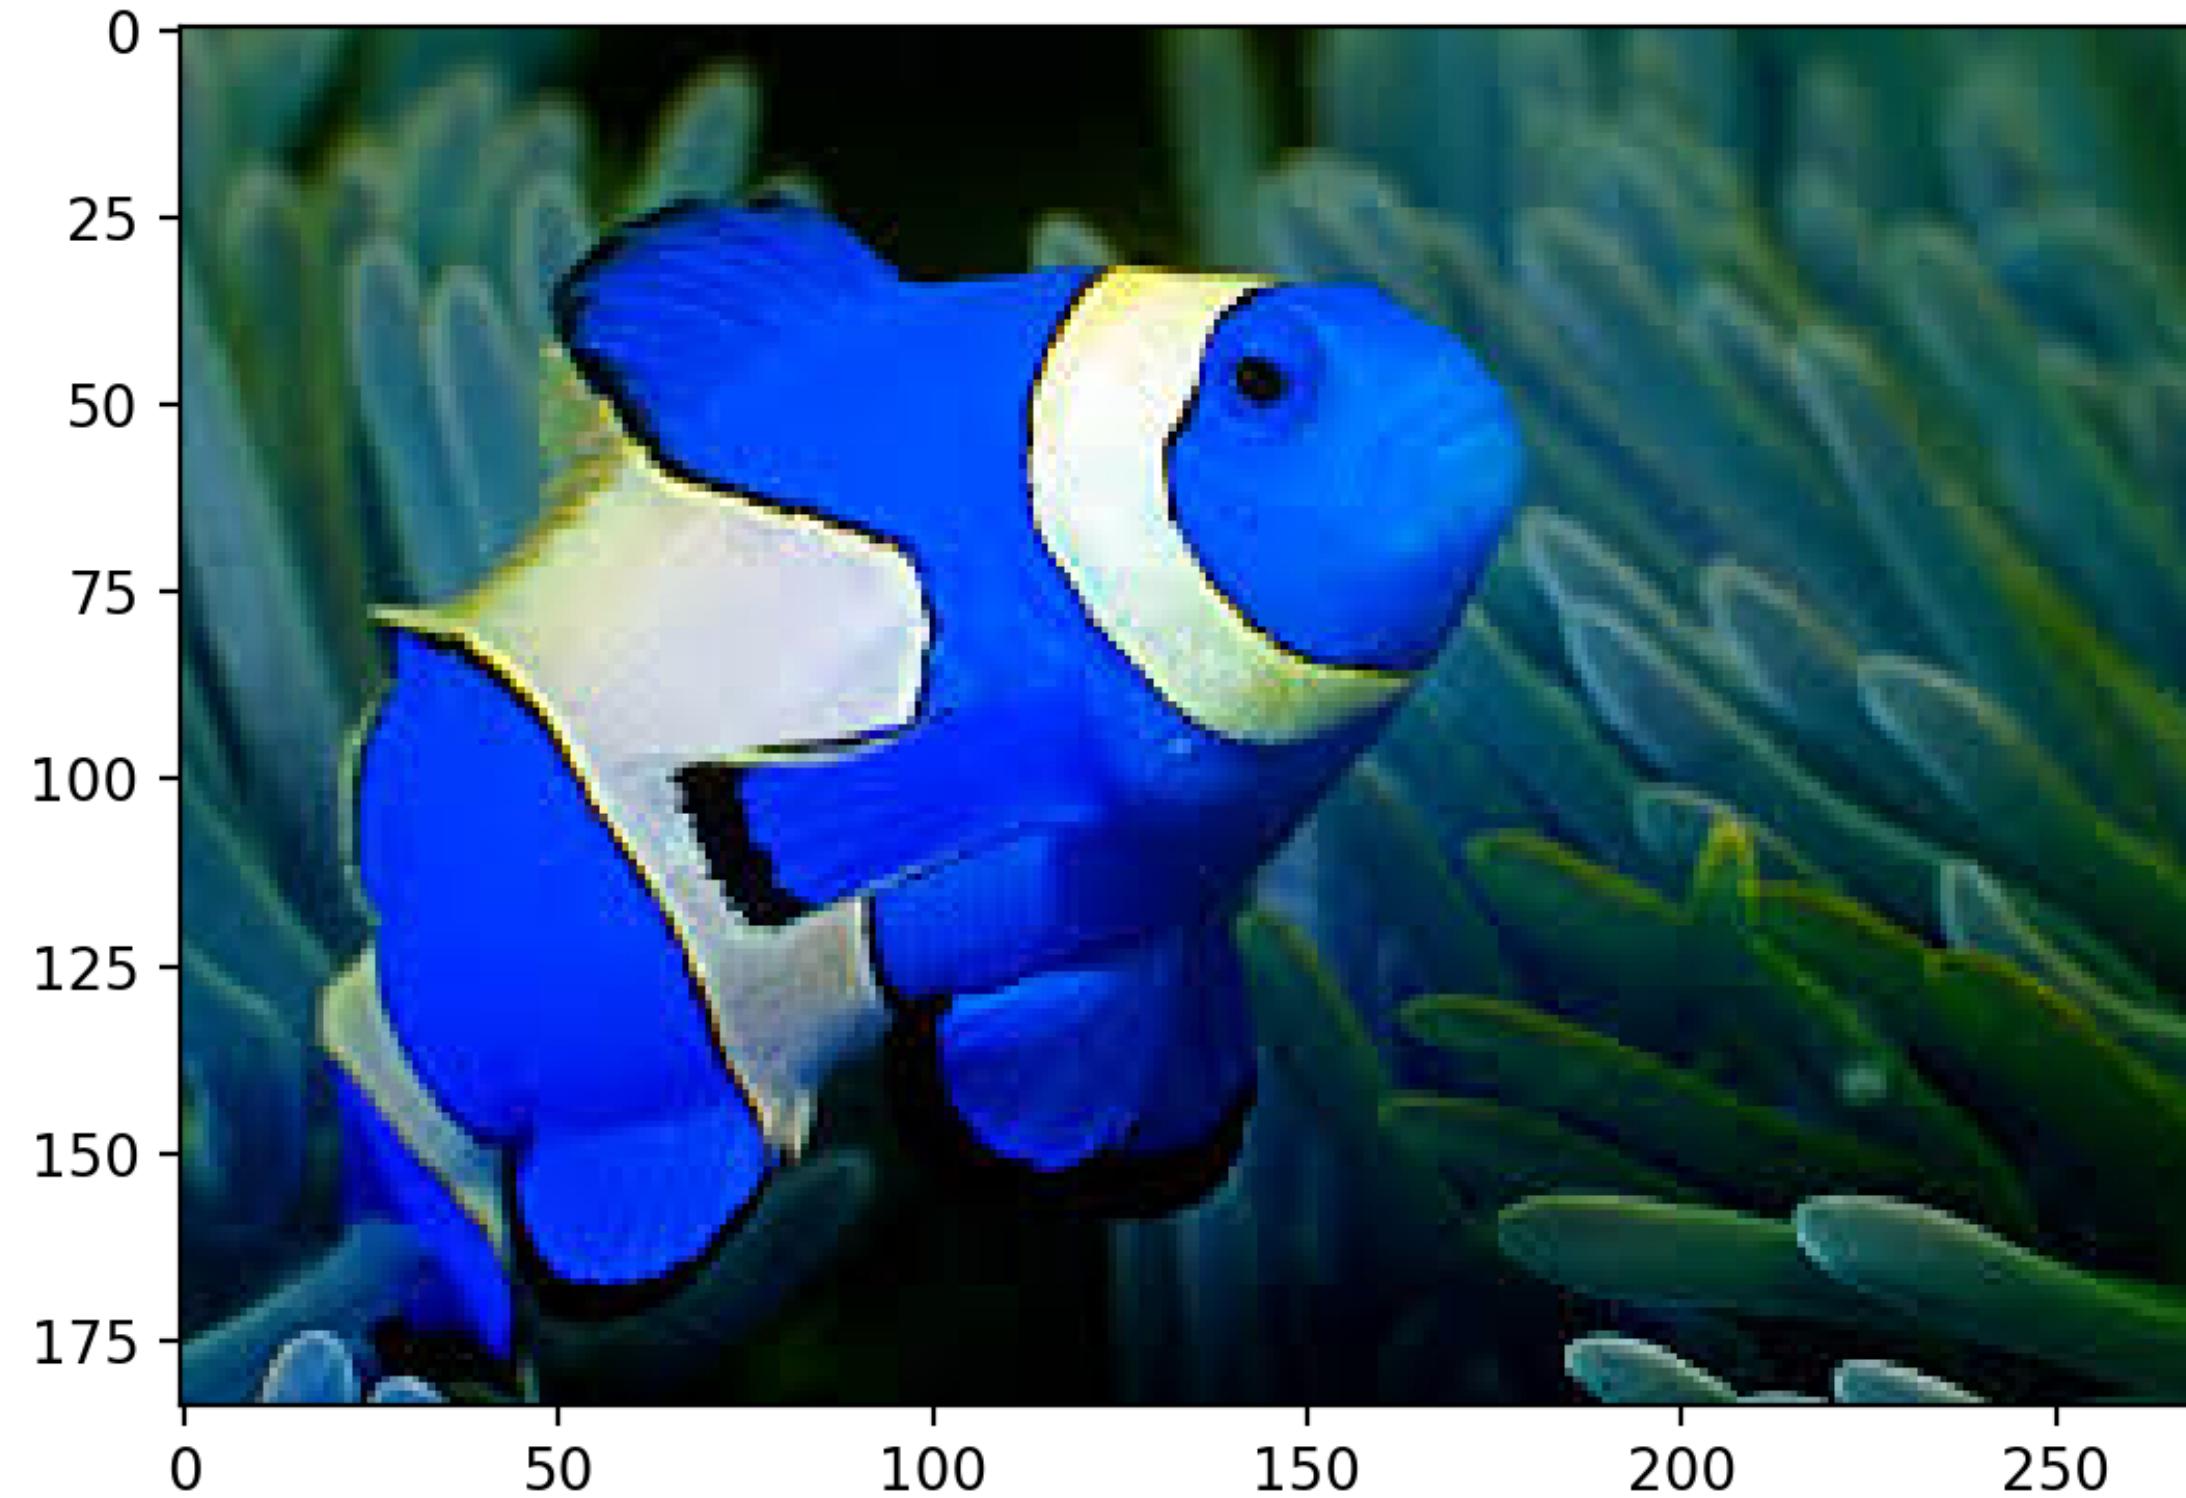
Image Segmentation



```
nemo = cv2.imread("nemo0.jpg")
plt.imshow(nemo)
plt.show()
```



```
nemo = cv2.cvtColor(nemo, cv2.COLOR_BGR2RGB)  
  
plt.imshow(nemo)  
plt.show()
```

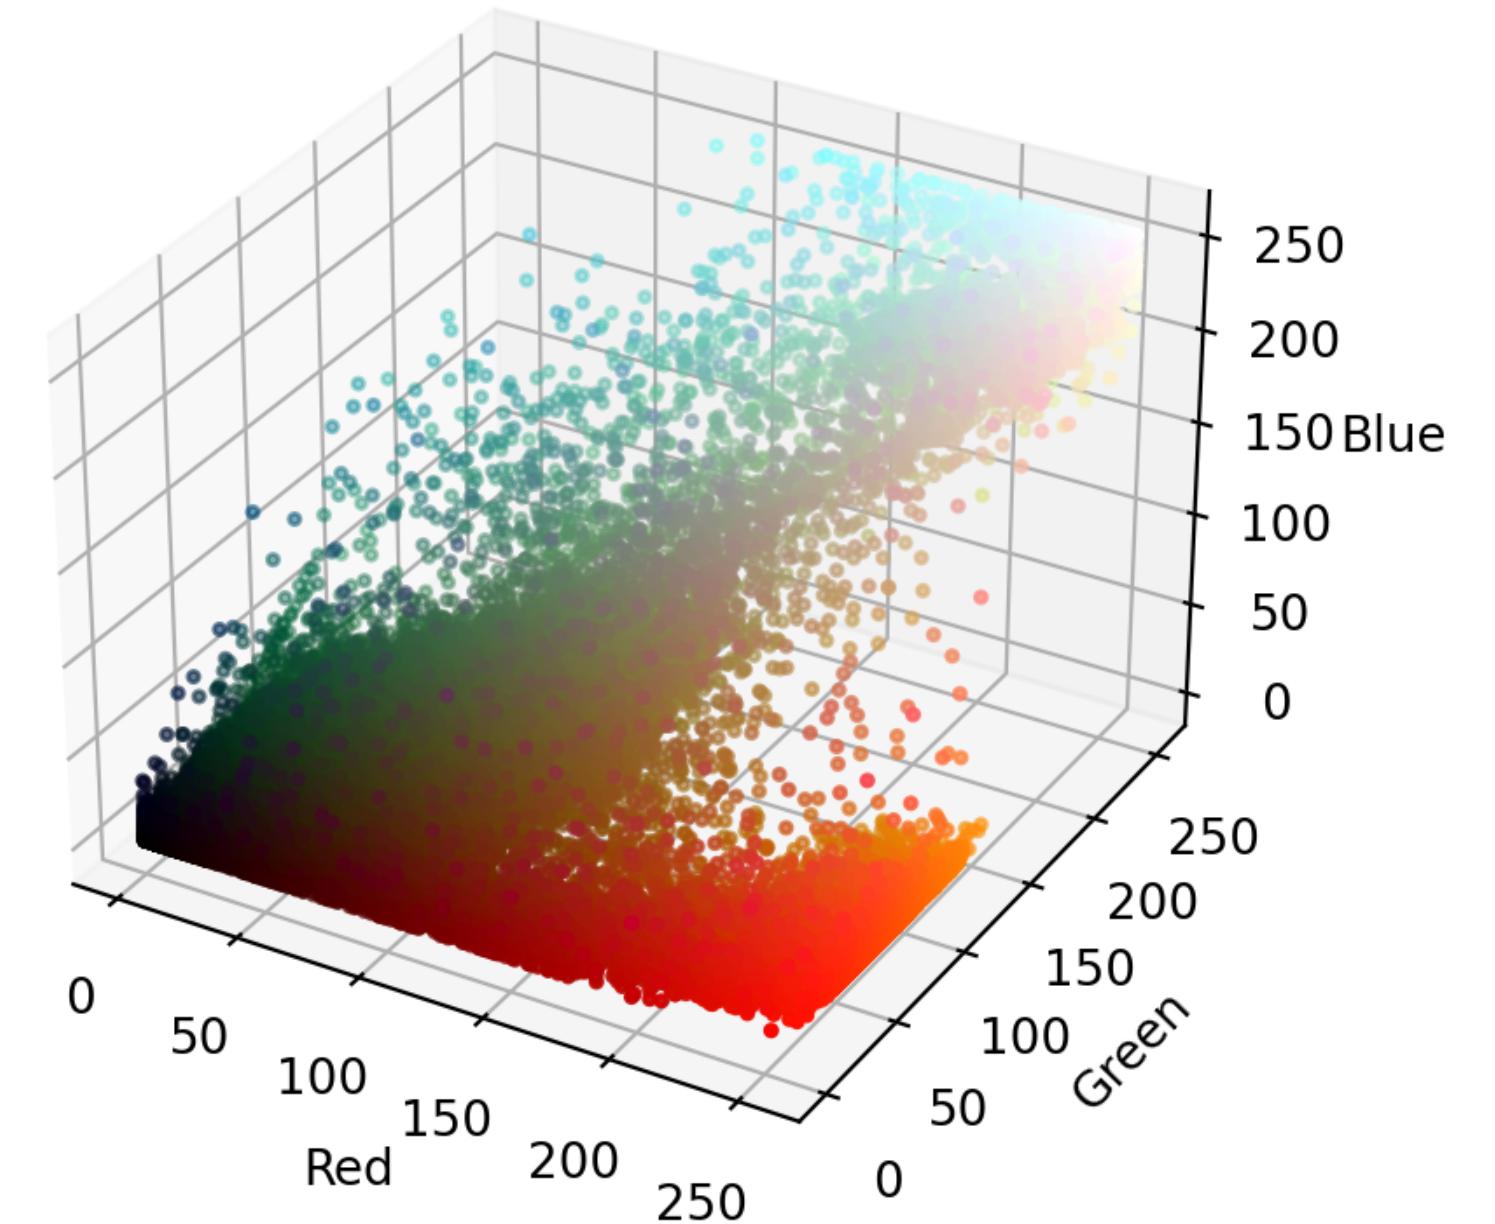


Visualizing in RGB Color Space

```
r, g, b = cv2.split(nemo)

fig = plt.figure()
axis = fig.add_subplot(1, 1, 1, projection="3d")
pixel_colors = nemo.reshape((np.shape(nemo)[0] * np.shape(nemo)[1], 3))
norm = colors.Normalize(vmin=-1.0, vmax=1.0)
norm.autoscale(pixel_colors)
pixel_colors = norm(pixel_colors).tolist()

axis.scatter(
    r.flatten(), g.flatten(), b.flatten(), facecolors=pixel_colors, marker="."
)
axis.set_xlabel("Red")
axis.set_ylabel("Green")
axis.set_zlabel("Blue")
plt.show()
```

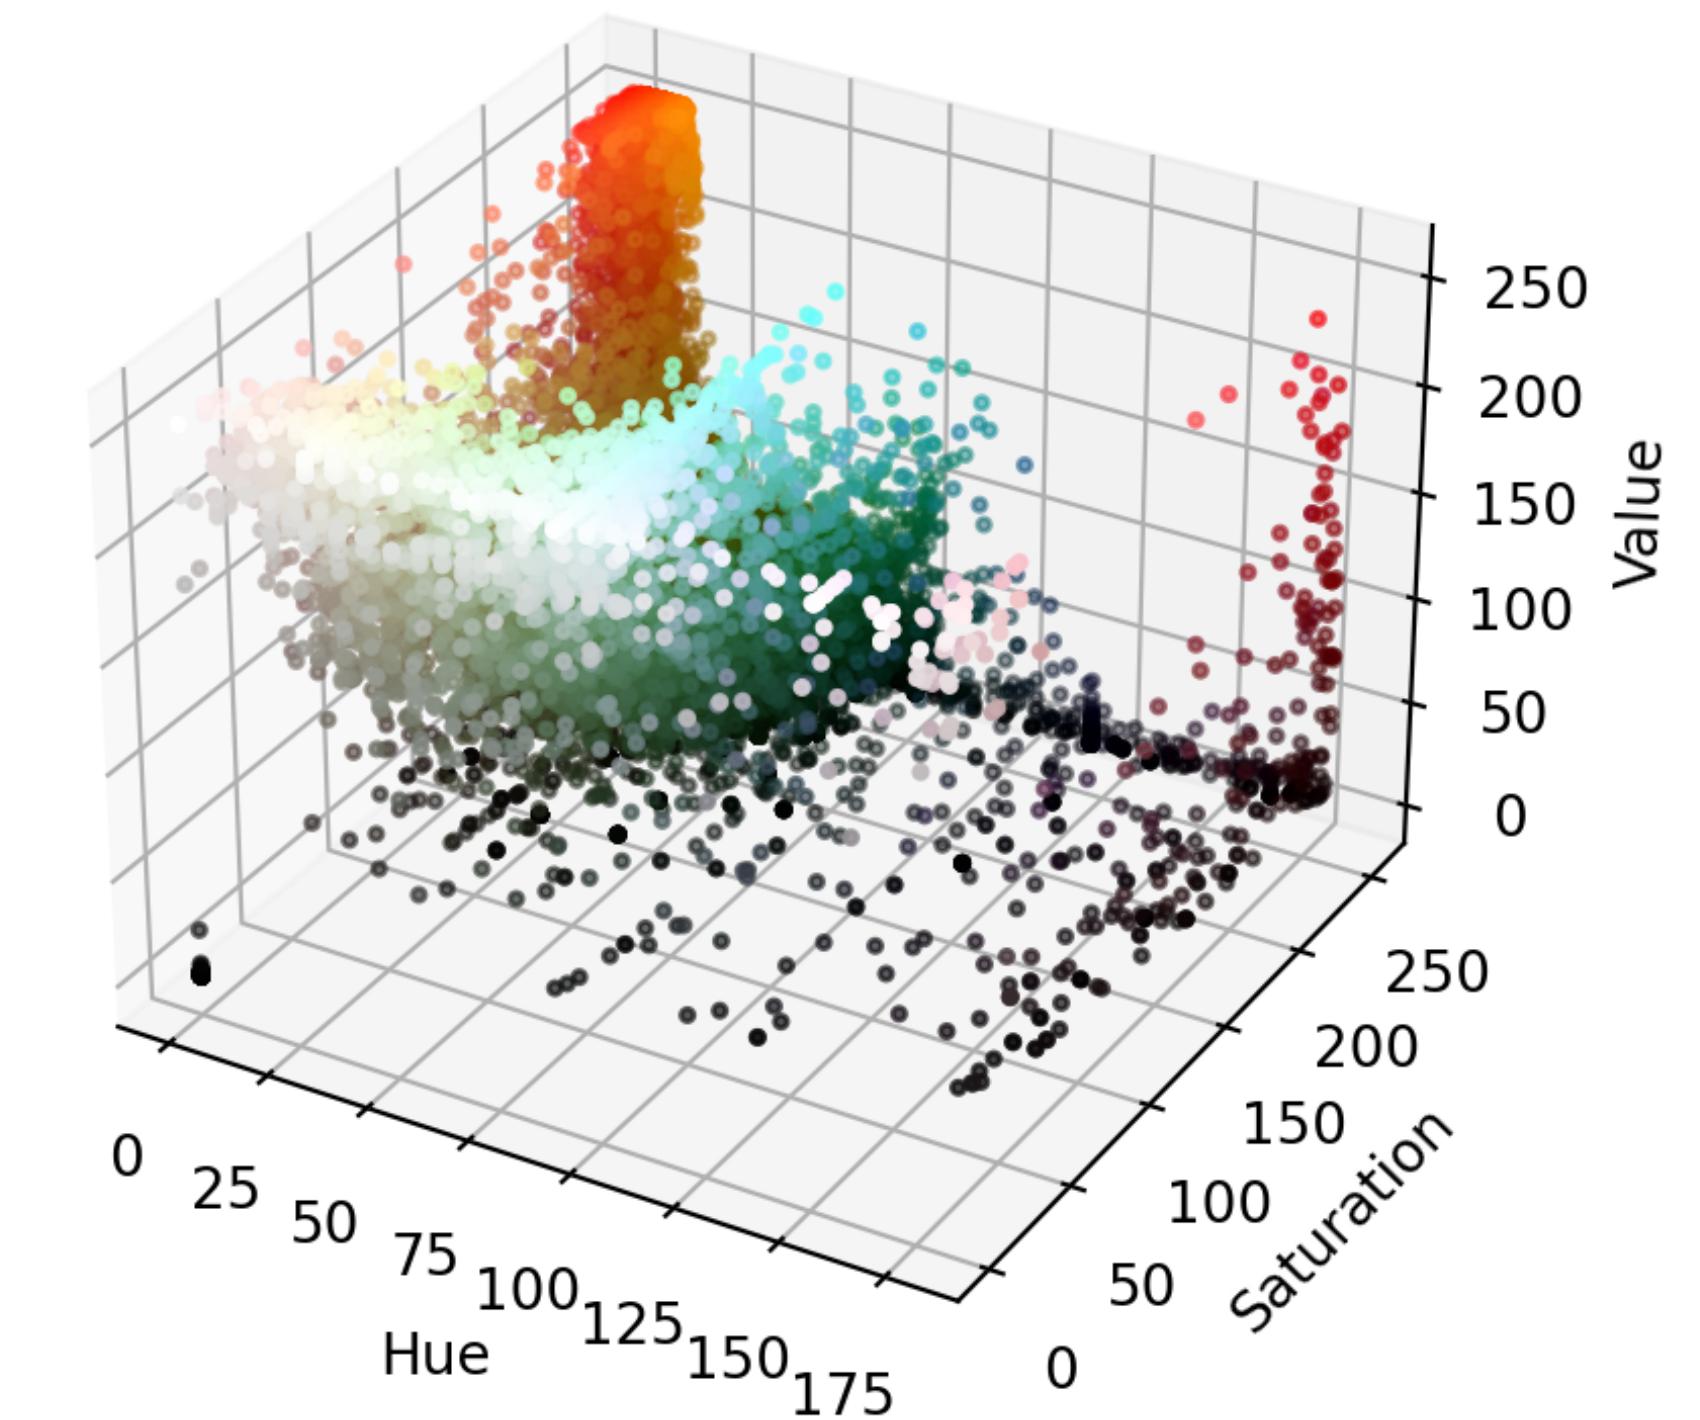


Visualizing in HSV Color Space

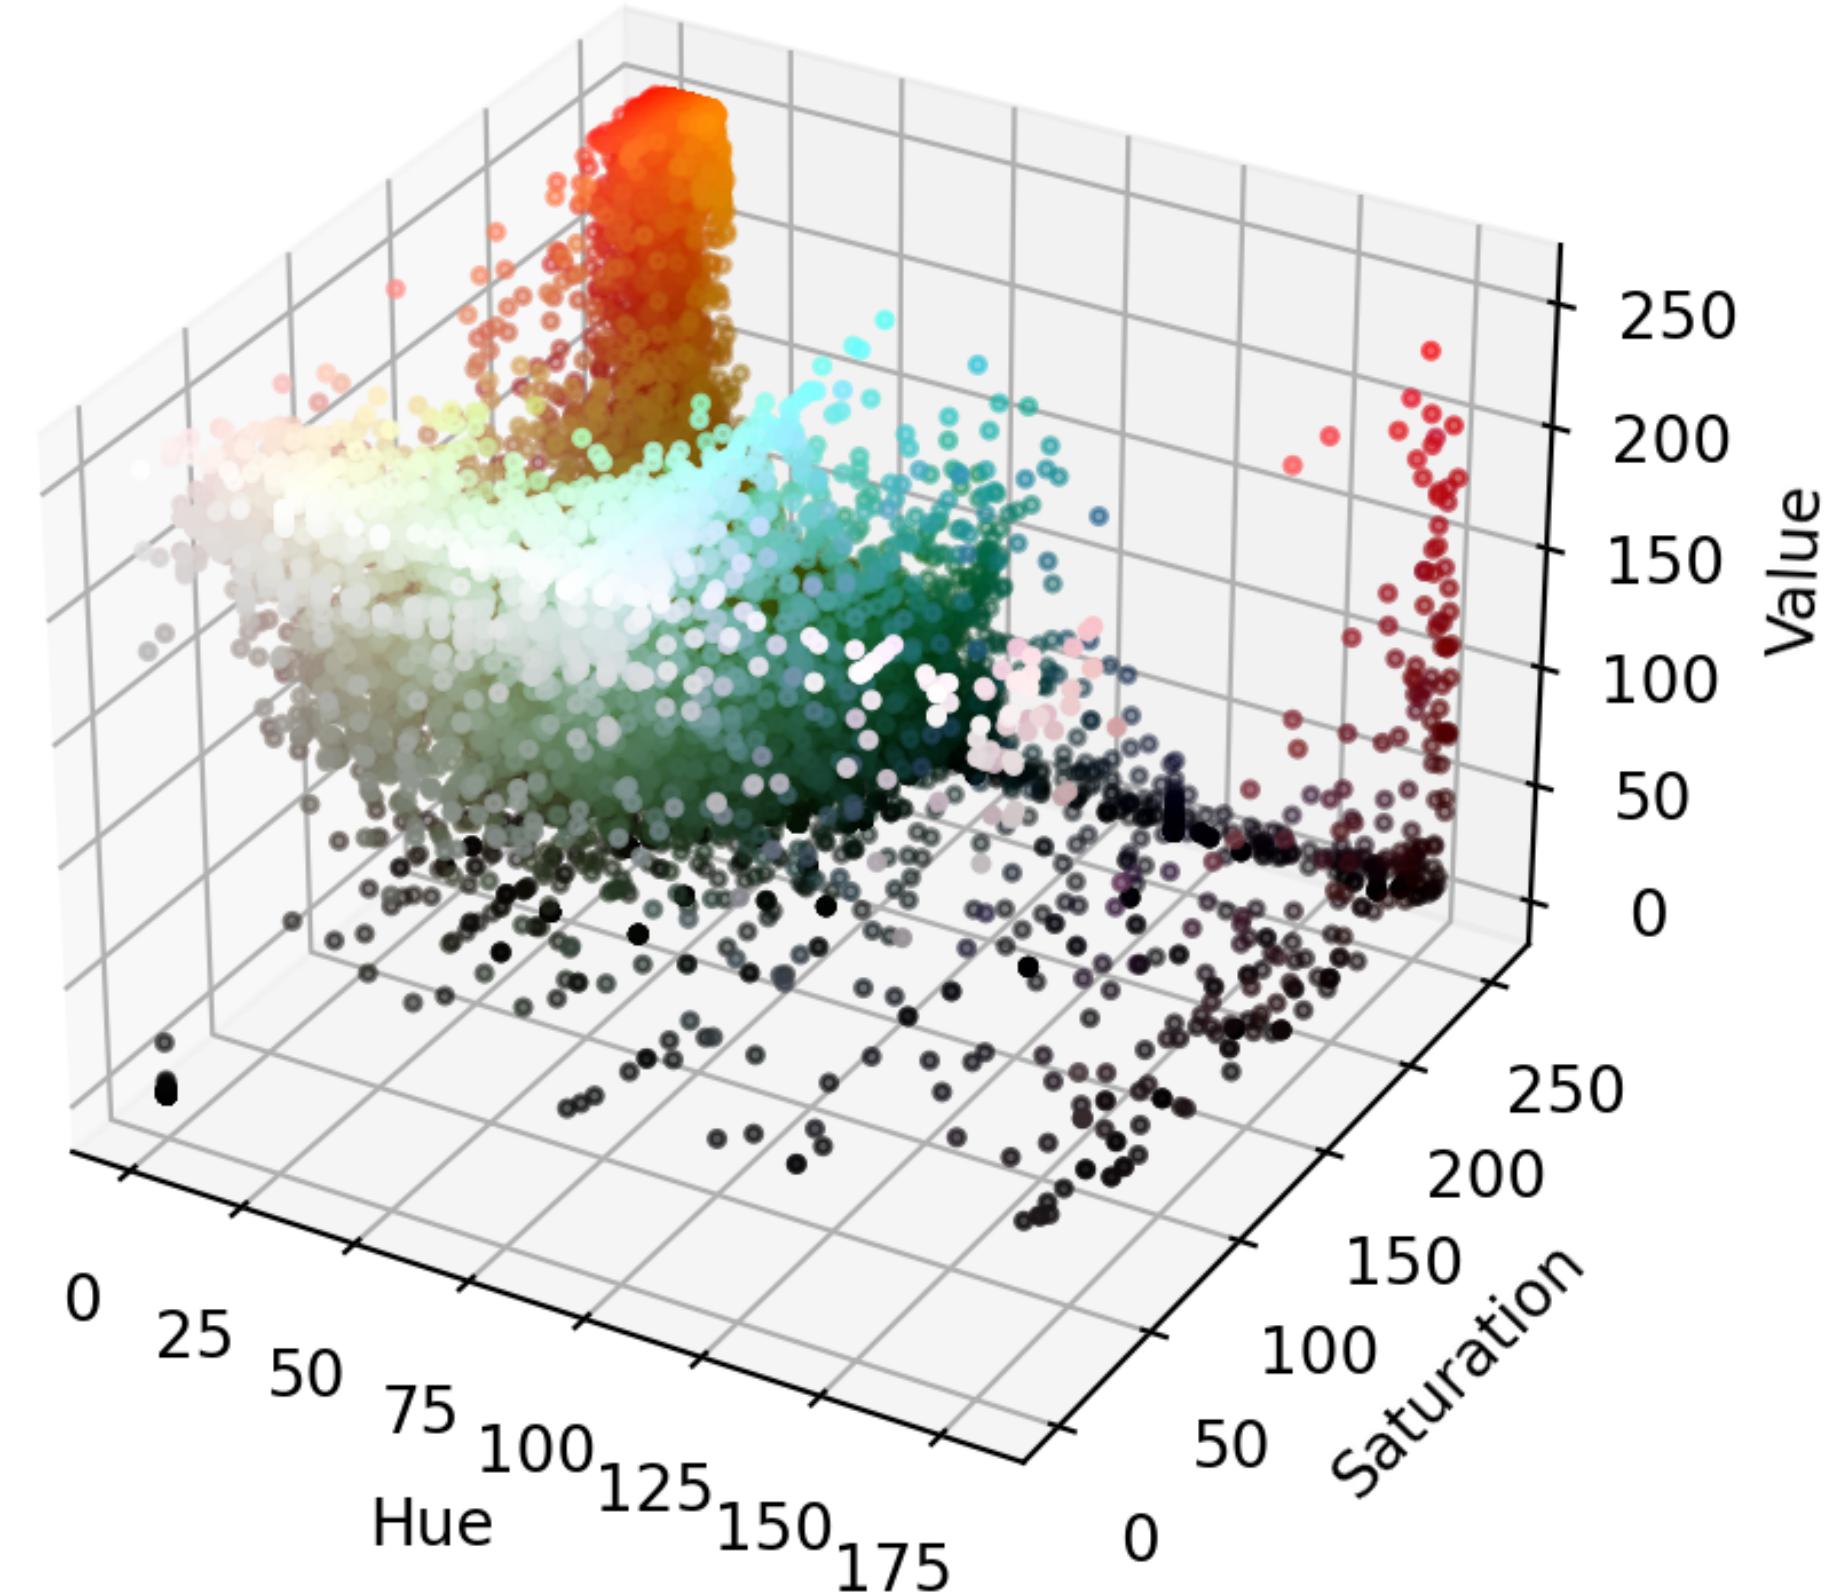
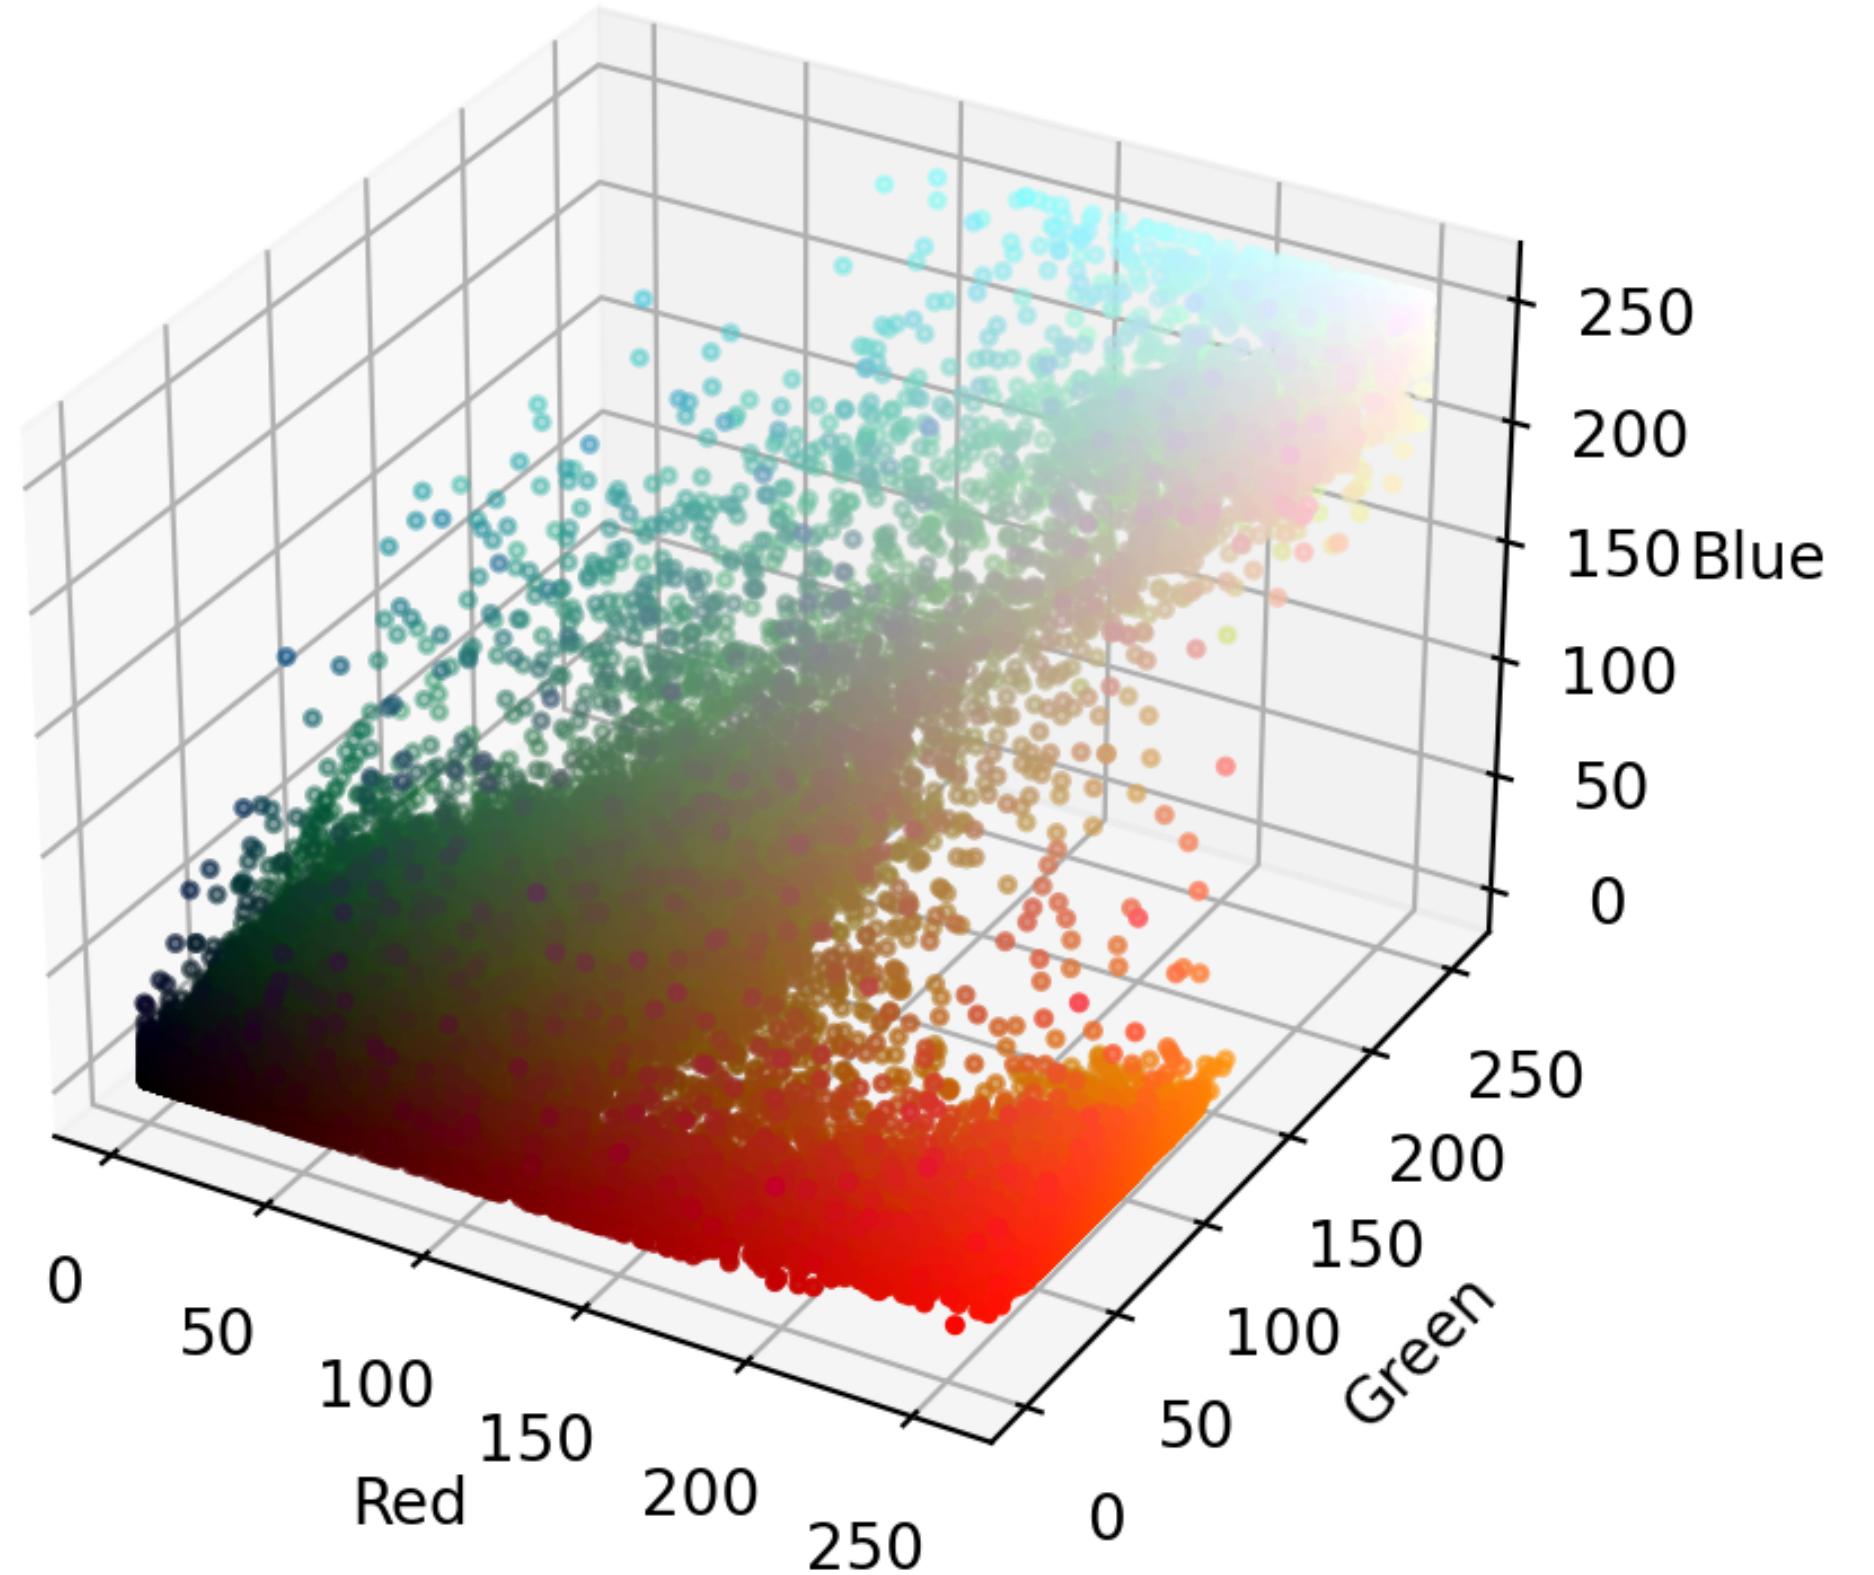
```
hsv_nemo = cv2.cvtColor(nemo, cv2.COLOR_RGB2HSV)
h, s, v = cv2.split(hsv_nemo)

fig = plt.figure()
axis = fig.add_subplot(1, 1, 1, projection="3d")

axis.scatter(
    h.flatten(), s.flatten(), v.flatten(), facecolors=pixel_colors, marker="."
)
axis.set_xlabel("Hue")
axis.set_ylabel("Saturation")
axis.set_zlabel("Value")
plt.show()
```



RGB vs HSL



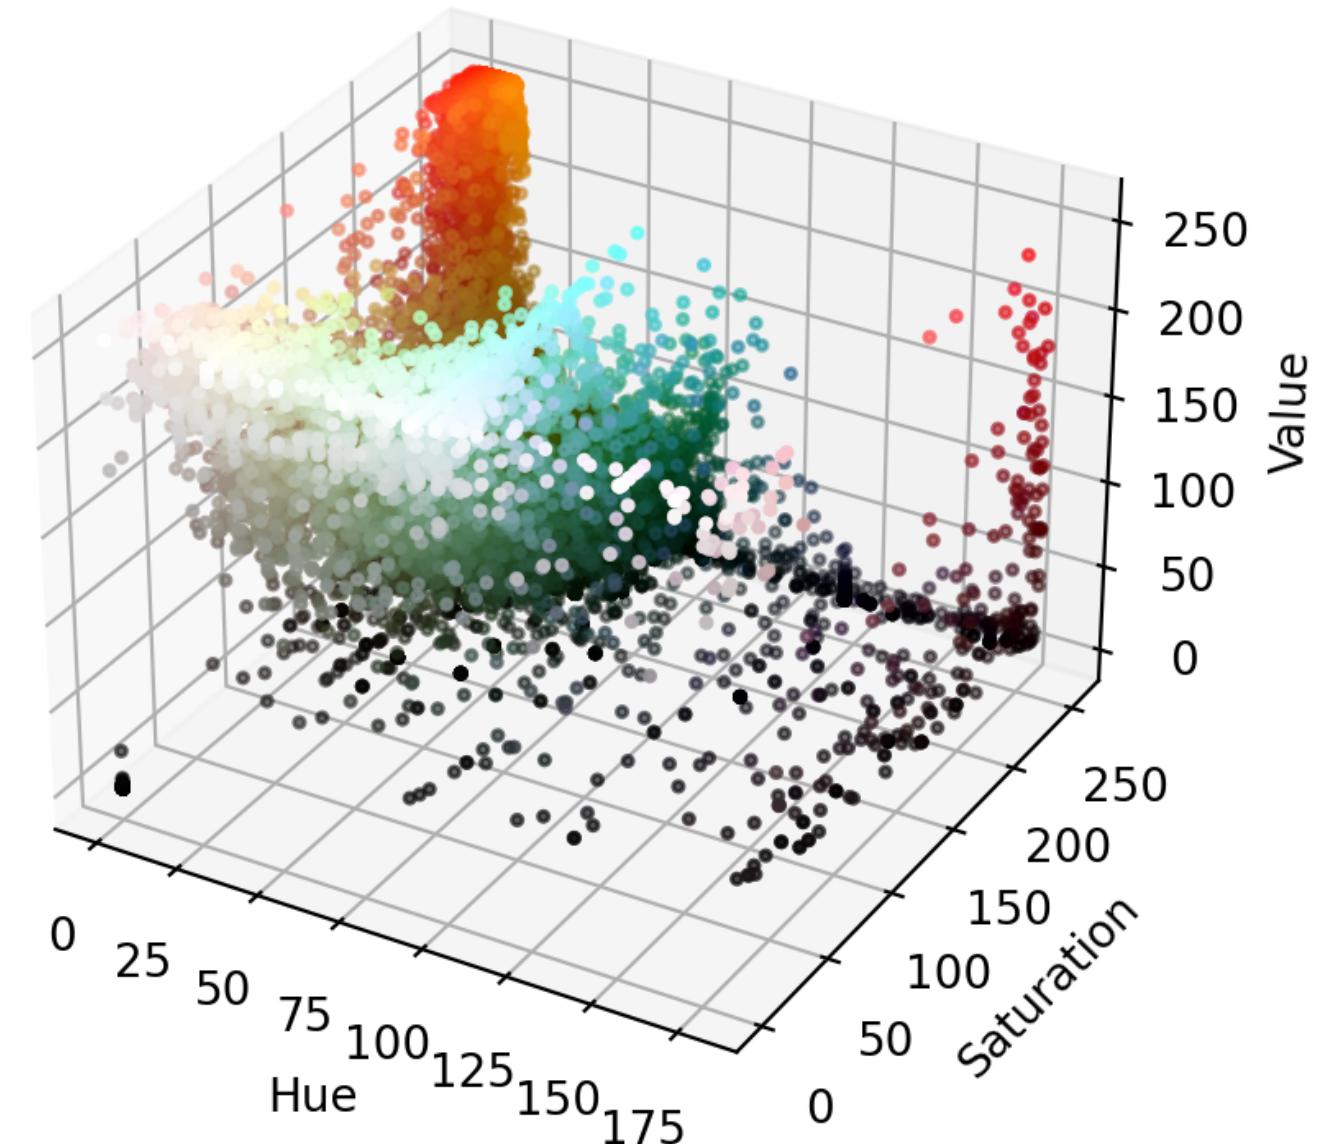
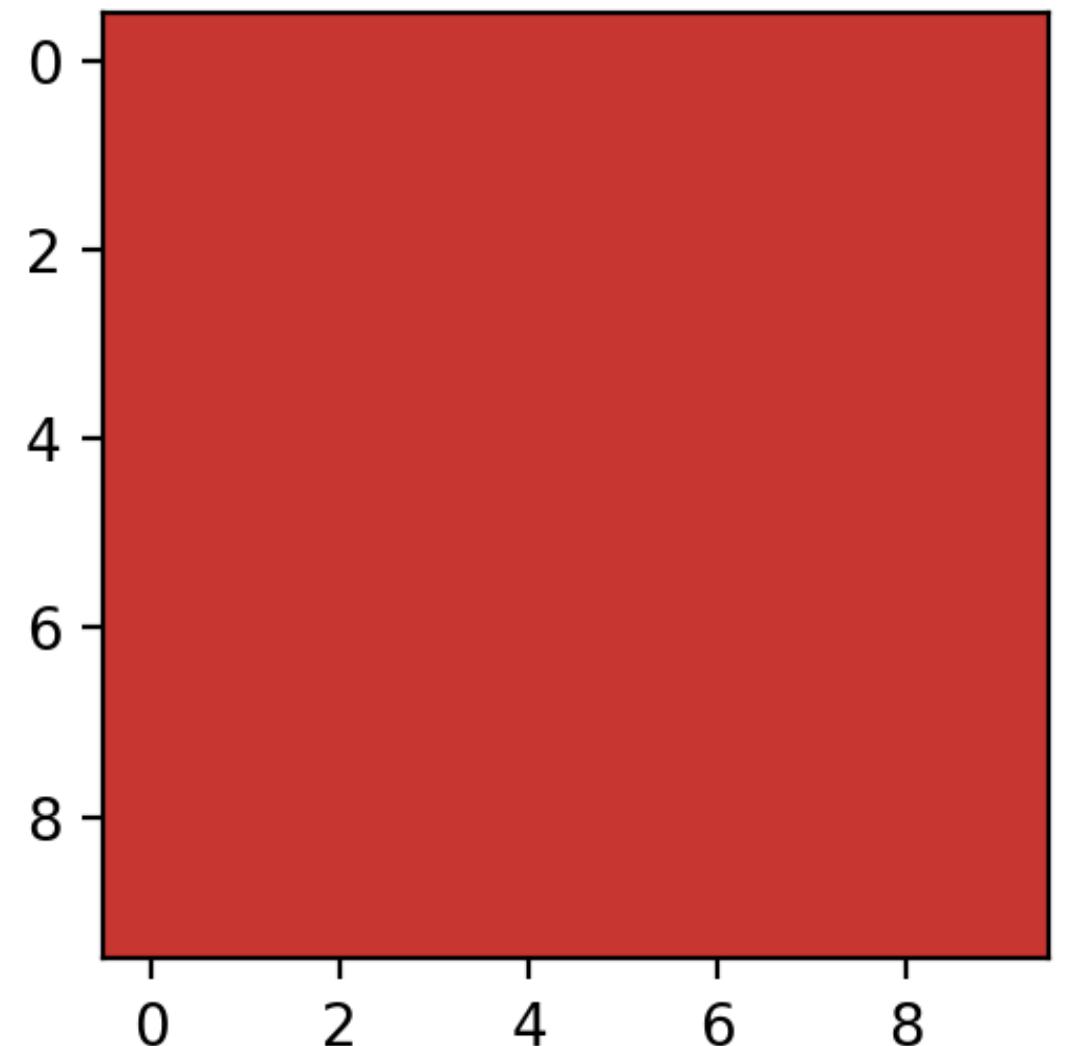
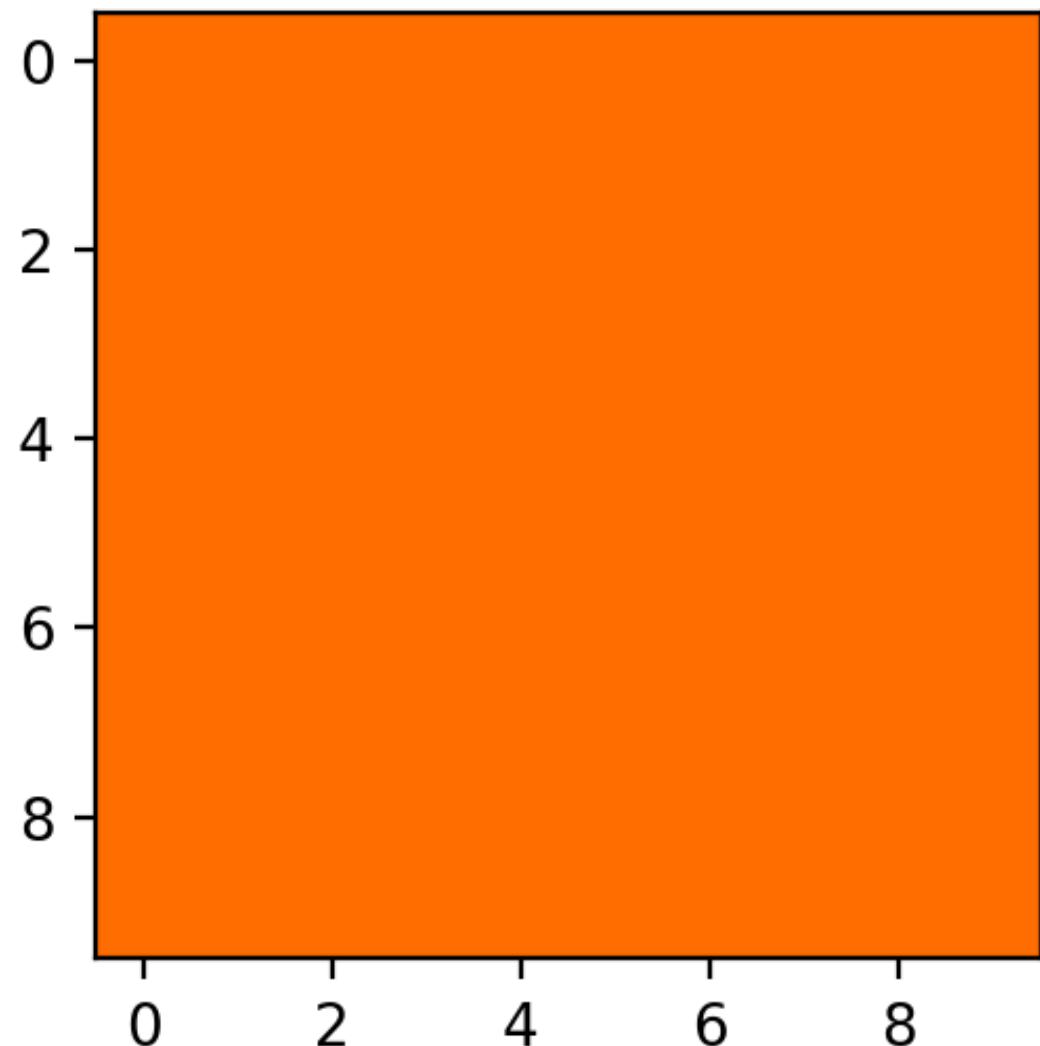
Picking Out a Range

```
light_orange = (1, 190, 200)
dark_orange = (18, 255, 255)

# Normalise to 0 - 1 range for viewing

lo_square = np.full((10, 10, 3), light_orange, dtype=np.uint8) / 255.0
do_square = np.full((10, 10, 3), dark_orange, dtype=np.uint8) / 255.0

plt.subplot(1, 2, 1)
plt.imshow(hsv_to_rgb(do_square))
plt.subplot(1, 2, 2)
plt.imshow(hsv_to_rgb(lo_square))
plt.show()
```



Picking Out a Range

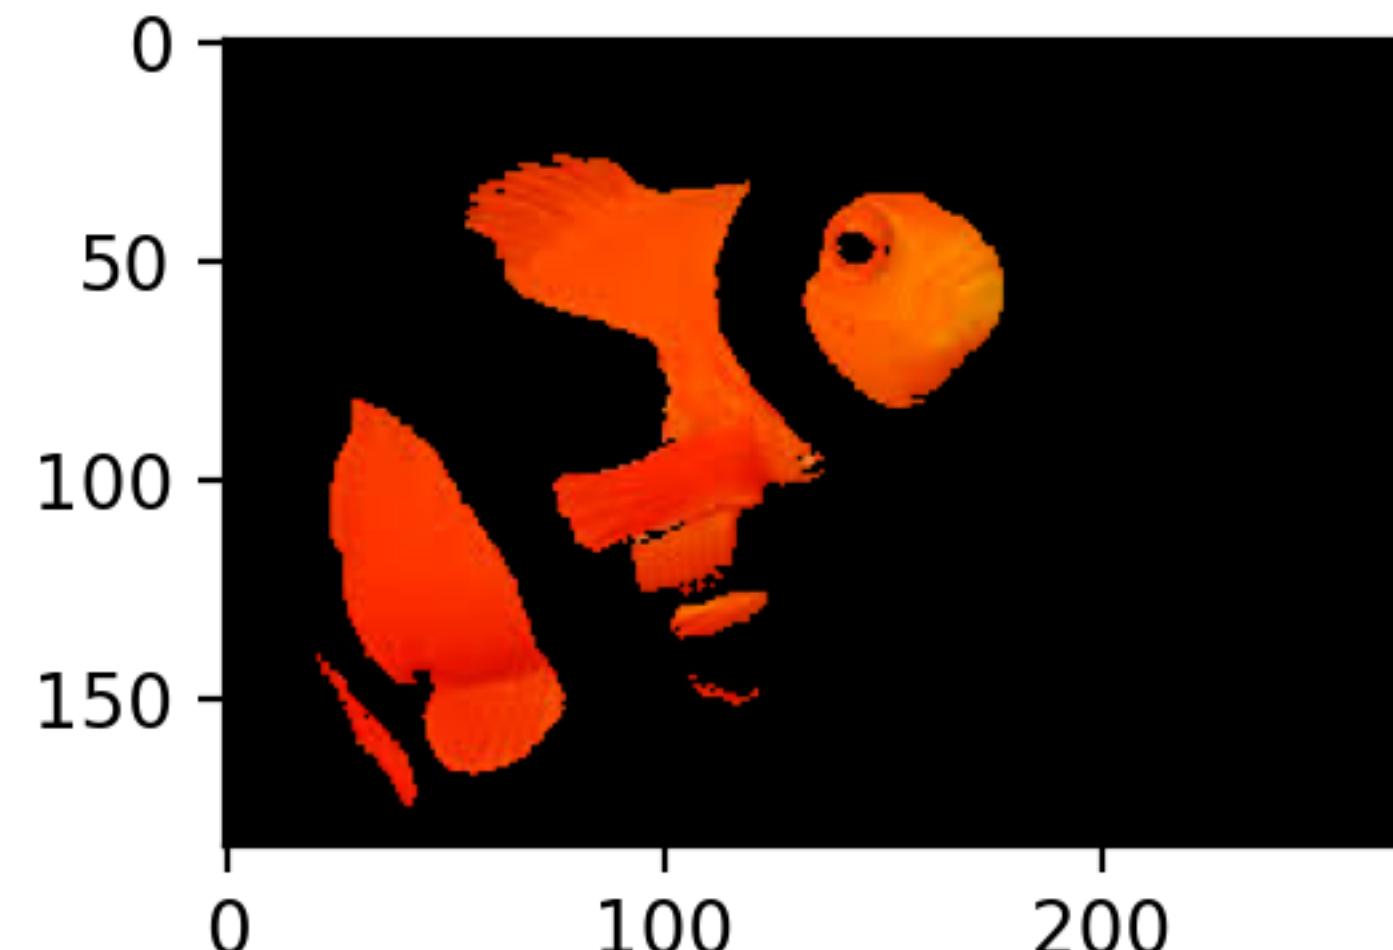
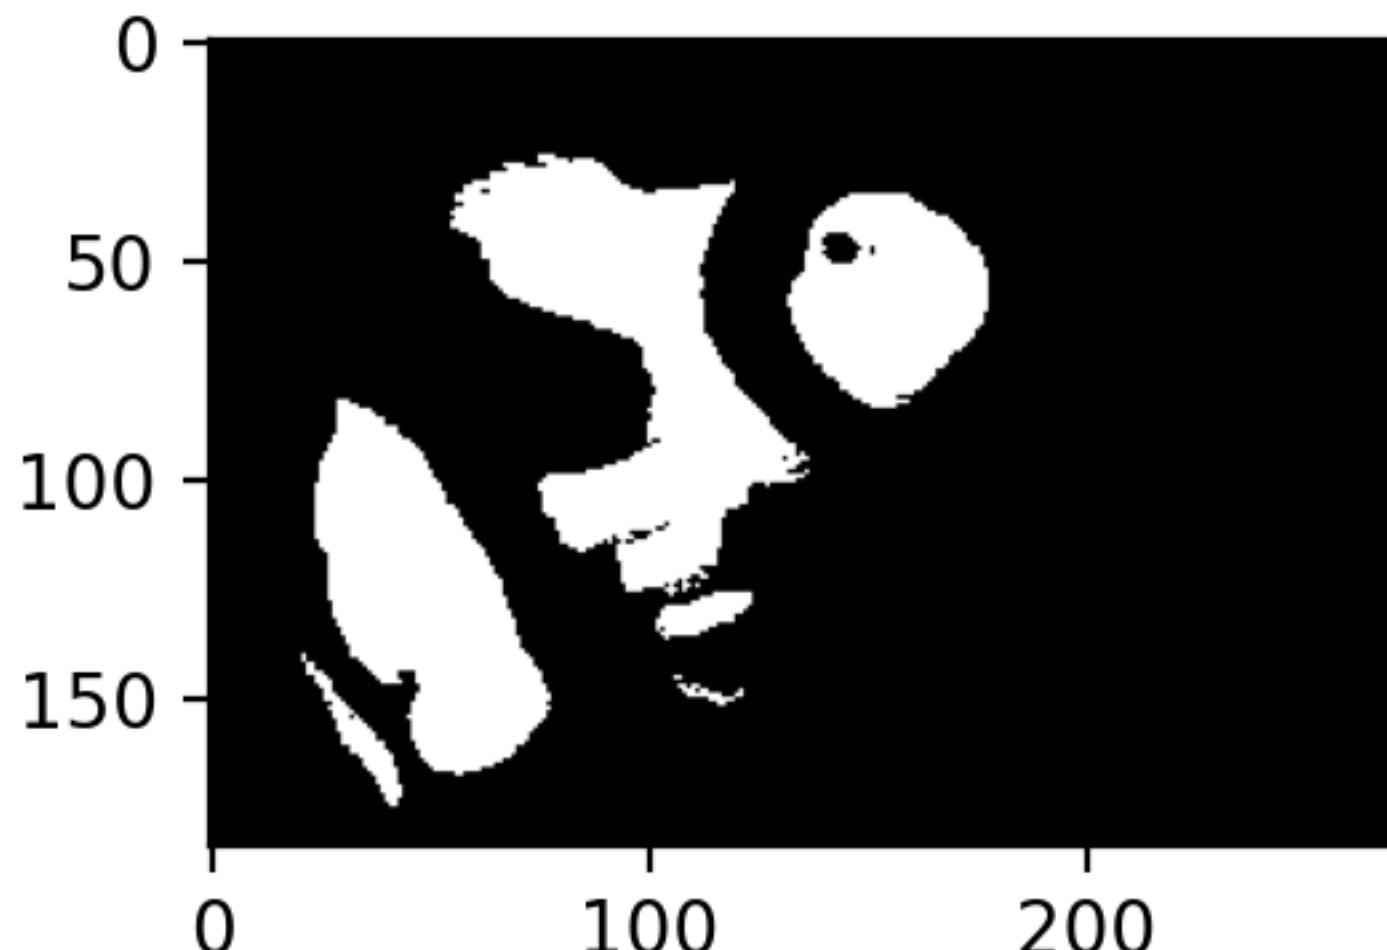
```
mask = cv2.inRange(hsv_nemo, light_orange, dark_orange)

# Bitwise-AND mask and original image

result = cv2.bitwise_and(nemo, nemo, mask=mask)

# Convert back to RGB in order to plot using `matplotlib.pyplot`

plt.subplot(1, 2, 1)
plt.imshow(mask, cmap="gray")
plt.subplot(1, 2, 2)
plt.imshow(result)
plt.show()
```

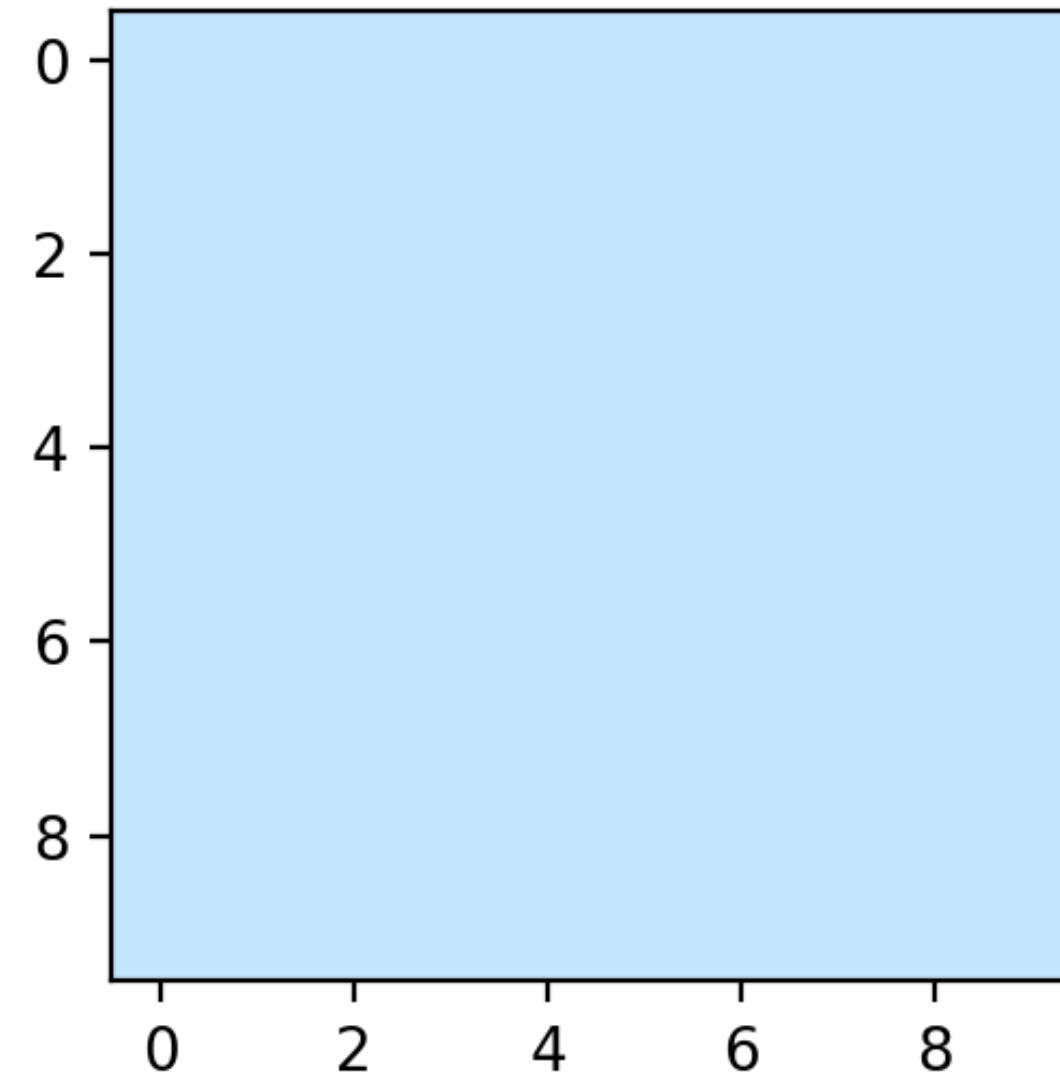
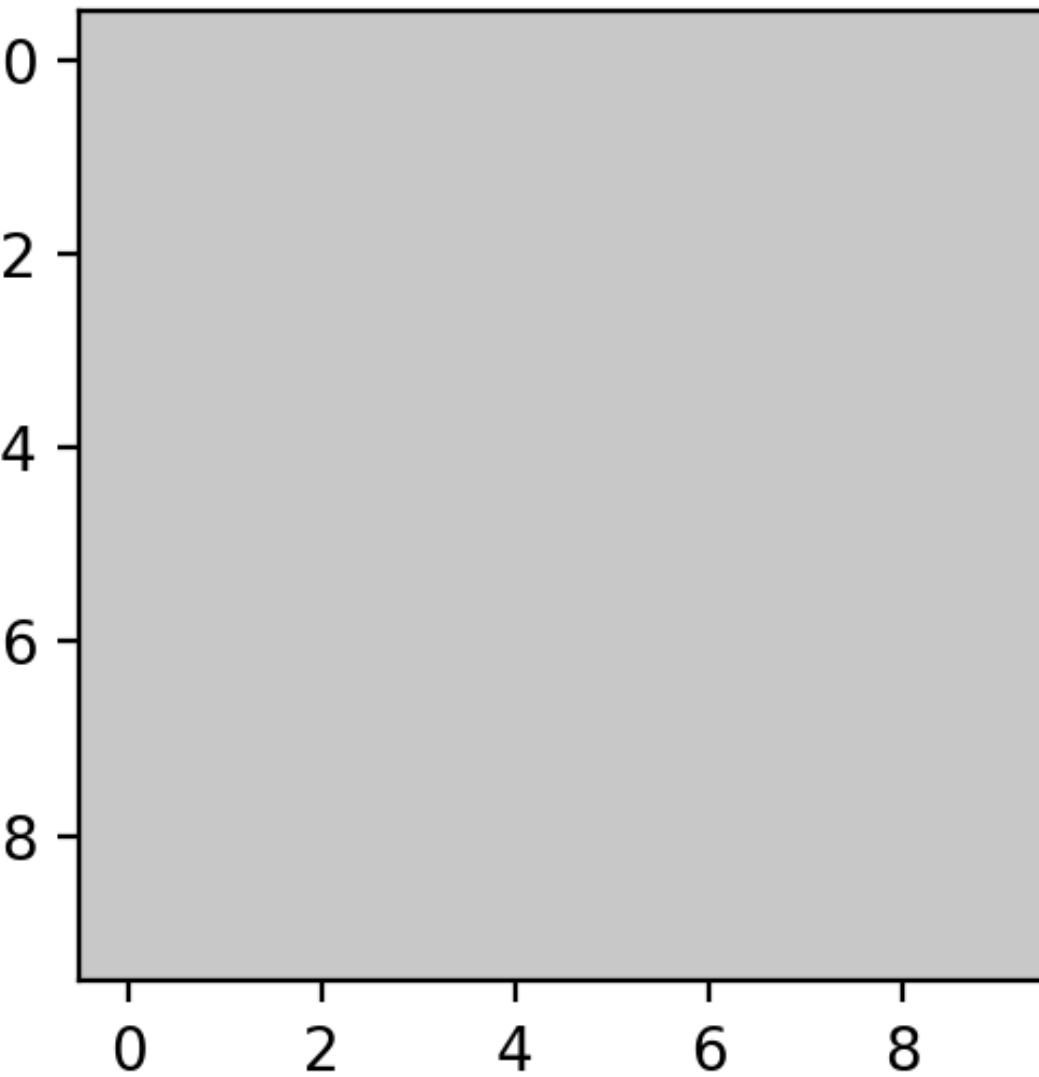


Adding a Second Mask

```
light_white = (0, 0, 200)
dark_white = (145, 60, 255)

lw_square = np.full((10, 10, 3), light_white, dtype=np.uint8) / 255.0
dw_square = np.full((10, 10, 3), dark_white, dtype=np.uint8) / 255.0

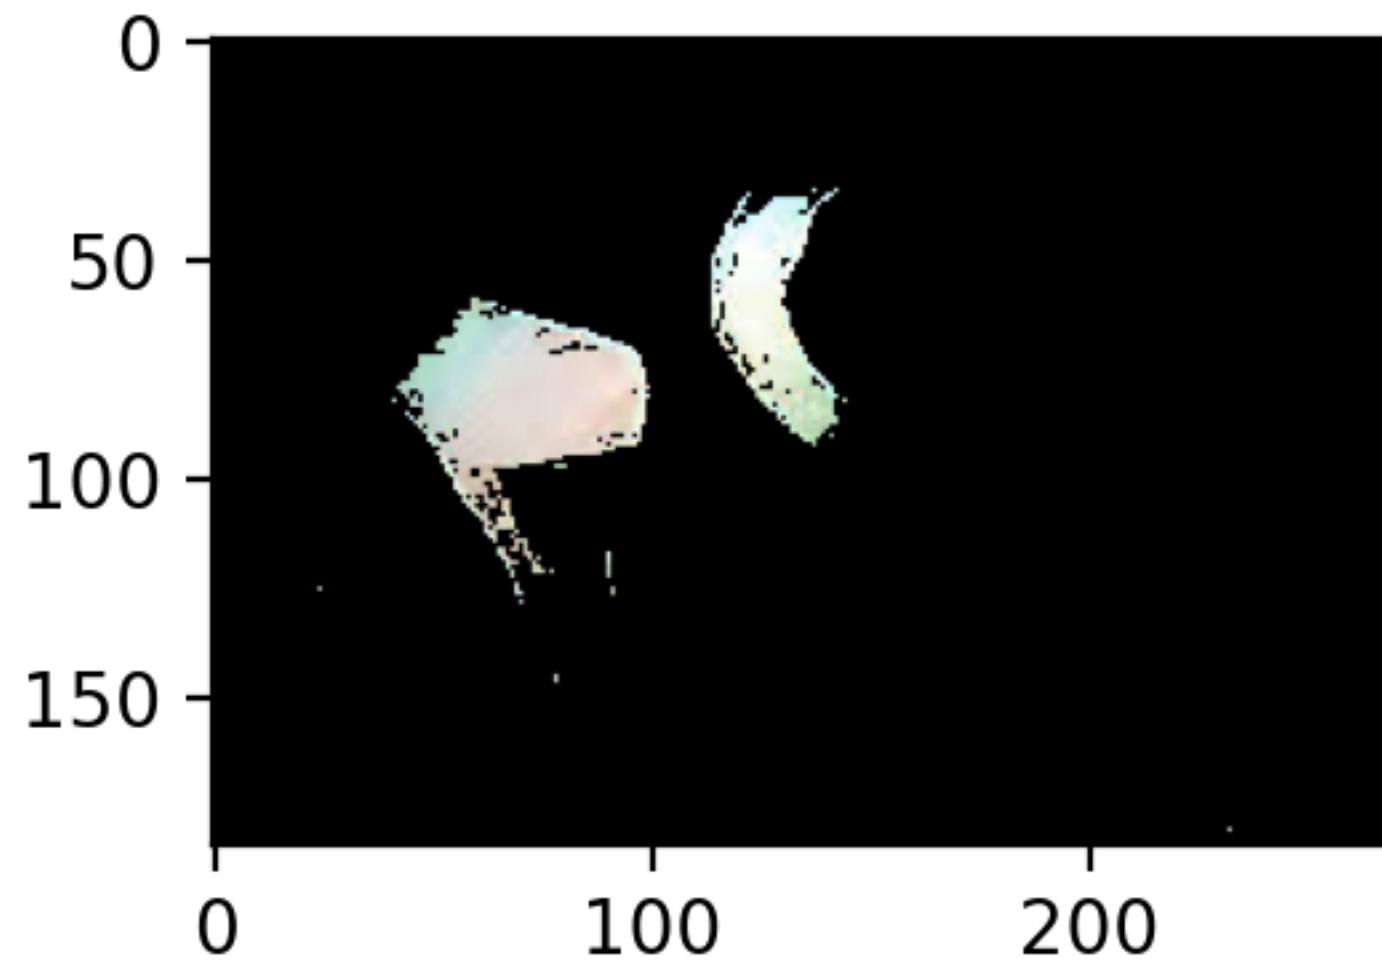
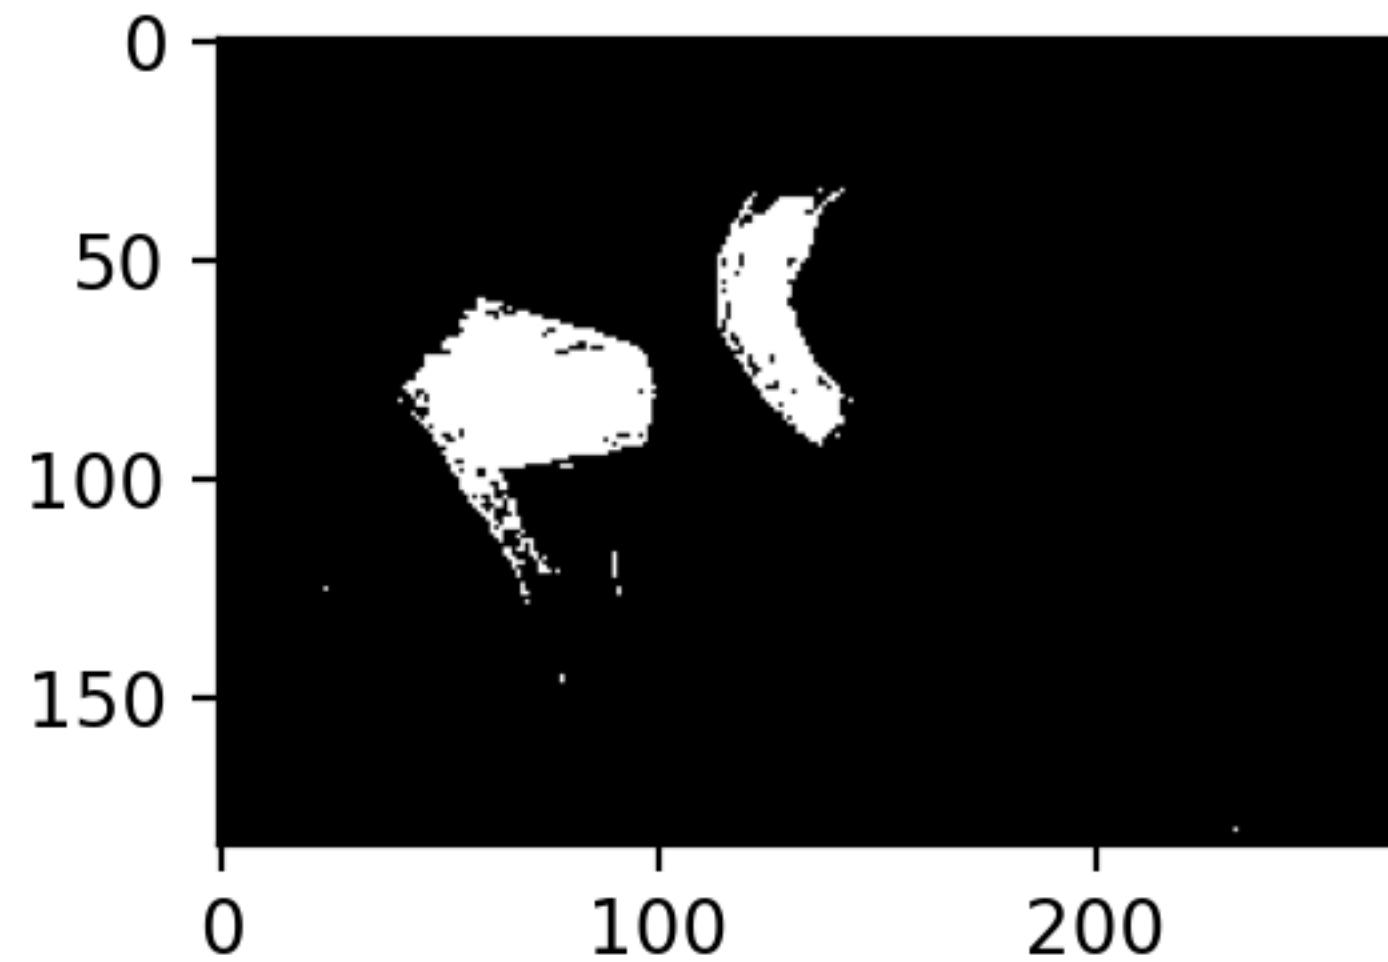
plt.subplot(1, 2, 1)
plt.imshow(hsv_to_rgb(lw_square))
plt.subplot(1, 2, 2)
plt.imshow(hsv_to_rgb(dw_square))
plt.show()
```



Adding a Second Mask

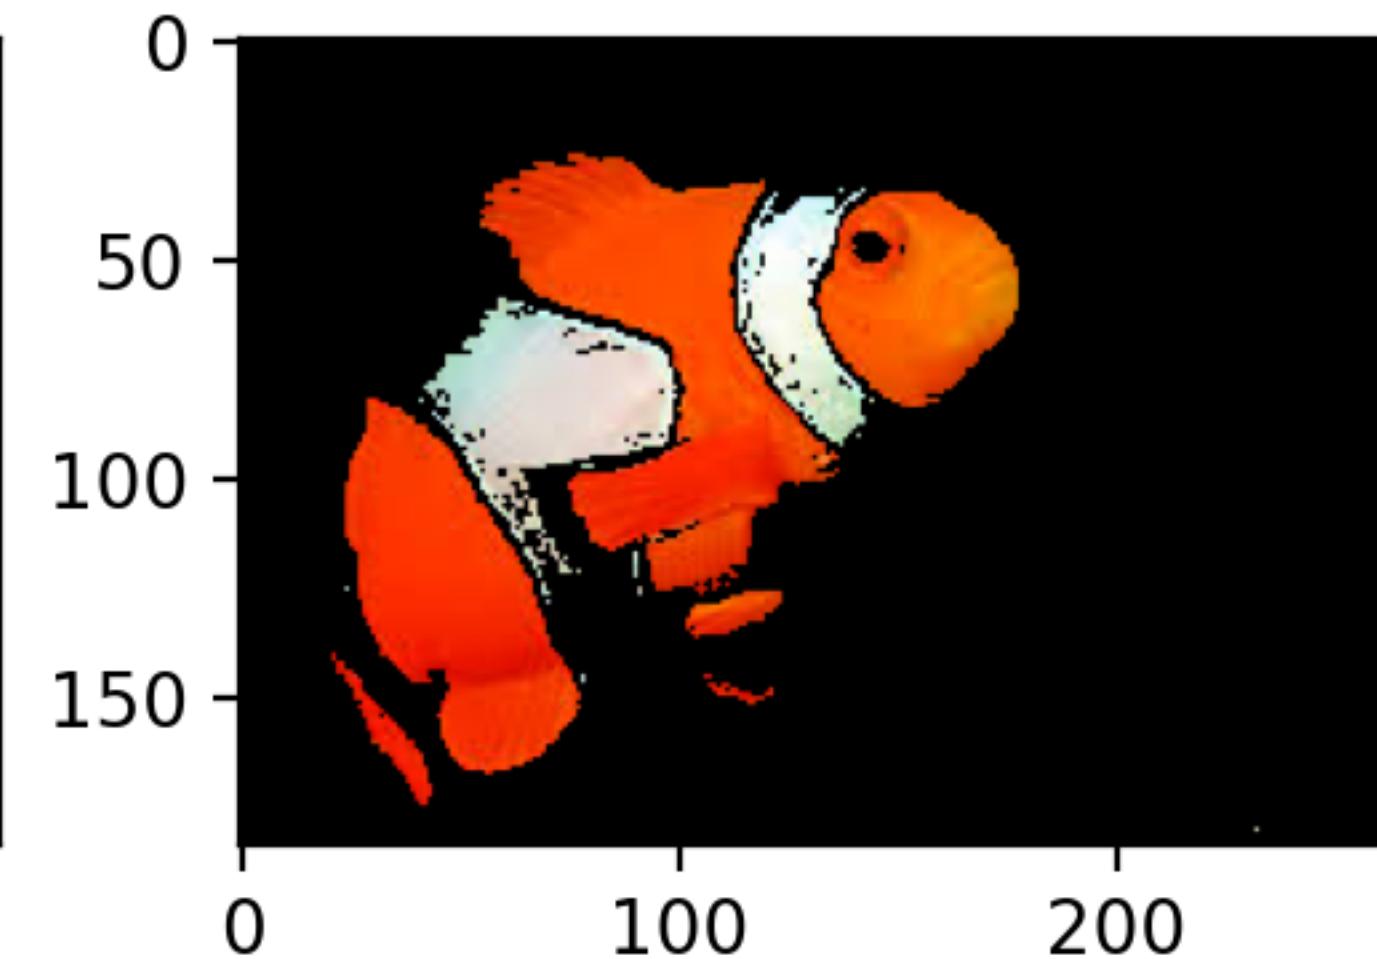
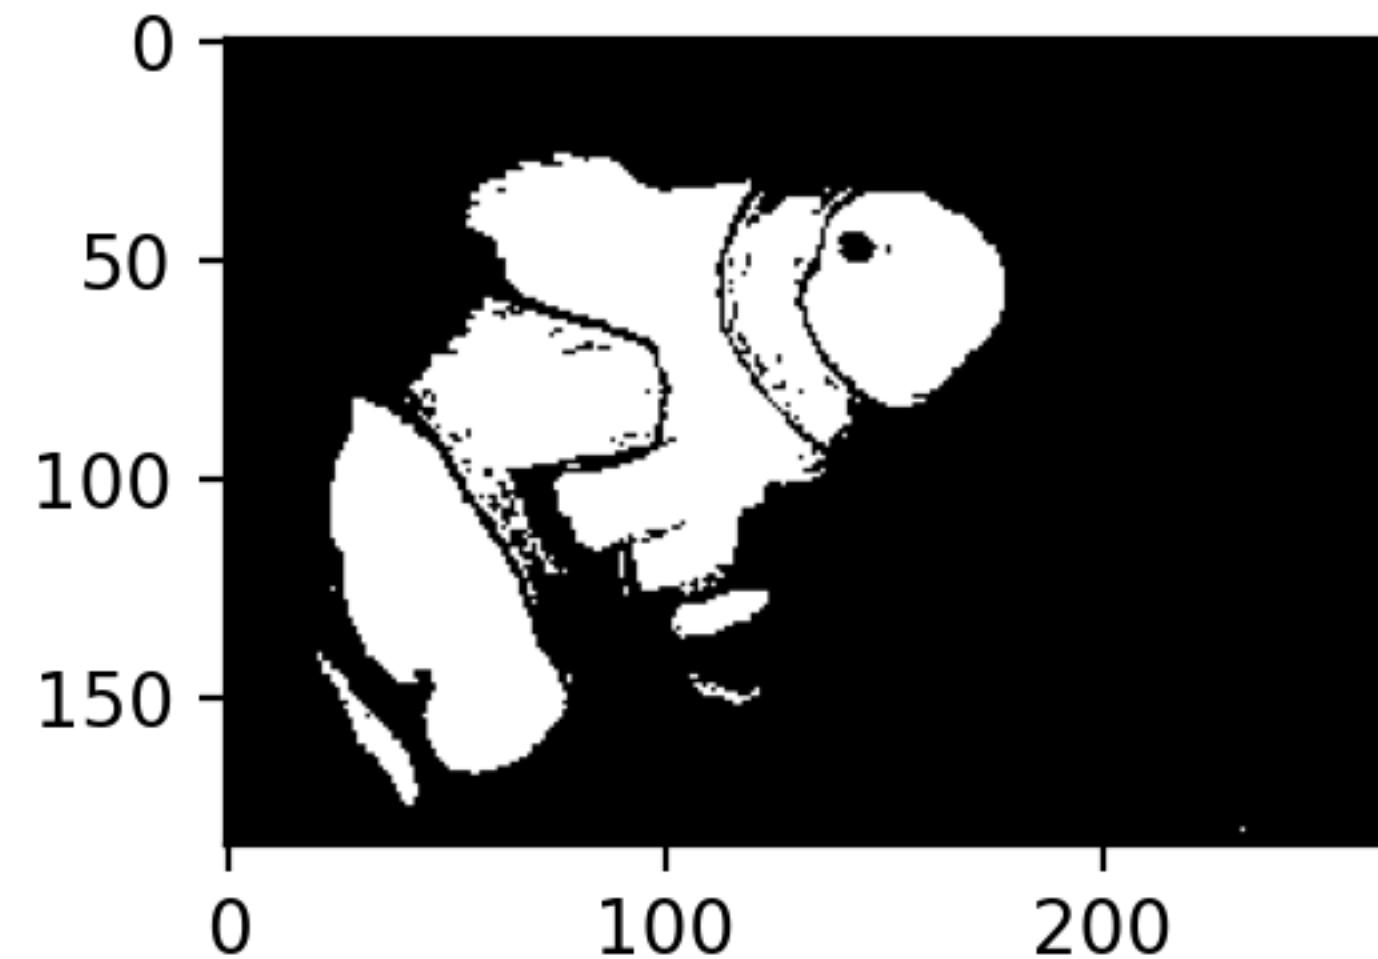
```
mask_white = cv2.inRange(hsv_nemo, light_white, dark_white)
result_white = cv2.bitwise_and(nemo, nemo, mask=mask_white)

plt.subplot(1, 2, 1)
plt.imshow(mask_white, cmap="gray")
plt.subplot(1, 2, 2)
plt.imshow(result_white)
plt.show()
```

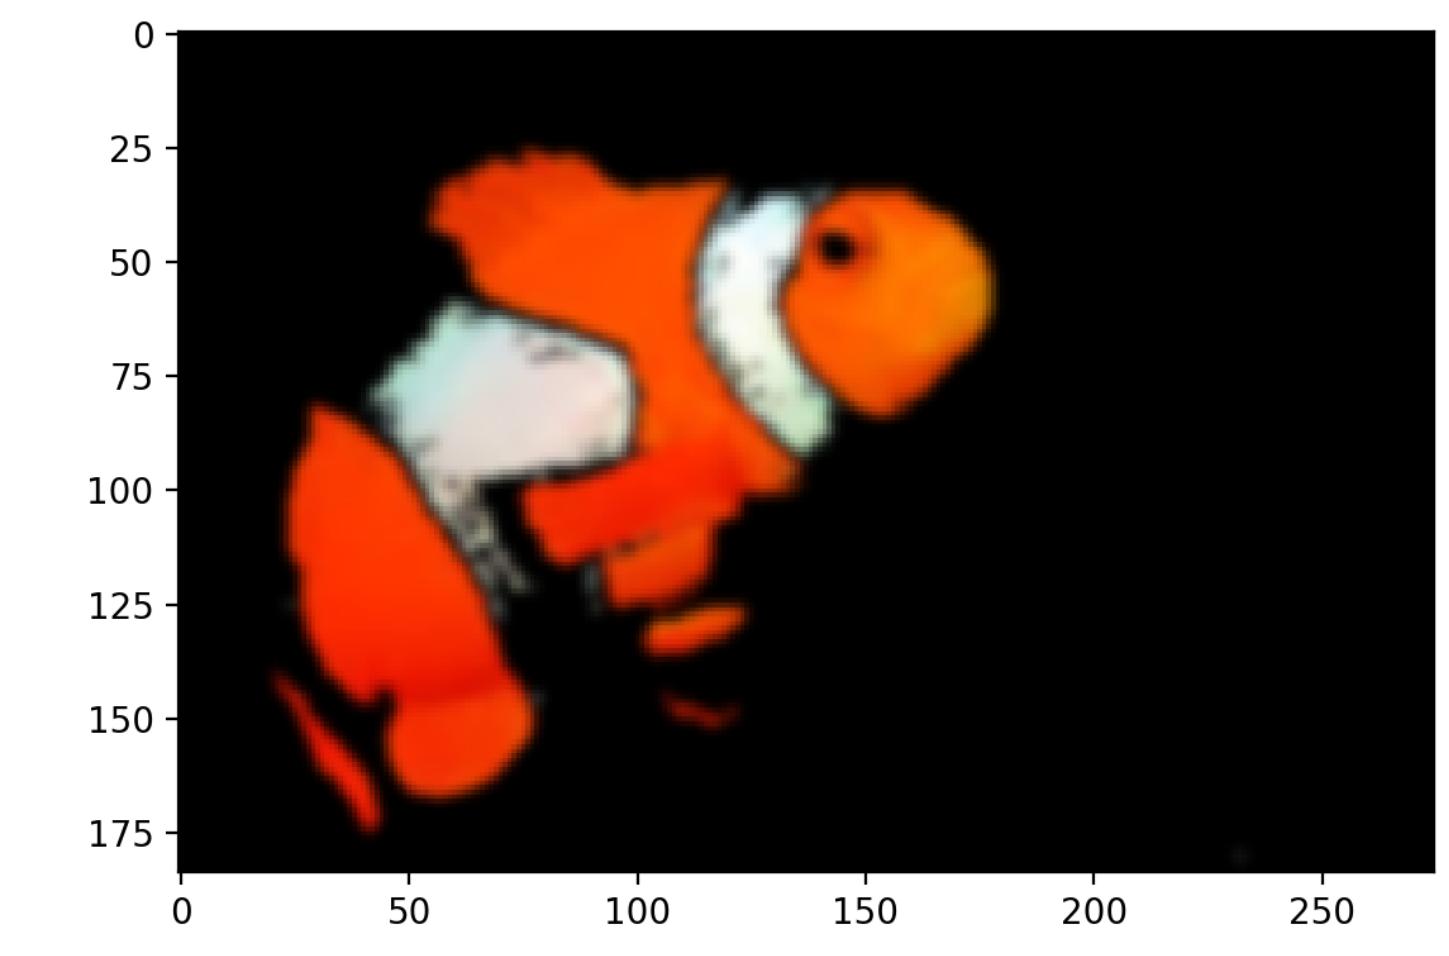
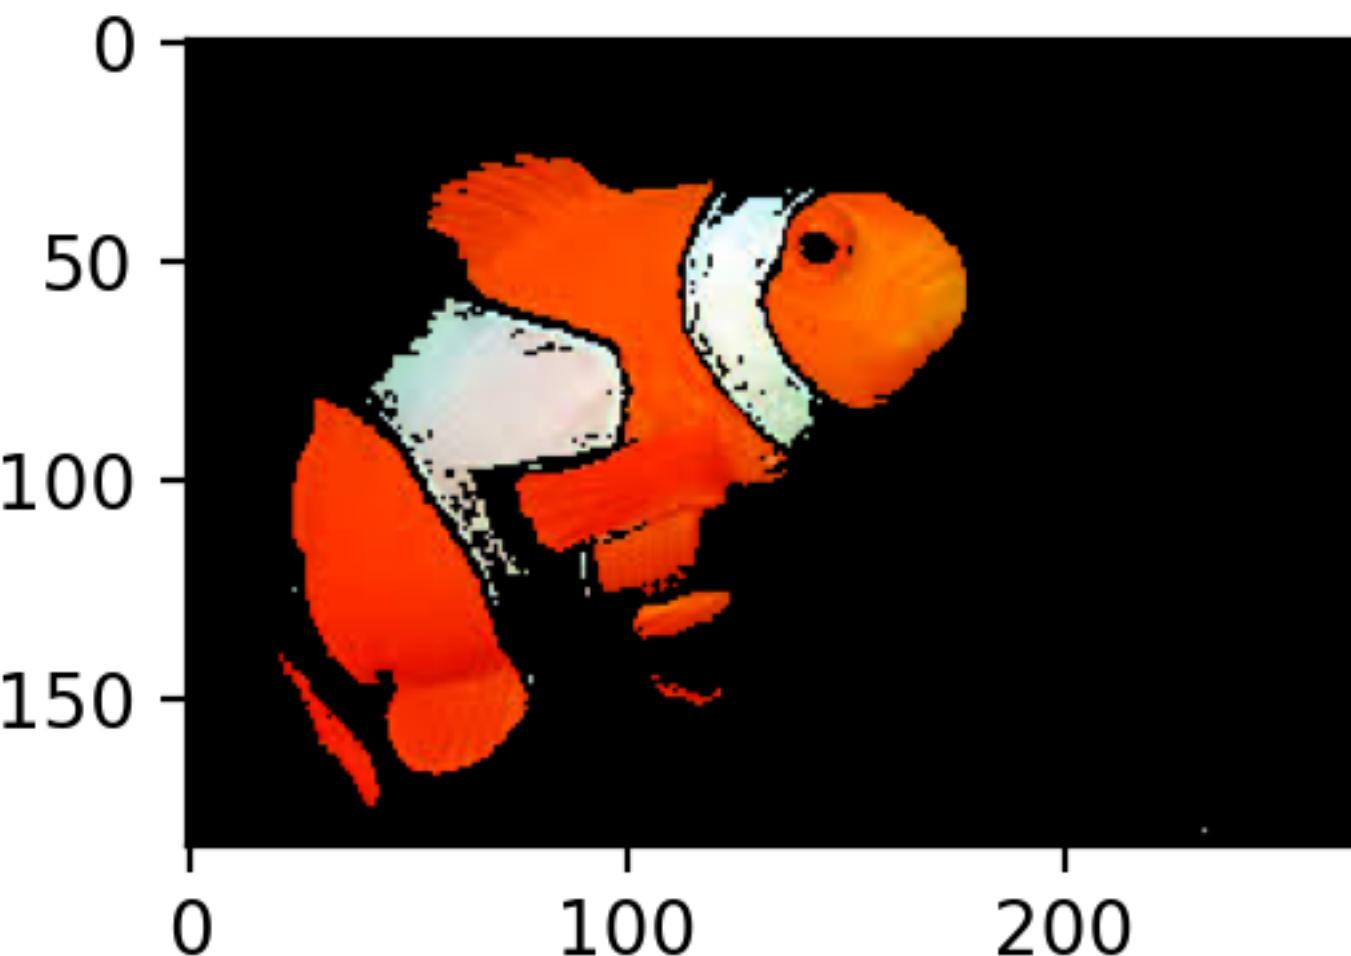
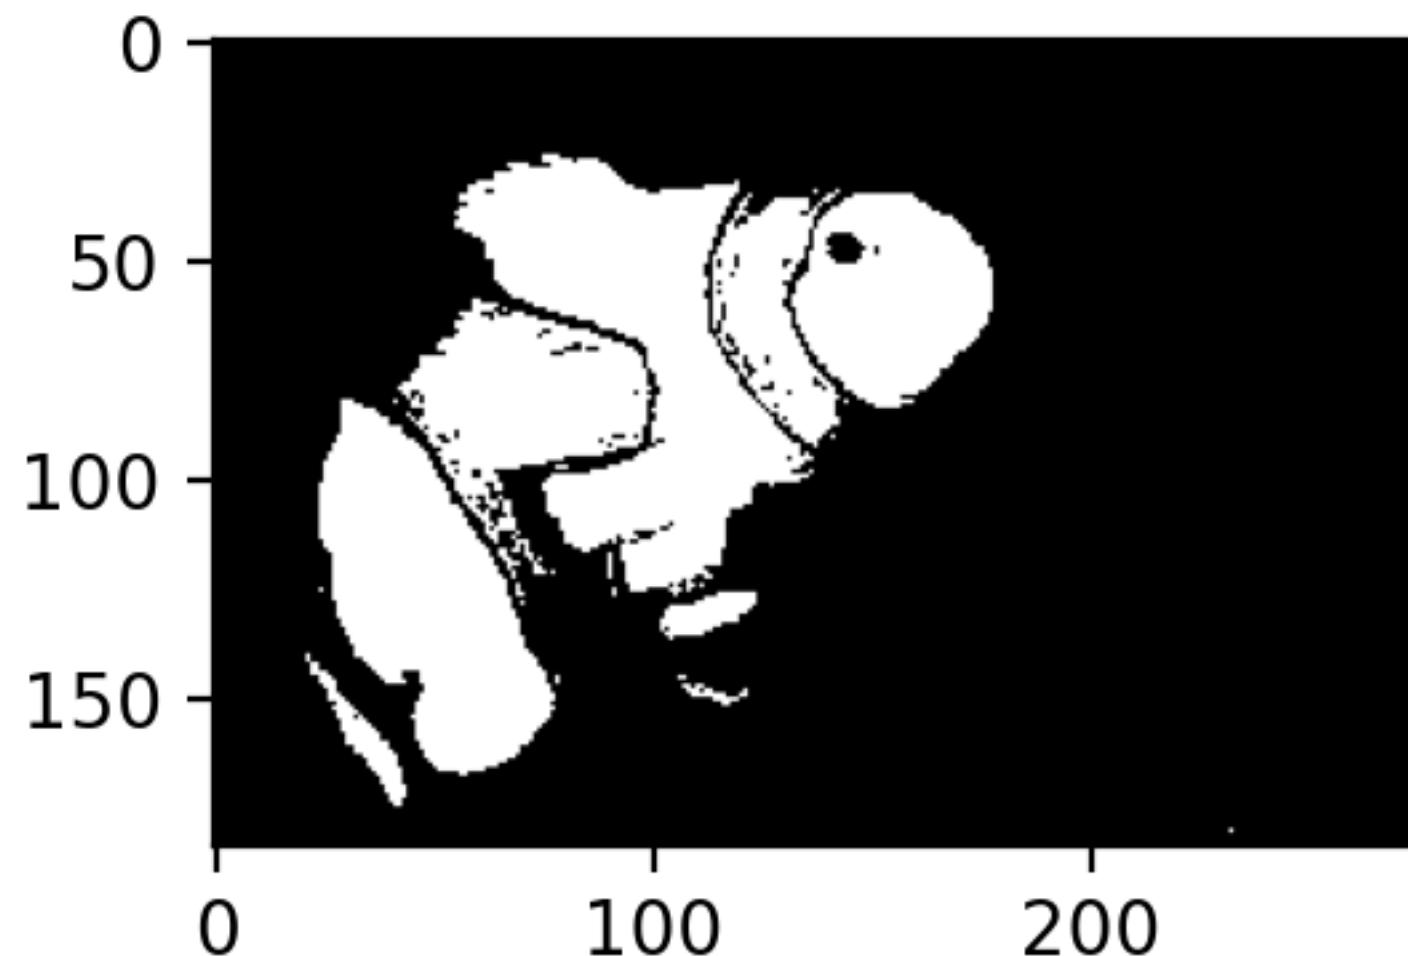


Combining Masks

```
final_mask = mask + mask_white  
  
final_result = cv2.bitwise_and(nemo, nemo, mask=final_mask)  
  
plt.subplot(1, 2, 1)  
plt.imshow(final_mask, cmap="gray")  
plt.subplot(1, 2, 2)  
plt.imshow(final_result)  
plt.show()
```



Gaussian Blur



```
blur = cv2.GaussianBlur(final_result, (7, 7), 0)

plt.imshow(blur)
plt.show()
```

Generalizing the Segmentation

```
path = "nemo"

nemos_friends = []
for i in range(6):
    friend = cv2.cvtColor(
        cv2.imread(path + str(i) + ".jpg"), cv2.COLOR_BGR2RGB
    )
    nemos_friends.append(friend)
```

```
def segment_fish(image):
    """Attempts to segment the clown fish out of the provided image."""
    hsv_image = cv2.cvtColor(image, cv2.COLOR_RGB2HSV)
    light_orange = (1, 190, 200)
    dark_orange = (18, 255, 255)
    mask = cv2.inRange(hsv_image, light_orange, dark_orange)
    light_white = (0, 0, 200)
    dark_white = (145, 60, 255)
    mask_white = cv2.inRange(hsv_image, light_white, dark_white)
    final_mask = mask + mask_white
    result = cv2.bitwise_and(image, image, mask=final_mask)
    result = cv2.GaussianBlur(result, (7, 7), 0)
    return result
```

```
results = [segment_fish(friend) for friend in nemos_friends]

for i in range(1, 6):
    plt.subplot(1, 2, 1)
    plt.imshow(nemos_friends[i])
    plt.subplot(1, 2, 2)
    plt.imshow(results[i])
    plt.show()
```

